

The background features a dark blue gradient on the left, transitioning into a large, abstract, curved shape on the right. This shape is composed of various shades of purple and blue, with a bright orange-yellow curved line running along its bottom edge. The overall design is modern and dynamic.

aws SUMMIT

CHICAGO | AUGUST 25, 2022

ANT301

Observing logs and traces with Amazon OpenSearch Service

Jon Handler

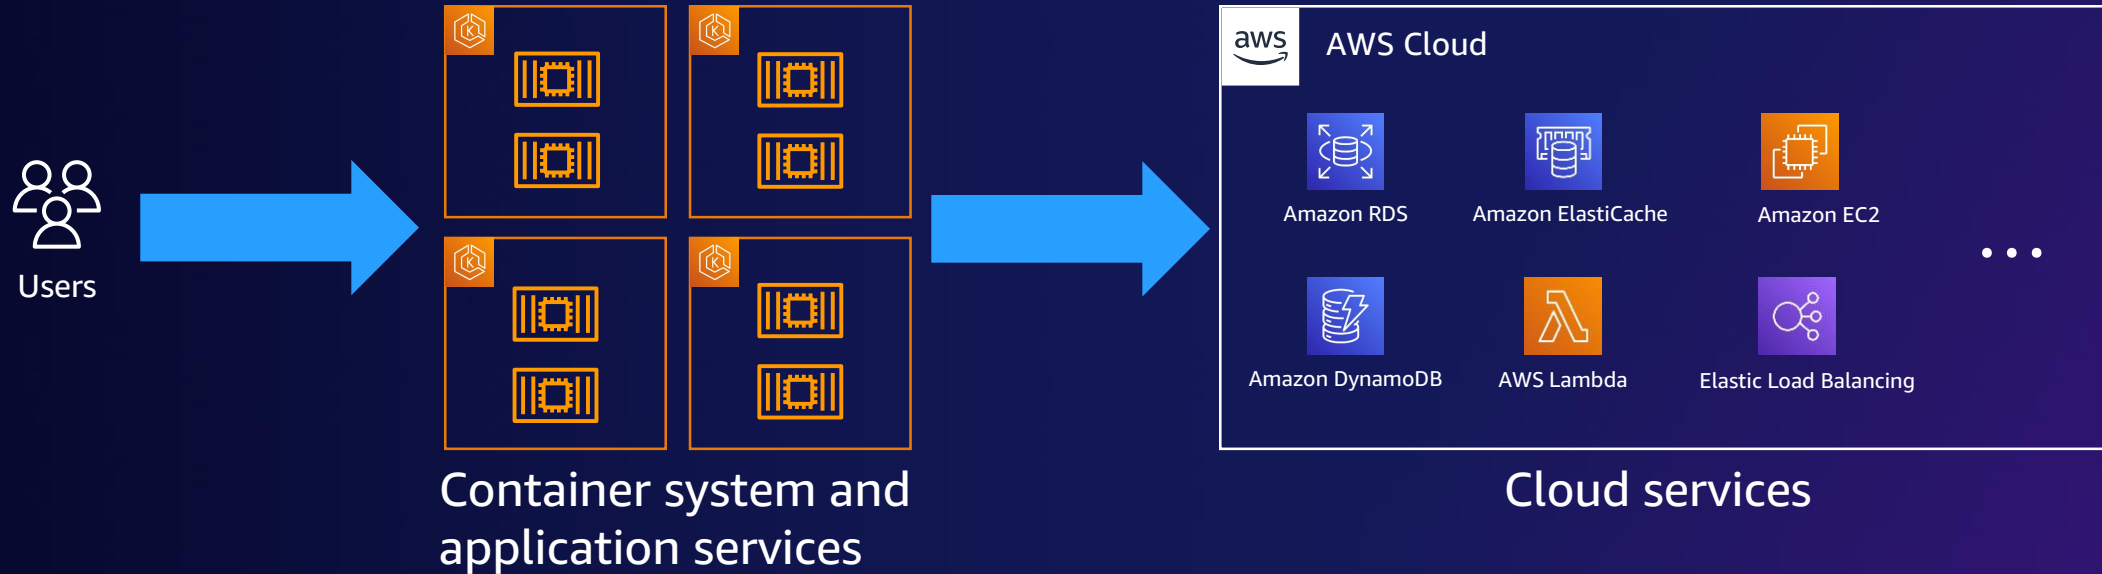
Sr. Principal Specialist Solutions Architect
AWS

Muthu Pitchaimani

Sr. Specialist Solutions Architect
AWS



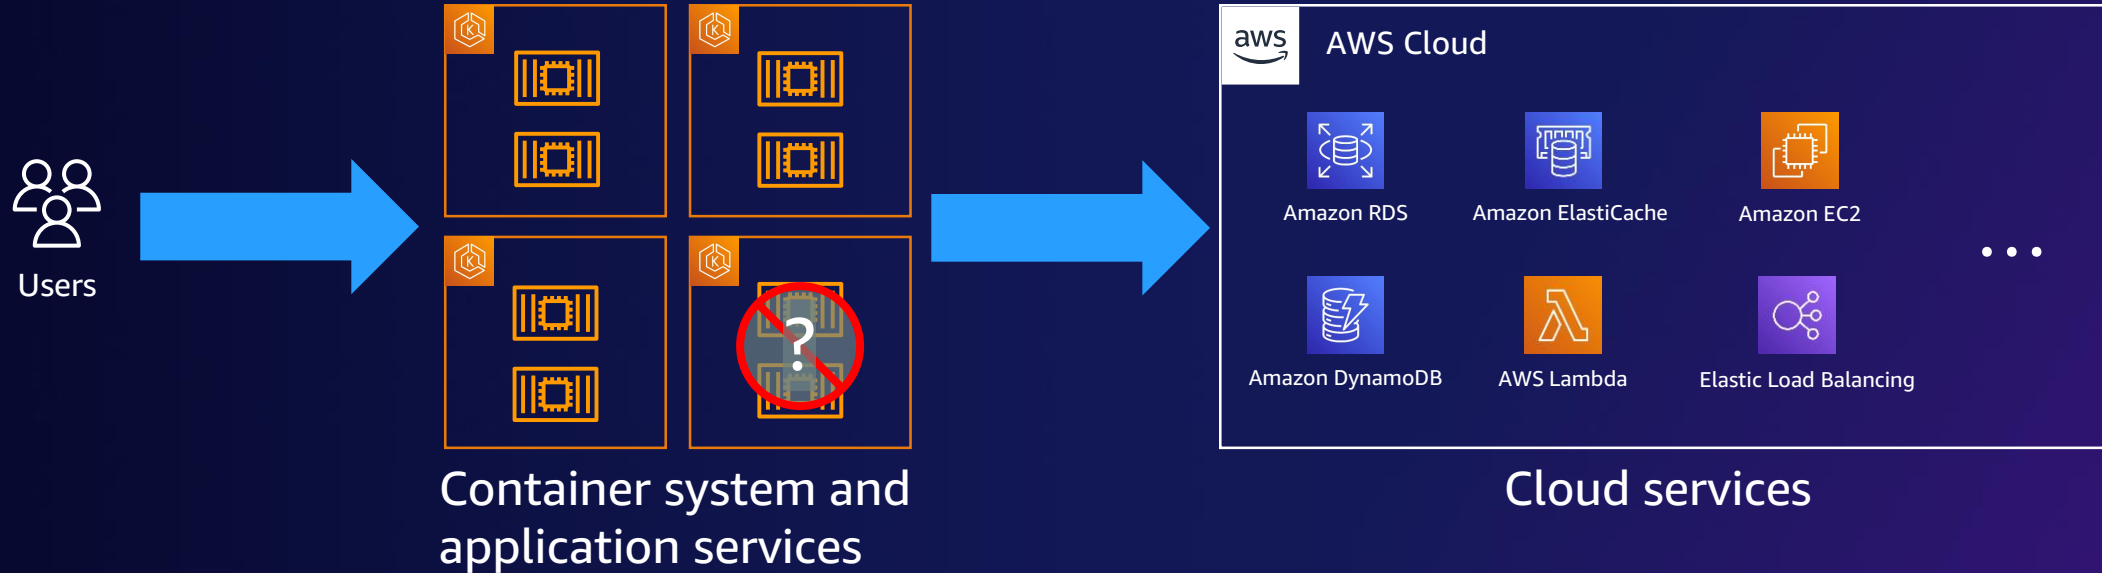
Modern applications



Distributed across services and microservices

Visibility can be low, especially for application microservice interaction and interaction with AWS services

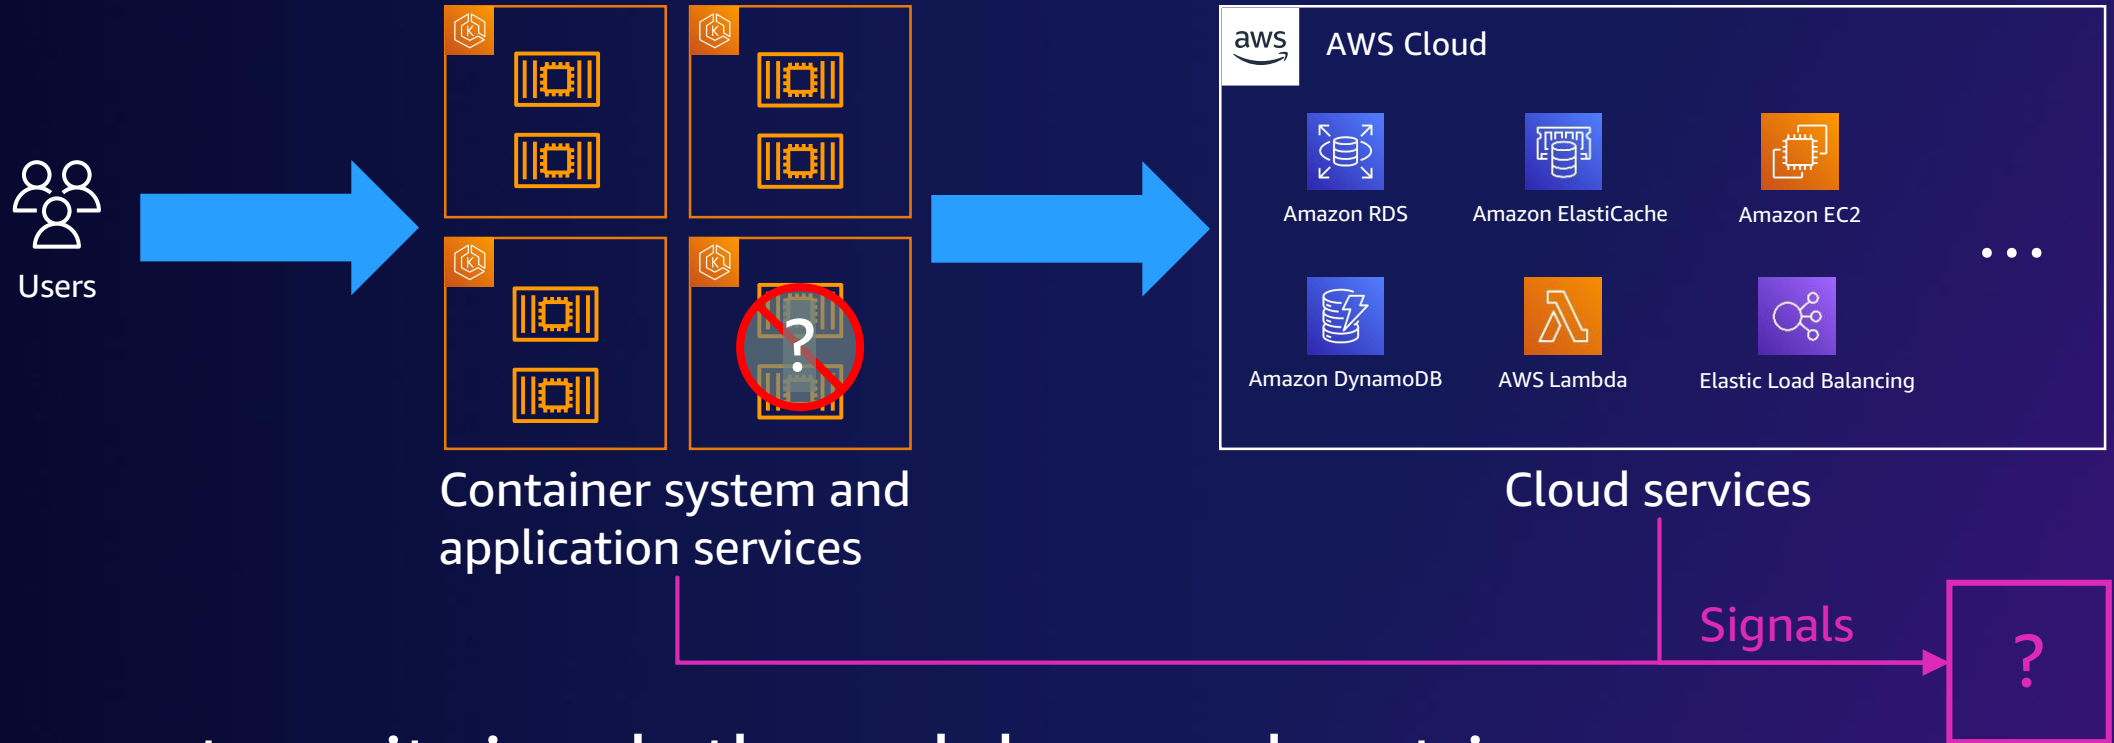
When things go wrong



Was it: a code bug? A service-API failure? A cloud service failure?

Decoupled code and services are hard to diagnose!

Gathering and analyzing signals



Components emit signals through logs and metrics

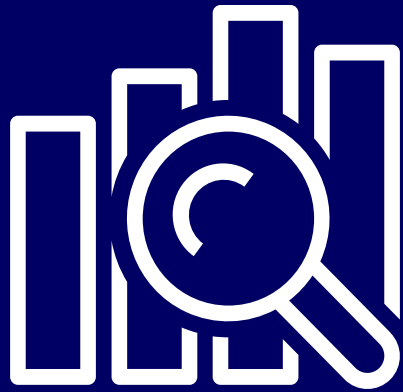
Remediate failures with analysis of interactions and code



OpenSearch

Community-driven, open-source **search and analytics suite** derived from Apache 2.0 licensed Elasticsearch 7.10.2 and Kibana 7.10.2

Consists of a search engine, OpenSearch; a visualization user interface, OpenSearch Dashboards; and a series functionality adding tools and plugins



Amazon OpenSearch Service

A managed service that makes it easy to **deploy, operate, and scale OpenSearch** and legacy Elasticsearch clusters in the AWS Cloud

Perform interactive log analytics, real-time application monitoring, website search, and more. Offers the latest versions of OpenSearch, support for 19 versions of Elasticsearch (1.5 to 7.10 versions), and visualization capabilities powered by OpenSearch Dashboards and Kibana (1.5 to 7.10 versions).

OpenSearch

Search

Commands

search

/search-plugins/ppl/commands/#s...

Search templates

/opensearch/search-template/

Asynchronous search

Use the **search** command to retrieve a document from an index

... You can only use the **search** command as the first command in the PPL query

search source=<index> [boolean-expression]

... your full-text queries into a **search** template to accept user input and dynamically insert it into

... use OpenSearch as a backend **search** engine for your application or website, you can take in user queries from a **search** bar or a form field and pass them as parameters into a **search** template

... can simplify your code with **search** templates

Searching large volumes of data can take a long

The screenshot shows the OpenSearch Dashboards interface. At the top, there's a 'Discover' tab with a search bar and a filter button. Below this is a histogram showing the count of hits over time, with a peak around 1,600 hits. The x-axis is labeled 'timestamp per 3 hours' and the y-axis is 'Count'. Below the histogram, there's a list of log entries with details like agent, host, index, ip, machine, request, response, and timestamp.

The screenshot shows the OpenSearch Dashboards interface. At the top, there's a bar chart showing anomalies by index and detector. The x-axis is 'Now' and the y-axis is 'Anomaly grade'. Below this is a pie chart showing the distribution of anomalies by index and detector. The legend indicates that the inner circle shows the anomaly distribution by your indices, and the outer circle shows the anomaly distribution by your detectors. To the right, there's a table titled 'Detectors and features' with columns for 'Detector' and 'Features'.

Detector	Features
D1	F1 F2 F3
D2	error
Demo-2	aaa bbb
Demo-FnF-prep	Total_order Avg_price
Demo-all-hands	Total_order Avg_price
Demo-one-feature	Total-Order
Demo-verizon	F1 F2 F3
Demo1	F1
Demo2	F1 F2
DemoDetector	F1 F2

Text search

Natural language

Boolean queries

Relevance

Streaming data

High-volume ingest

Near real time

Distributed storage

Analysis

Time-based visualizations

Nestable statistics

Time series tools

aws

© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Foundation for observability: Data drives decisions



Logs

Amazon OpenSearch Service



Metrics

Amazon Managed Service for Prometheus
Amazon Managed Grafana



Traces

AWS Distro for OpenTelemetry
Fluent Bit



AWS monitoring and observability services help you maintain SLAs
by **detecting, investigating, and remediating problems**
to achieve

Availability

Reliability

Performance

Analyzing logs and metrics

```
199.72.81.55 -- [01/Jul/1995:00:00:01 -0400] "GET /history/apollo/ HTTP/1.0" 200 6245
uncomp6.uncomp.net -- [01/Jul/1995:00:00:06 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
199.120.110.21 -- [01/Jul/1995:00:00:09 -0400] "GET /shuttle/missions/sts-73/mission-sts-73.html HTTP/1.0" 200 4085
burger.letters.com -- [01/Jul/1995:00:00:11 -0400] "GET /shuttle/countdown/liftoff.html HTTP/1.0" 304 0
199.120.110.21 -- [01/Jul/1995:00:00:11 -0400] "GET /shuttle/missions/sts-73/sts-73-patch-small.gif HTTP/1.0" 200 41
burger.letters.com -- [01/Jul/1995:00:00:12 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 304 0
burger.letters.com -- [01/Jul/1995:00:00:12 -0400] "GET /shuttle/countdown/video/livevideo.gif HTTP/1.0" 200 0
205.212.115.106 -- [01/Jul/1995:00:00:12 -0400] "GET /shuttle/countdown/countdown.html HTTP/1.0" 200 3985
d104.aa.net -- [01/Jul/1995:00:00:13 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
129.94.144.152 -- [01/Jul/1995:00:00:13 -0400] "GET / HTTP/1.0" 200 7074
uncomp6.uncomp.net -- [01/Jul/1995:00:00:14 -0400] "GET /shuttle/countdown/count.gif HTTP/1.0" 200 40310
uncomp6.uncomp.net -- [01/Jul/1995:00:00:14 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 200 786
uncomp6.uncomp.net -- [01/Jul/1995:00:00:14 -0400] "GET /images/KSC-logosmall.gif HTTP/1.0" 200 1204
d104.aa.net -- [01/Jul/1995:00:00:15 -0400] "GET /shuttle/countdown/count.gif HTTP/1.0" 200 40310
d104.aa.net -- [01/Jul/1995:00:00:15 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 200 786
d104.aa.net -- [01/Jul/1995:00:00:15 -0400] "GET /images/KSC-logosmall.gif HTTP/1.0" 200 1204
129.94.144.152 -- [01/Jul/1995:00:00:17 -0400] "GET /images/ksclogo-medium.gif HTTP/1.0" 304 0
199.120.110.21 -- [01/Jul/1995:00:00:17 -0400] "GET /images/launch-logo.gif HTTP/1.0" 200 1713
ppptky391.asahi-net.or.jp -- [01/Jul/1995:00:00:18 -0400] "GET /facts/about_ksc.html HTTP/1.0" 200 3977
net-1-141.eden.com -- [01/Jul/1995:00:00:19 -0400] "GET /shuttle/missions/sts-71/images/KSC-95EC-0916.jpg HTTP/1.0"
ppptky391.asahi-net.or.jp -- [01/Jul/1995:00:00:19 -0400] "GET /images/launchpalms-small.gif HTTP/1.0" 200 11473
205.189.154.54 -- [01/Jul/1995:00:00:24 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
waters-gw.starway.net.au -- [01/Jul/1995:00:00:25 -0400] "GET /shuttle/missions/51-l/mission-51-l.html HTTP/1.0" 200
ppp-mia-30.shadow.net -- [01/Jul/1995:00:00:27 -0400] "GET / HTTP/1.0" 200 7074
205.189.154.54 -- [01/Jul/1995:00:00:29 -0400] "GET /shuttle/countdown/count.gif HTTP/1.0" 200 40310
alyssa.prodigy.com -- [01/Jul/1995:00:00:33 -0400] "GET /shuttle/missions/sts-71/sts-71-patch-small.gif HTTP/1.0" 20
ppp-mia-30.shadow.net -- [01/Jul/1995:00:00:35 -0400] "GET /images/ksclogo-medium.gif HTTP/1.0" 200 5866
dial22.lloyd.com -- [01/Jul/1995:00:00:37 -0400] "GET /shuttle/missions/sts-71/images/KSC-95EC-0613.jpg HTTP/1.0" 20
smth-pc.moorecap.com -- [01/Jul/1995:00:00:38 -0400] "GET /history/apollo/apollo-13/images/70HC314.GIF HTTP/1.0" 20
205.189.154.54 -- [01/Jul/1995:00:00:40 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 200 786
ix-orl2-01.ix.netcom.com -- [01/Jul/1995:00:00:41 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
```

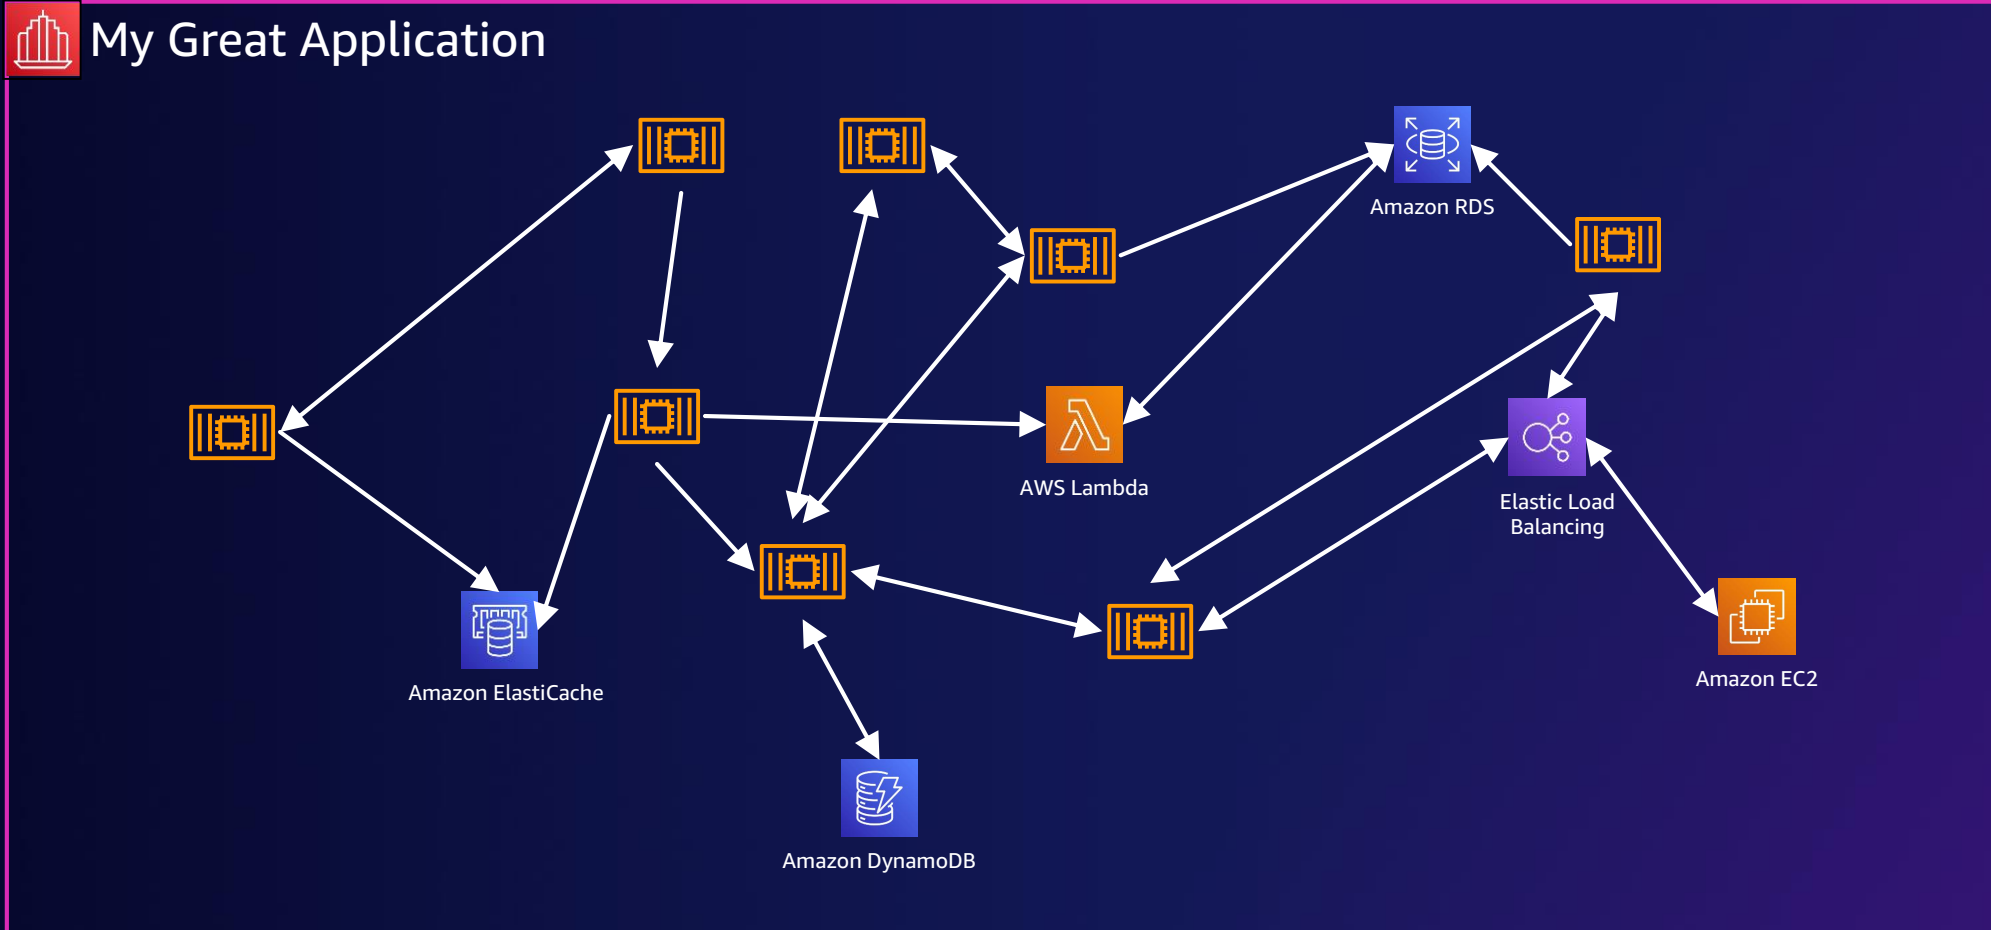


System-level view, from metrics and information in log lines

Error location through messages in logs

Missing: interaction between subsystems

What about service interaction?



What a mess!

OpenTelemetry



- OpenTelemetry is a **community-driven, open-source** project designed for the standardization, creation, and management of **telemetry data** such as **traces, metrics, and logs**
- Supports many popular open-source wire formats including **Jaeger, Zipkin, and Prometheus**
- **Currently supports** traces (GA), metrics (preview), and logs (experimental)
- Traces are emitted as a log line, either intra- or inter-service

AWS Distro for OpenTelemetry



OpenTelemetry

A Cloud Native Computing Foundation project

Open source observability agents and libraries

Supports all 3 data signals in 11 languages

AWS Distro for OpenTelemetry

Secure, production-ready, open-source distribution supported by AWS

Code contributions are upstream in OpenTelemetry

Certified by AWS for security and predictability



Data Prepper



- Data Prepper is a **community-driven, open-source** data collector for processing observability data
- Provides features to **filter, enrich, transform, normalize, and aggregate data** for downstream analytics and visualization
- Currently supports **processing of distributed trace data and log ingestion** with plans to support metric data in the future
- Integrations with **Jeager, Zipkin, OpenTelemetry, and Fluent Bit**

Traces and spans

A span is a unit of work

Spans have

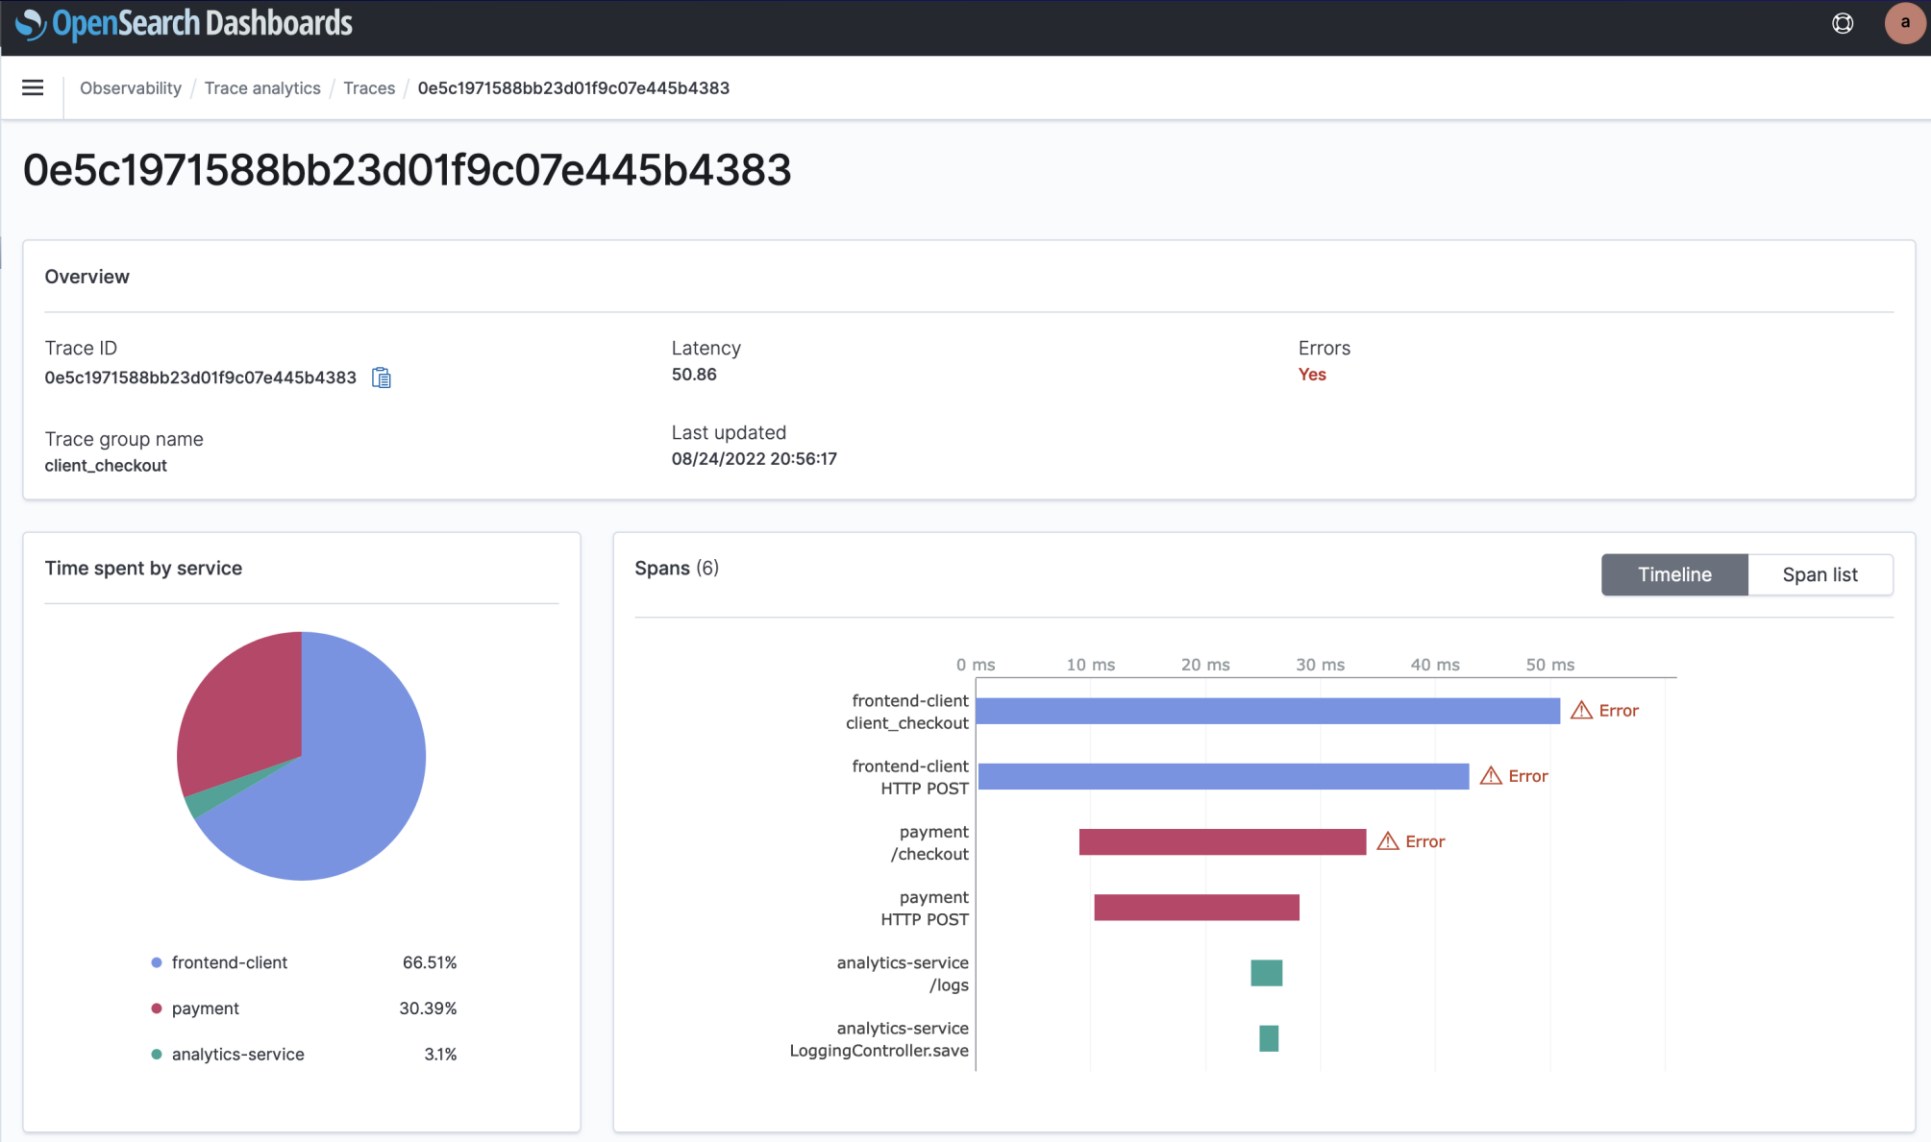
- Duration
- Methods
- Status codes

A trace is a particular user request

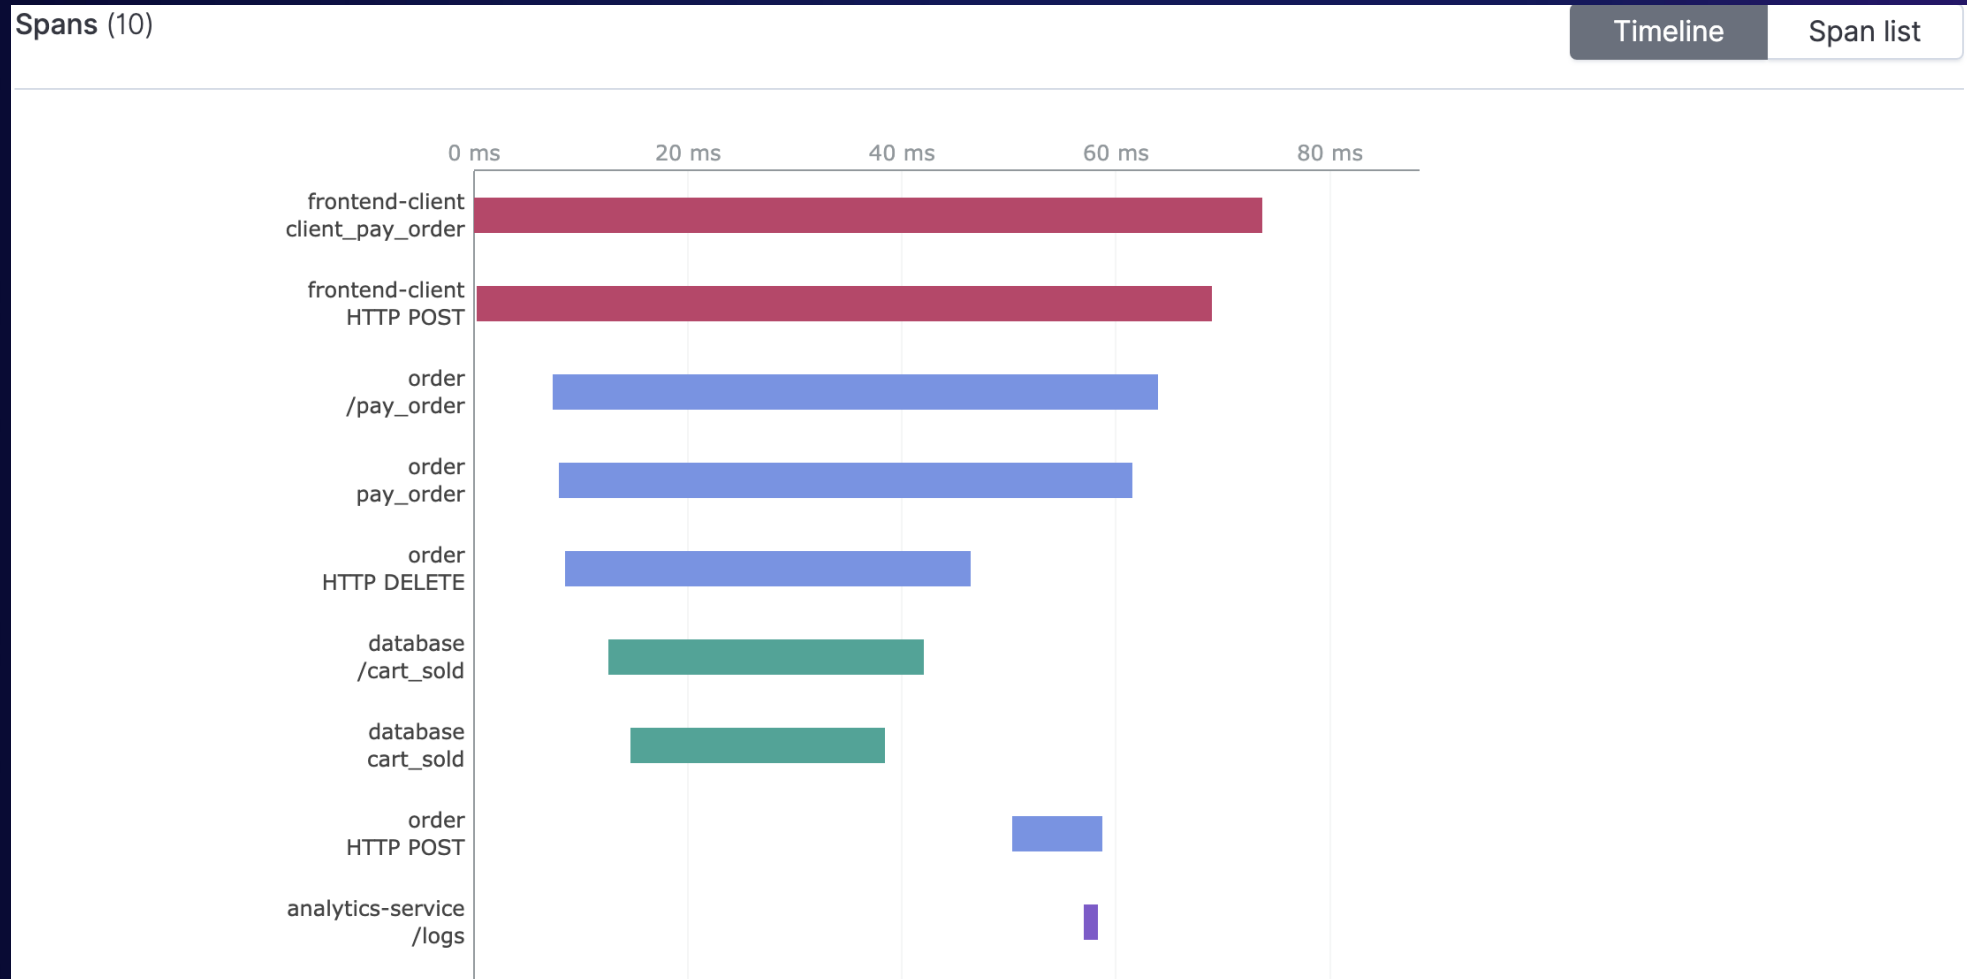
Traces have spans

```
{
  "_index" : "otel-v1-apm-span-000001",
  "_type" : "_doc",
  "_id" : "520f74be0ad94c8c",
  "_score" : 1.0,
  "_source" : {
    "traceId" : "348c4305a3418a1ea5631de518bb5cf7",
    "spanId" : "520f74be0ad94c8c",
    "traceState" : "",
    "parentSpanId" : "89e79a790afc003a",
    "name" : "HTTP PUT",
    "kind" : "SPAN_KIND_CLIENT",
    "startTime" : "2022-08-23T01:33:47.482462751Z",
    "endTime" : "2022-08-23T01:33:47.532808718Z",
    "durationInNanos" : 50345967,
    "serviceName" : "order",
    "events" : [ ],
    "links" : [ ],
    "droppedAttributesCount" : 0,
    "droppedEventsCount" : 0,
    "droppedLinksCount" : 0,
    "traceGroup" : "client_cancel_order",
    "traceGroupFields.endTime" : "2022-08-23T01:33:47.578544903Z",
    "traceGroupFields.statusCode" : 0,
    "traceGroupFields.durationInNanos" : 107274133,
    "span.attributes.http.method" : "PUT",
    "span.attributes.http.url" : "http://database-service.database-service.svc.cluster.local:80/cart_empty",
    "resource.attributes.telemetry.sdk.name" : "opentelemetry",
    "instrumentationLibrary.version" : "0.28b1",
    "resource.attributes.telemetry.sdk.language" : "python",
    "resource.attributes.telemetry.sdk.version" : "1.9.1",
    "resource.attributes.service.instance.id" : "139947915721696",
    "resource.attributes.service.name" : "order",
    "resource.attributes.host.hostname" : "order-service-6bc467f446-6kmt",
    "span.attributes.http.status_code" : 200,
    "status.code" : 0,
    "instrumentationLibrary.name" : "opentelemetry.instrumentation.requests"
  }
},
```


A single trace **contains** multiple spans



A single span can have MULTIPLE children



A single span can be of different kinds

- **SERVER**
 - Describes server-side handling of a synchronous remote request
 - Is the child of a remote CLIENT span
- **CLIENT**
 - Describes a request to some remote service
 - Is the parent of a remote SERVER span and does not end until the response is received
- **PRODUCER**
 - Describes initiators of an asynchronous request
 - Will often end before the corresponding child CONSUMER span
- **CONSUMER**
 - Describes a child of an asynchronous PRODUCER request
- **INTERNAL (Default)**
 - Describes internal operation within an process

Amazon OpenSearch Service



The what and why of search



If you have an **e-commerce platform**, you want customers to find the product they are looking for quickly



If you have a **document portal** with documents like scientific research articles, investment analyses, or health records, you want to enable a speedy and relevant search experience for your users



You may want to increase user engagement on your platform by **delivering personalized recommendations**, like a weekly music playlist or food recipes



Beyond these examples, you may have other parts of your tech stack where you want to add an **easy-to-use and snappy search experience**, especially with the option to integrate machine learning capabilities to power a personalized experience

Logs are ubiquitous



Applications and infrastructure

Services/microservices

Web applications

Business applications

APIs



IT and DevOps

Databases

Load balancers

Networking

Servers



IoT and wireless

Automotive

Home devices

Manufacturing

Mobile applications

Turning logs into insights



Applications

Is my infrastructure working?

What is the latency and error rate?

What caused my application issue?



Security

Is there any suspicious authentication activity?

What data was accessed by this IP address?

Are there instances of fraud?



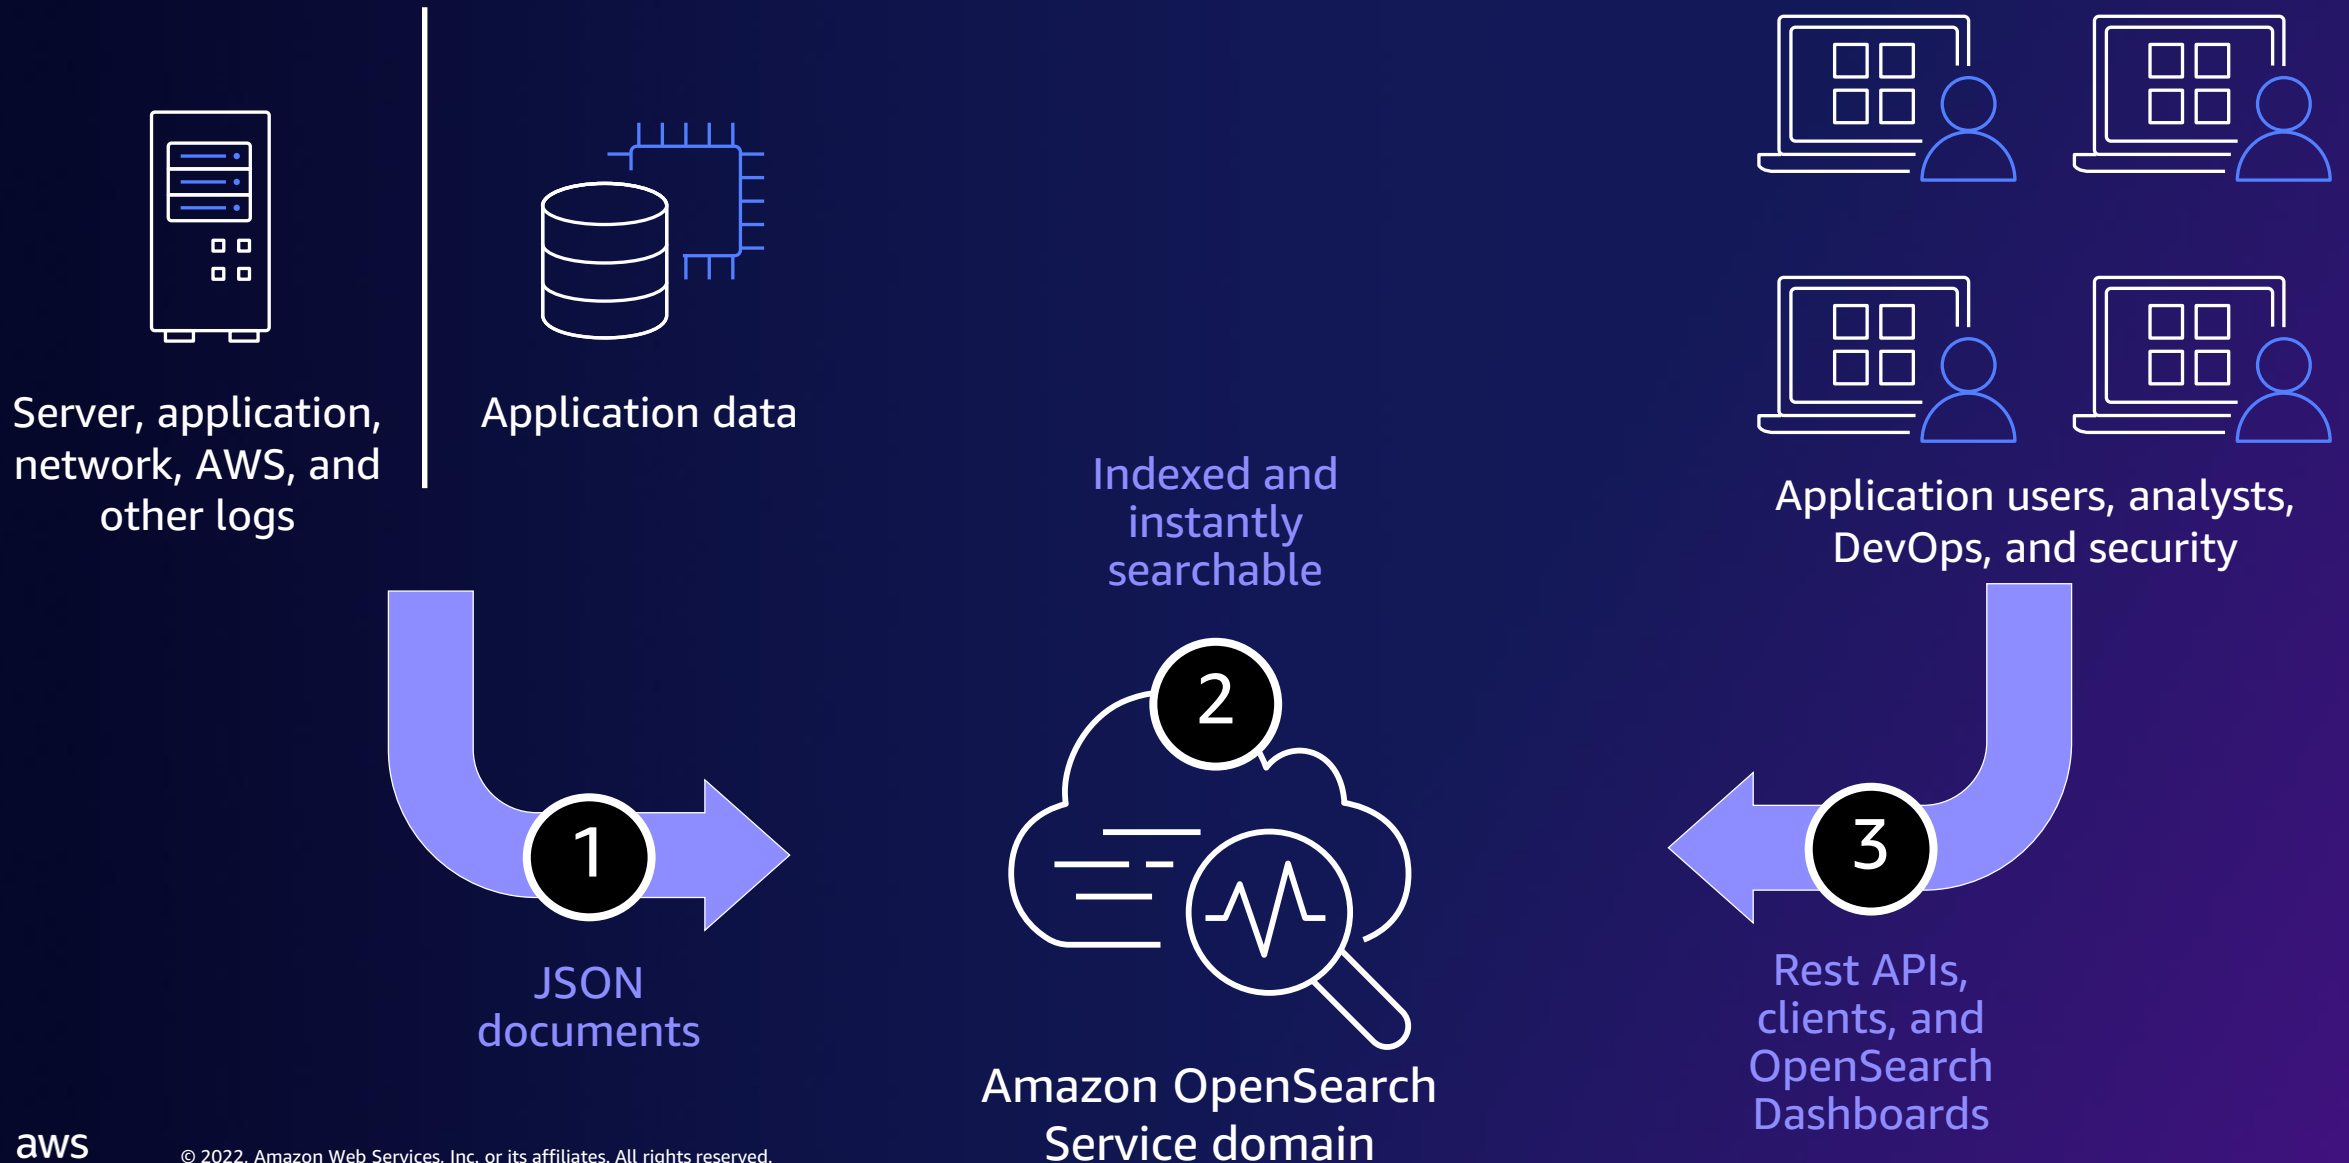
Business insights

What content/products are my users interested in?

Which features are used most or least?

What users are most active and why?

How does it work?



Amazon OpenSearch Service multi-layer security



Integrate with SAML and Amazon Cognito for OpenSearch Dashboards login

IAM to control access to the endpoint

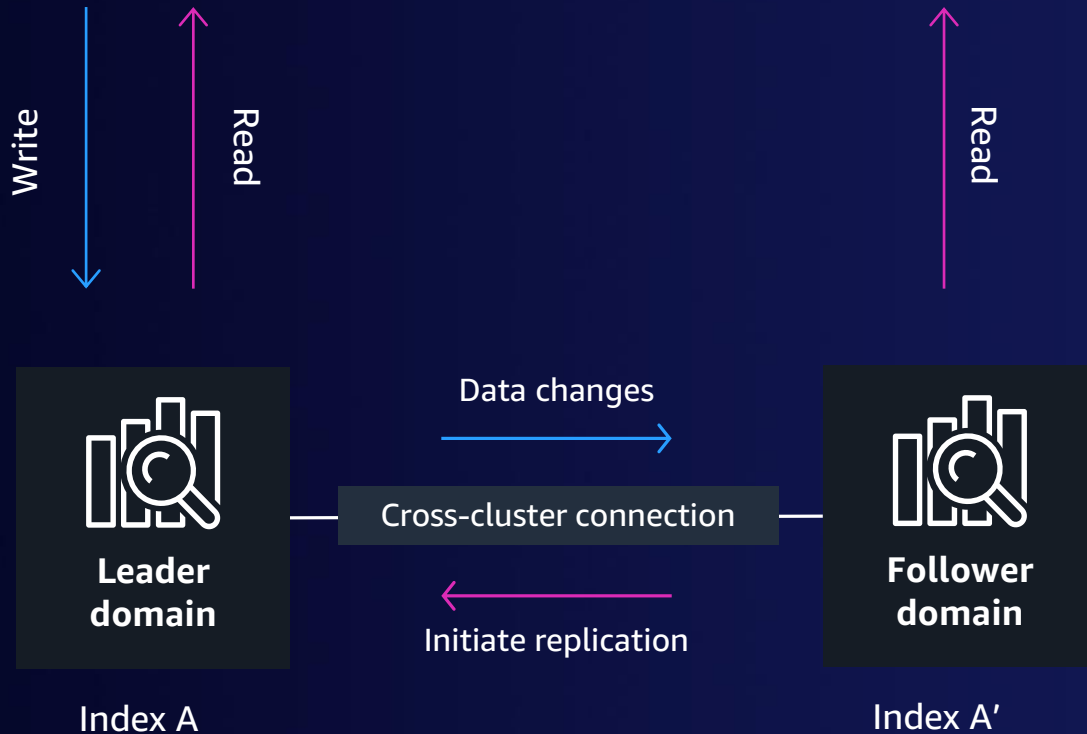
Use a private endpoint to deploy into your VPC and security groups for traffic control

Use OpenSearch fine-grained access control to secure your data and dashboards

Encrypt your data, in flight and at rest

Cross-cluster replication

HIGH AVAILABILITY, DISASTER RECOVERY, AND DATA PROXIMITY



Make your data highly available

- Active-passive replication
- Granular control at index level
- Cross-account and cross-Region support
- Sequential consistency and sub-minute global replication
- One leader/many followers

Machine learning innovations in OpenSearch

ANOMALY DETECTION FOR TIME SERIES

Mitigate issues faster with anomaly detection in streaming data

Performant at scale – machine learning models are distributed and processed across nodes

Easy to use – machine learning expertise is not required to use the service

Based on Random Cut Forest (RCF):
Guha, Mishra, Roy, Schrijvers ICML 2016

Dashboard

[Create detector](#)

All detectors

All detector states

All indices

Live anomalies Live

Live anomaly results across detectors for the last 30 minutes. The results refresh every 1 minute. For each detector, if an anomaly occurrence is detected at the end of the detector interval, you will see a bar representing its anomaly grade.

[View full screen](#)

Last updated time

05/13/2020 06:46 PM

Detector with most recent anomaly occurrence

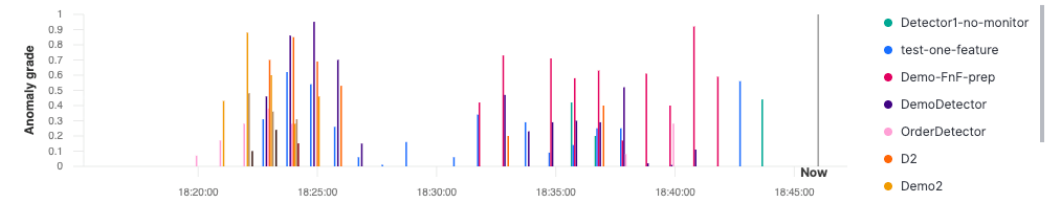
Detector1-no-monitor

Most recent anomaly grade

0.44

You are viewing 10 detectors with the most recent anomaly occurrences.

10 detectors with the most recent anomalies are shown on the chart. Adjust filters if there are specific detectors you would like to monitor.



Anomalies by index and detector

The inner circle shows the anomaly distribution by your indices. The outer circle shows the anomaly distribution by your detectors.

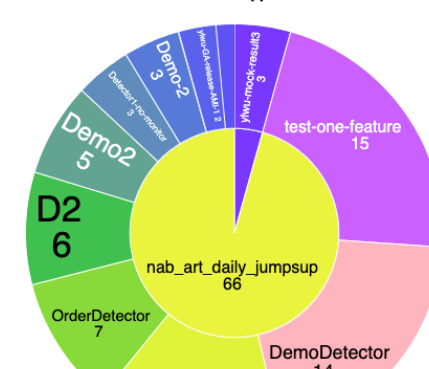
Last 24 hours

Indices with anomalies

2

Detectors with anomalies

11



Detectors and features

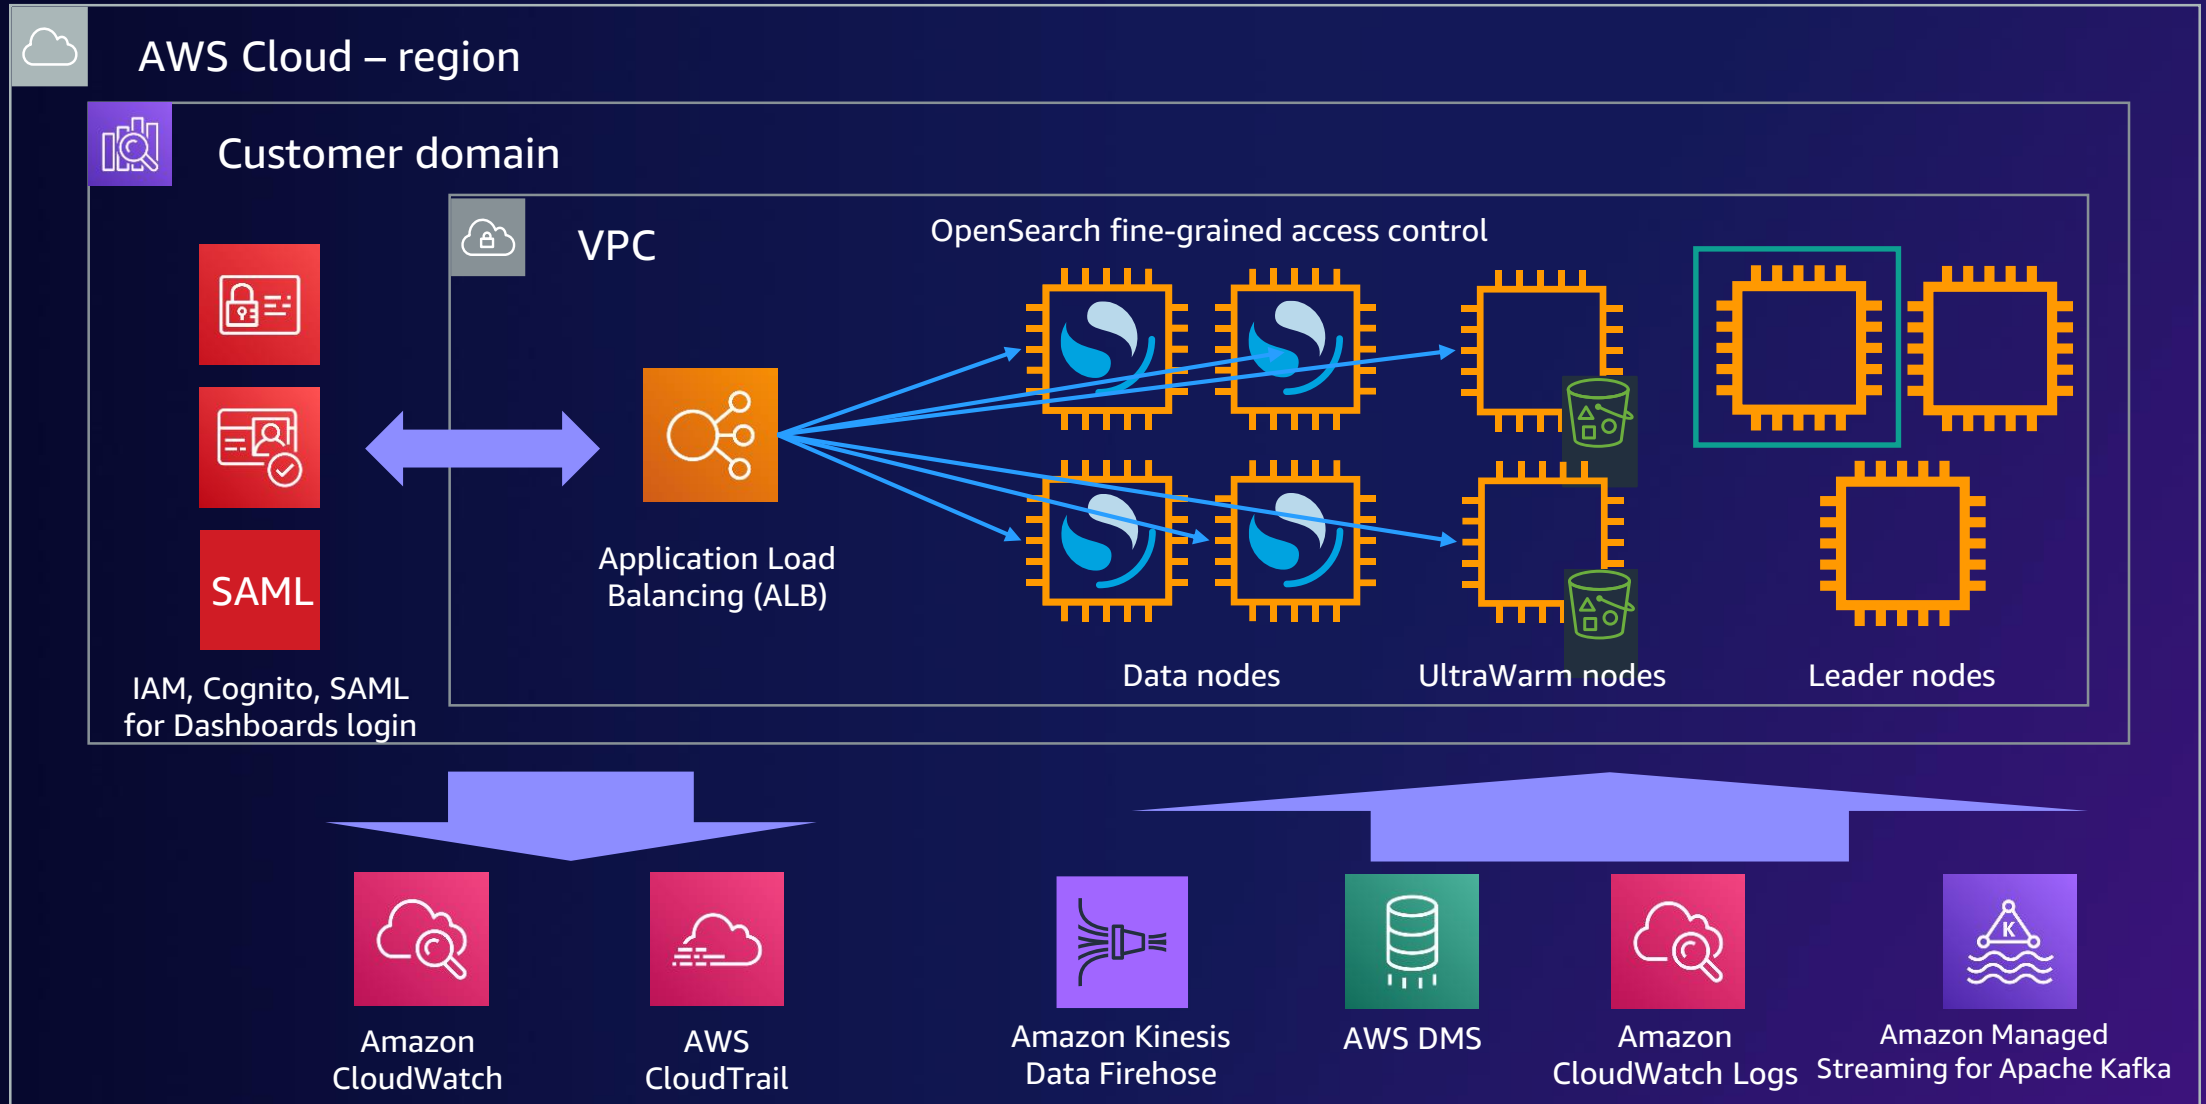
Detector	Features
D1	F1 F2 F3
D2	error
Demo-2	aaa bbb
Demo-FnF-prep	Total_order Avg_price
Demo-all-hands	Total_order Avg_price
Demo-one-feature	Total-Order
Demo-verizon	F1 F2 F3
Demo1	F1
Demo2	f1 f2 f3



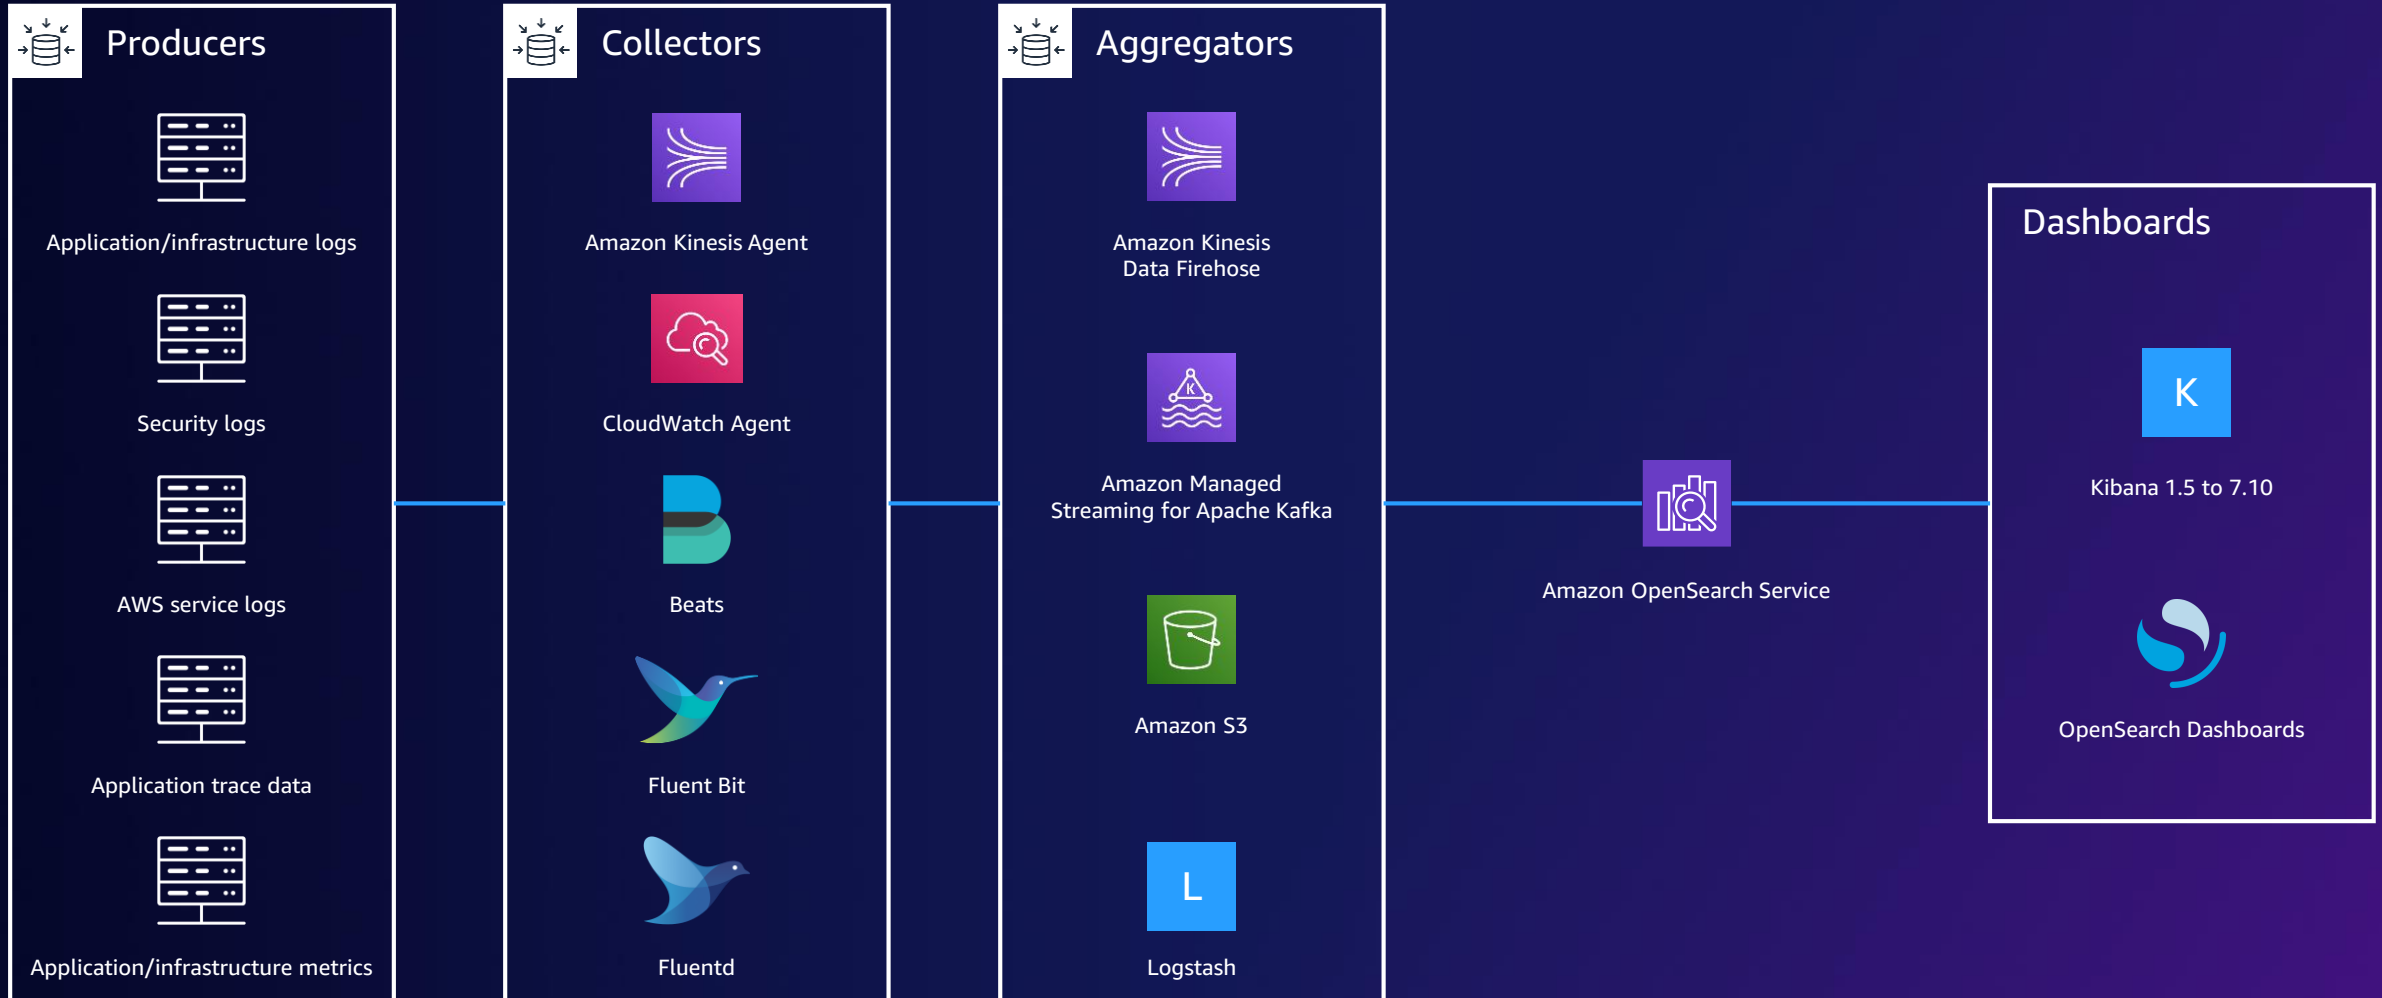
Architectures



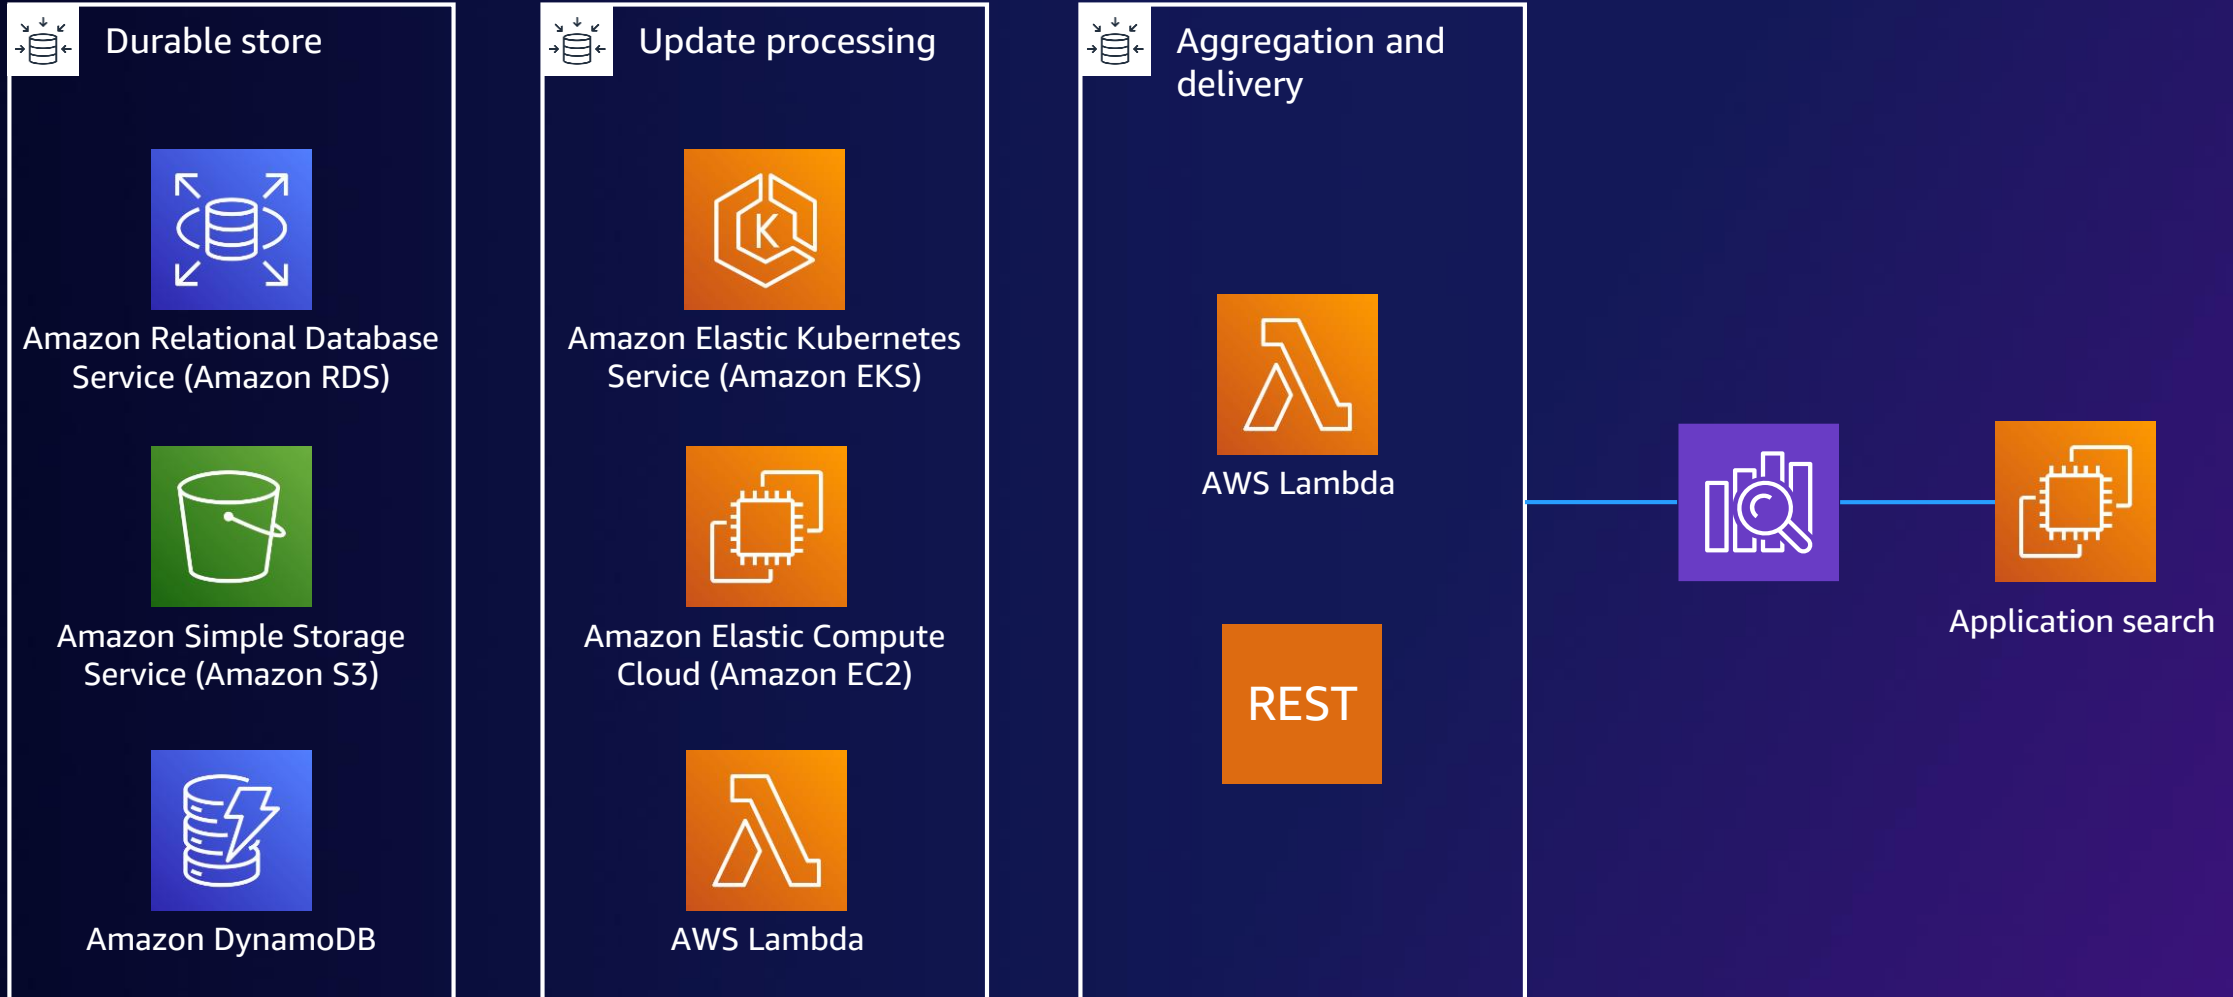
Amazon OpenSearch Service architecture



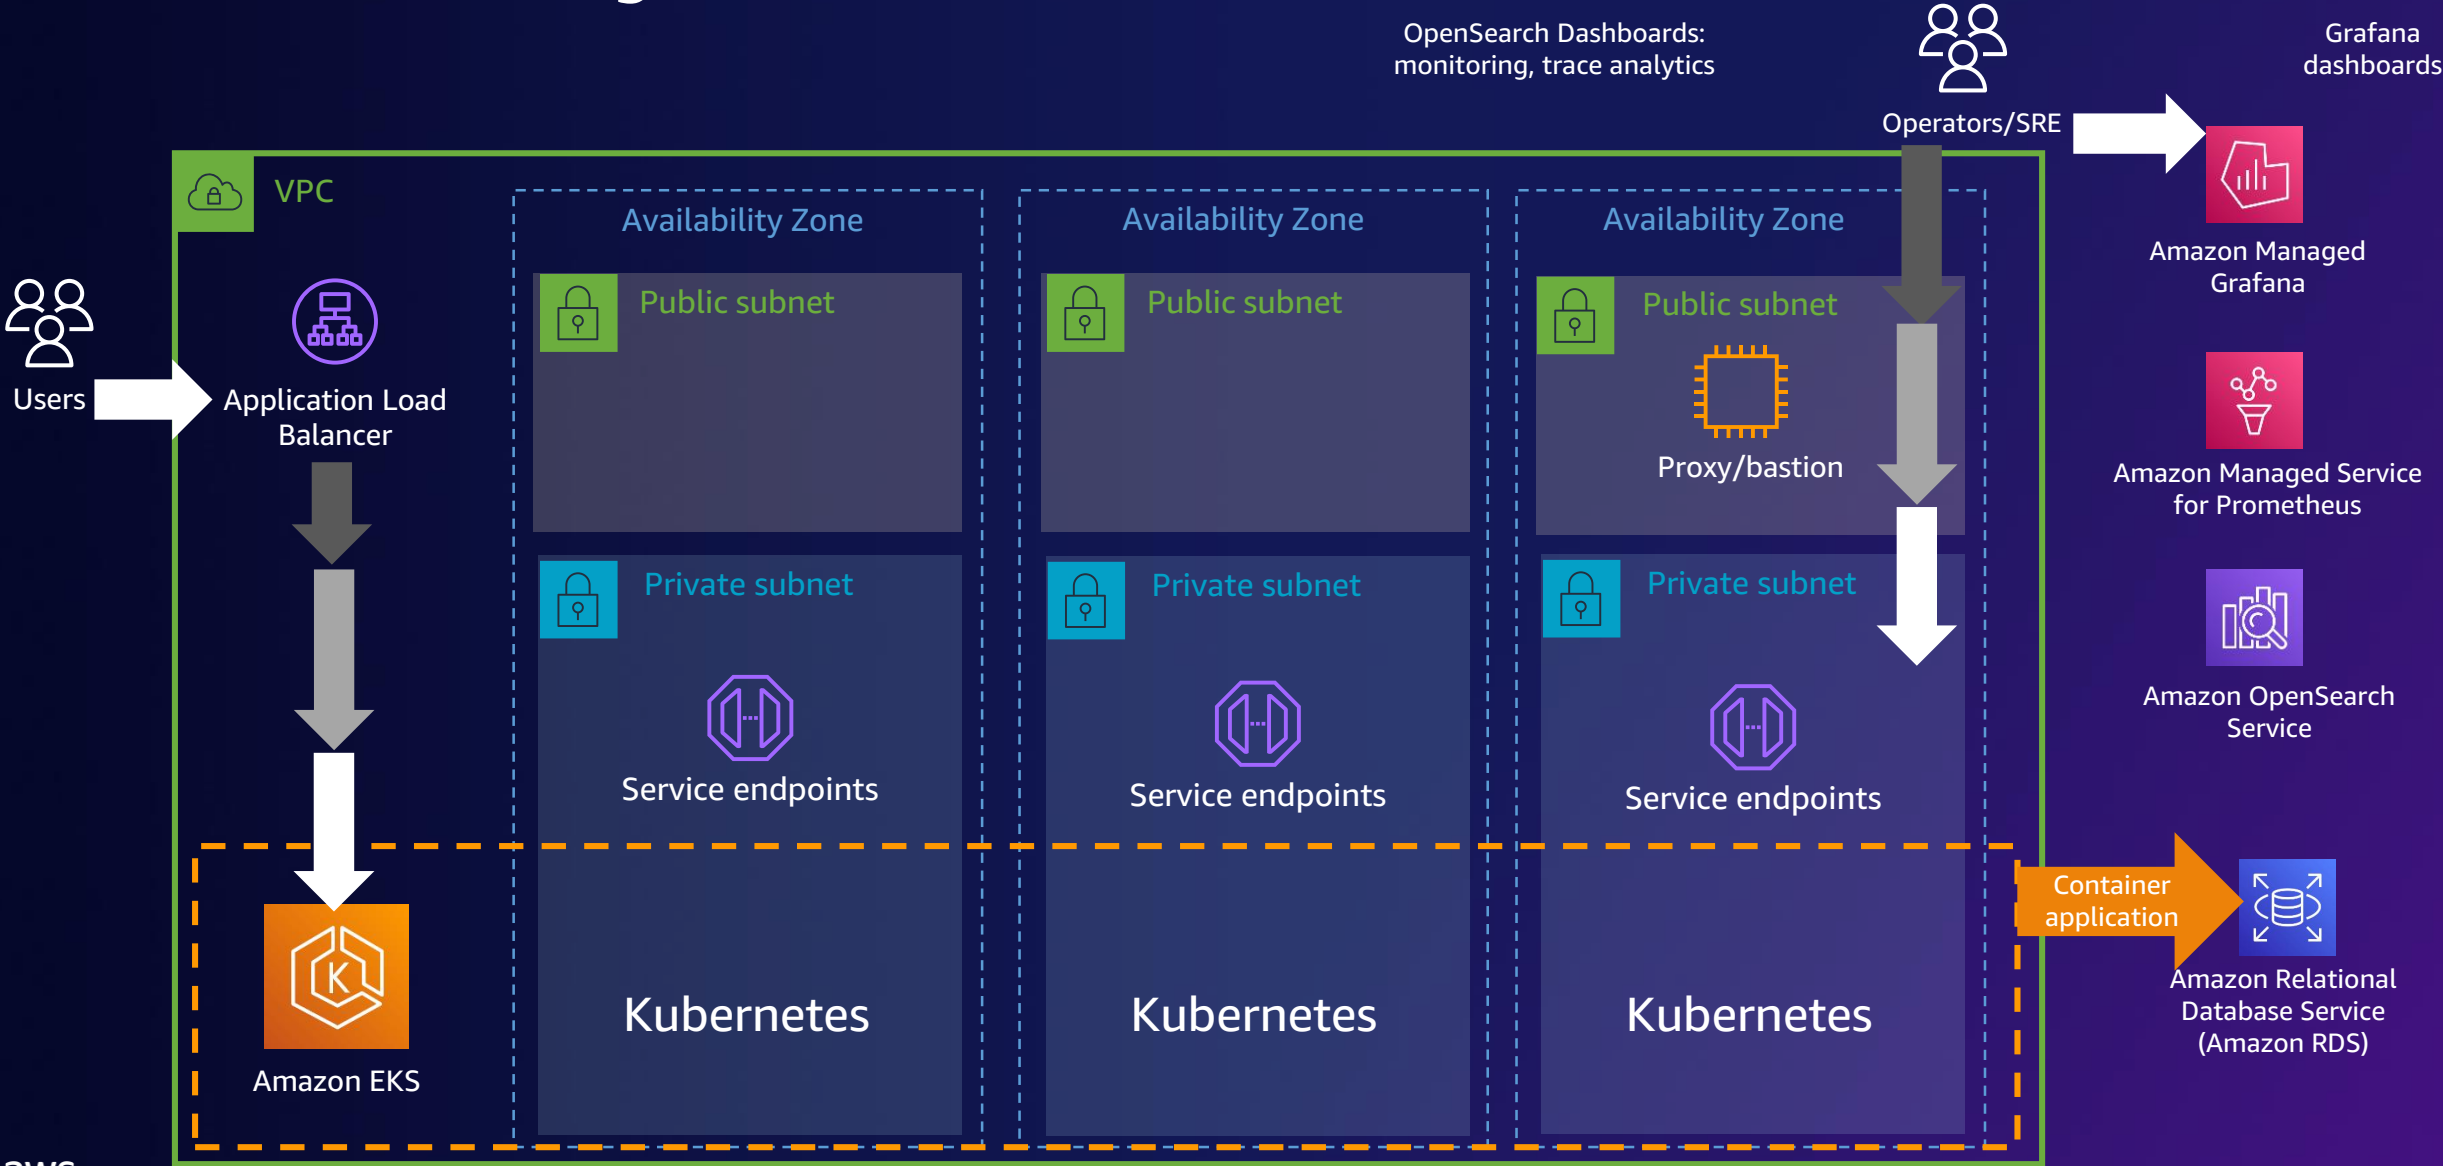
Amazon OpenSearch Service logs ingestion flow



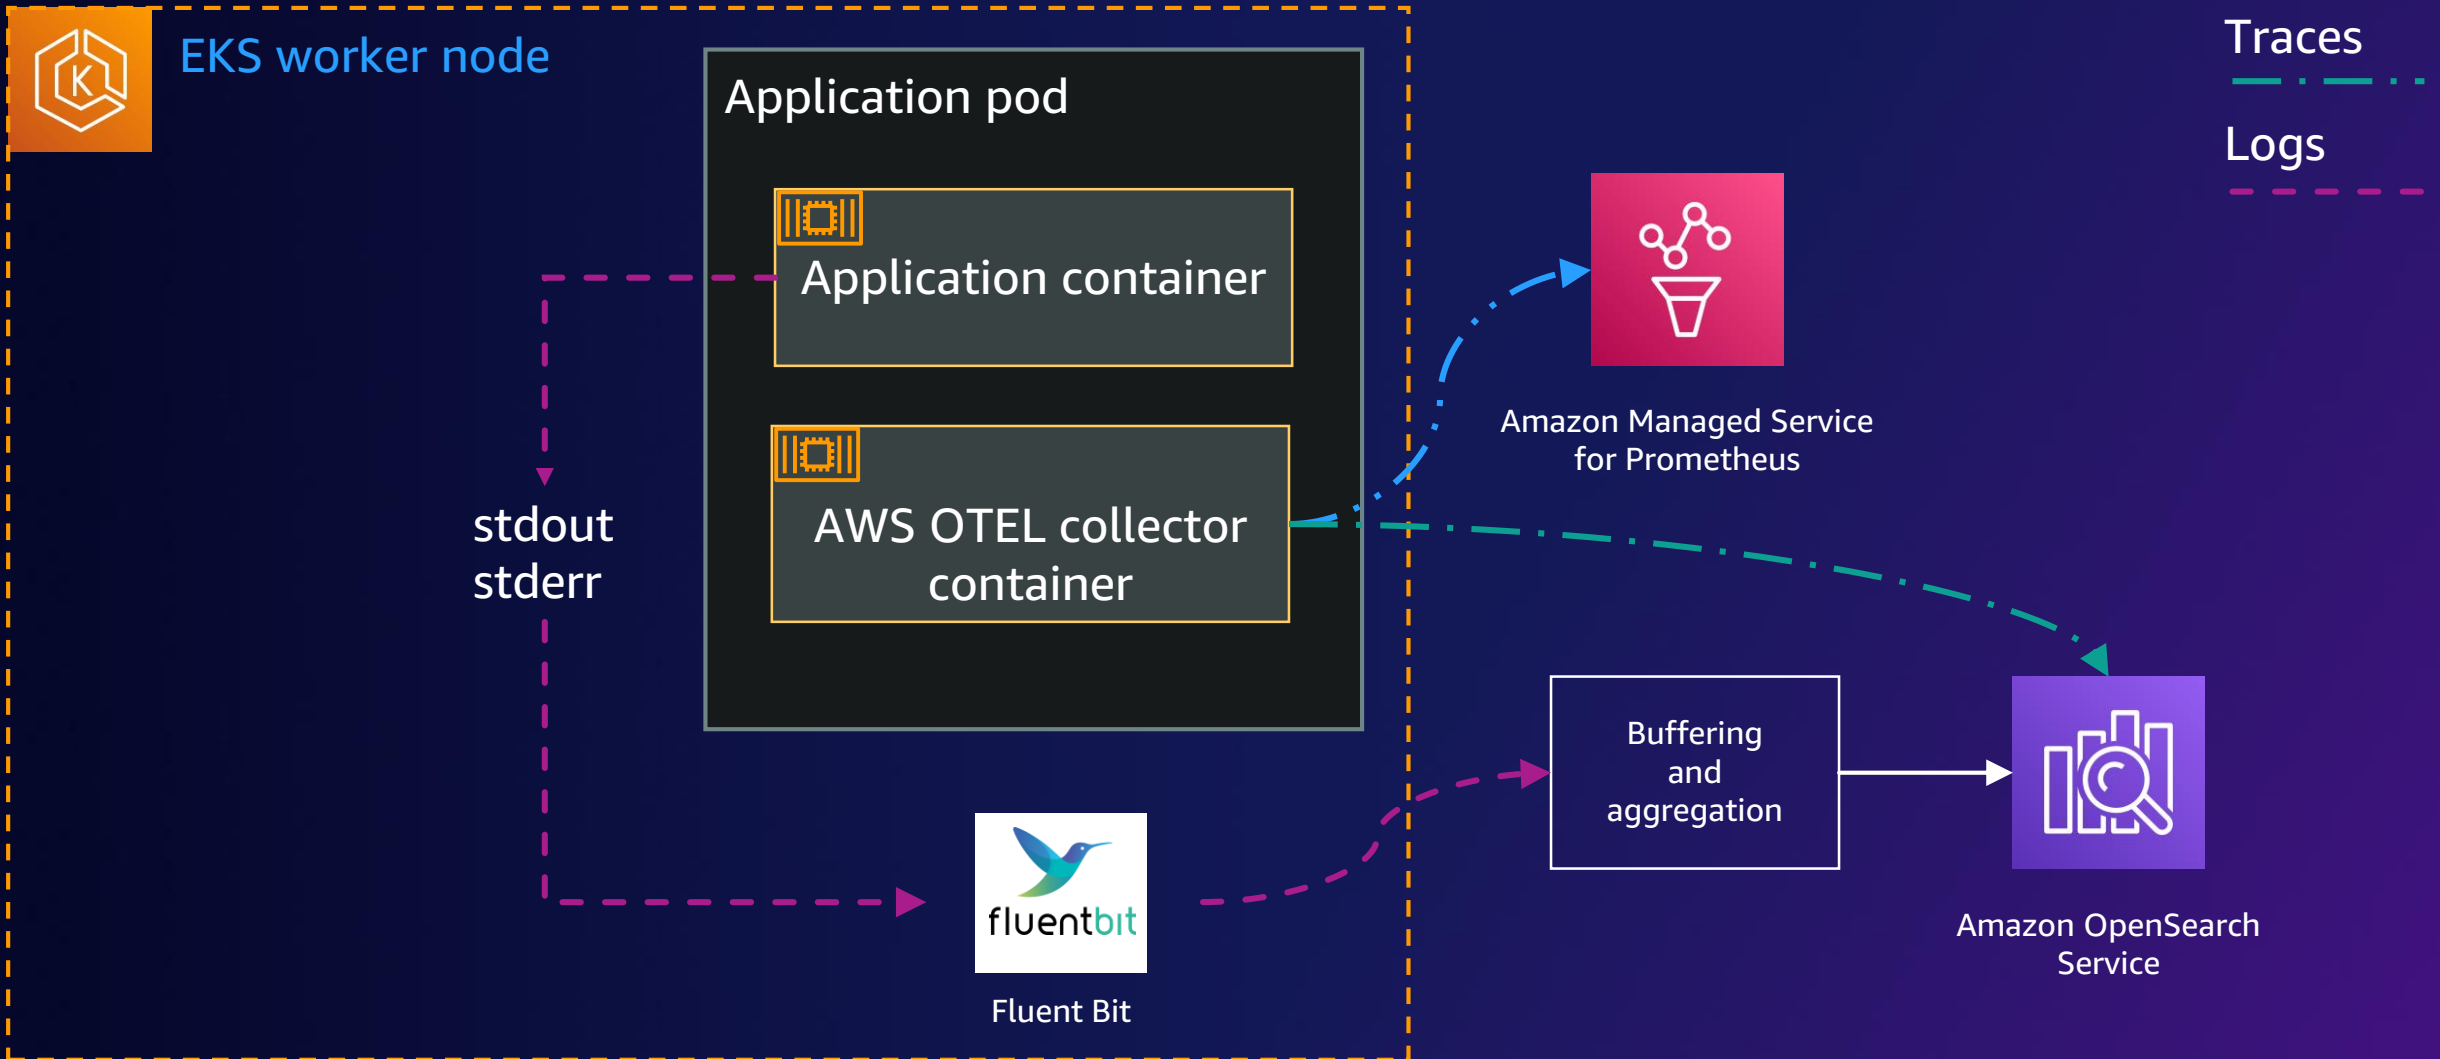
Amazon OpenSearch Service search ingestion flow



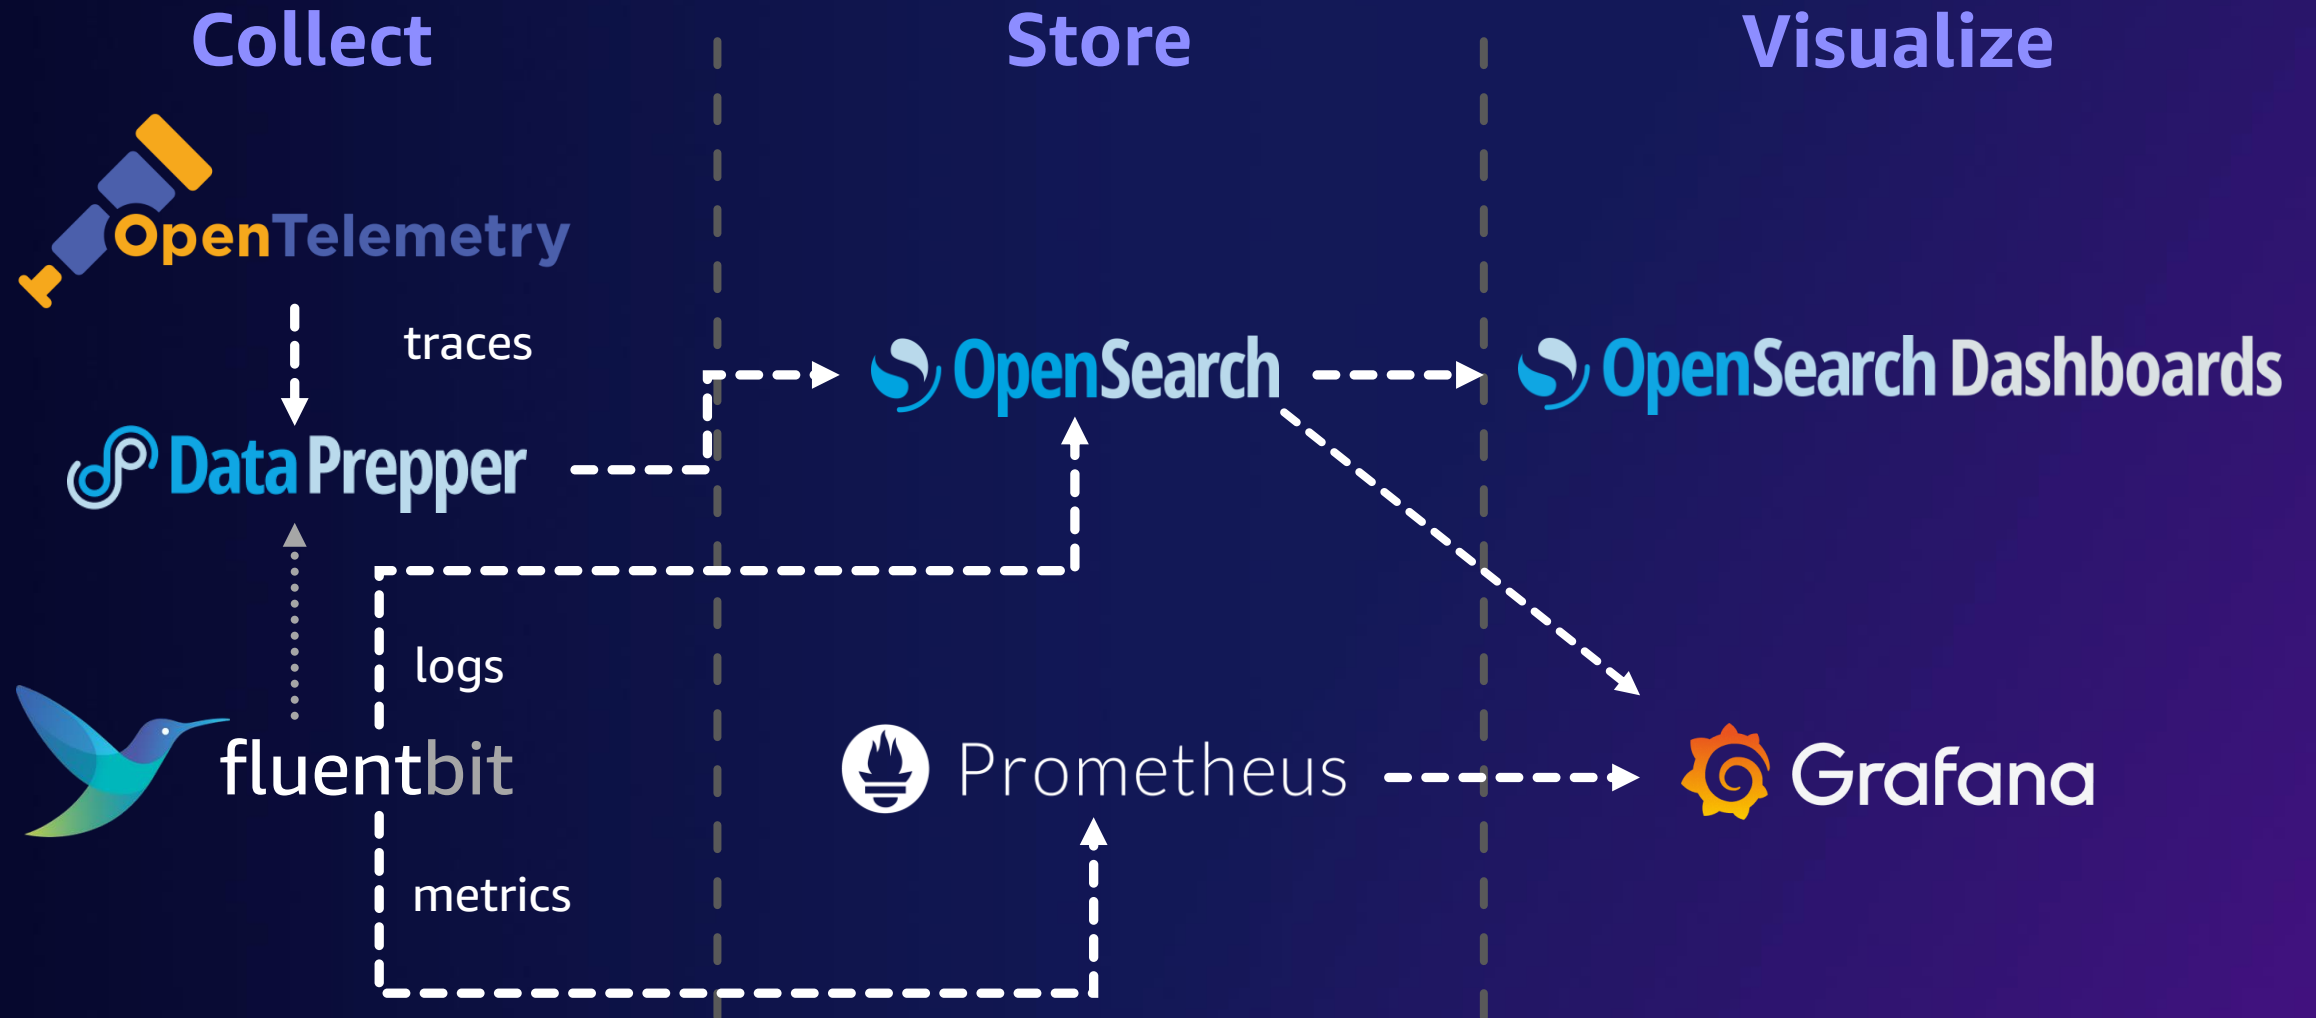
Observability architecture with AWS services



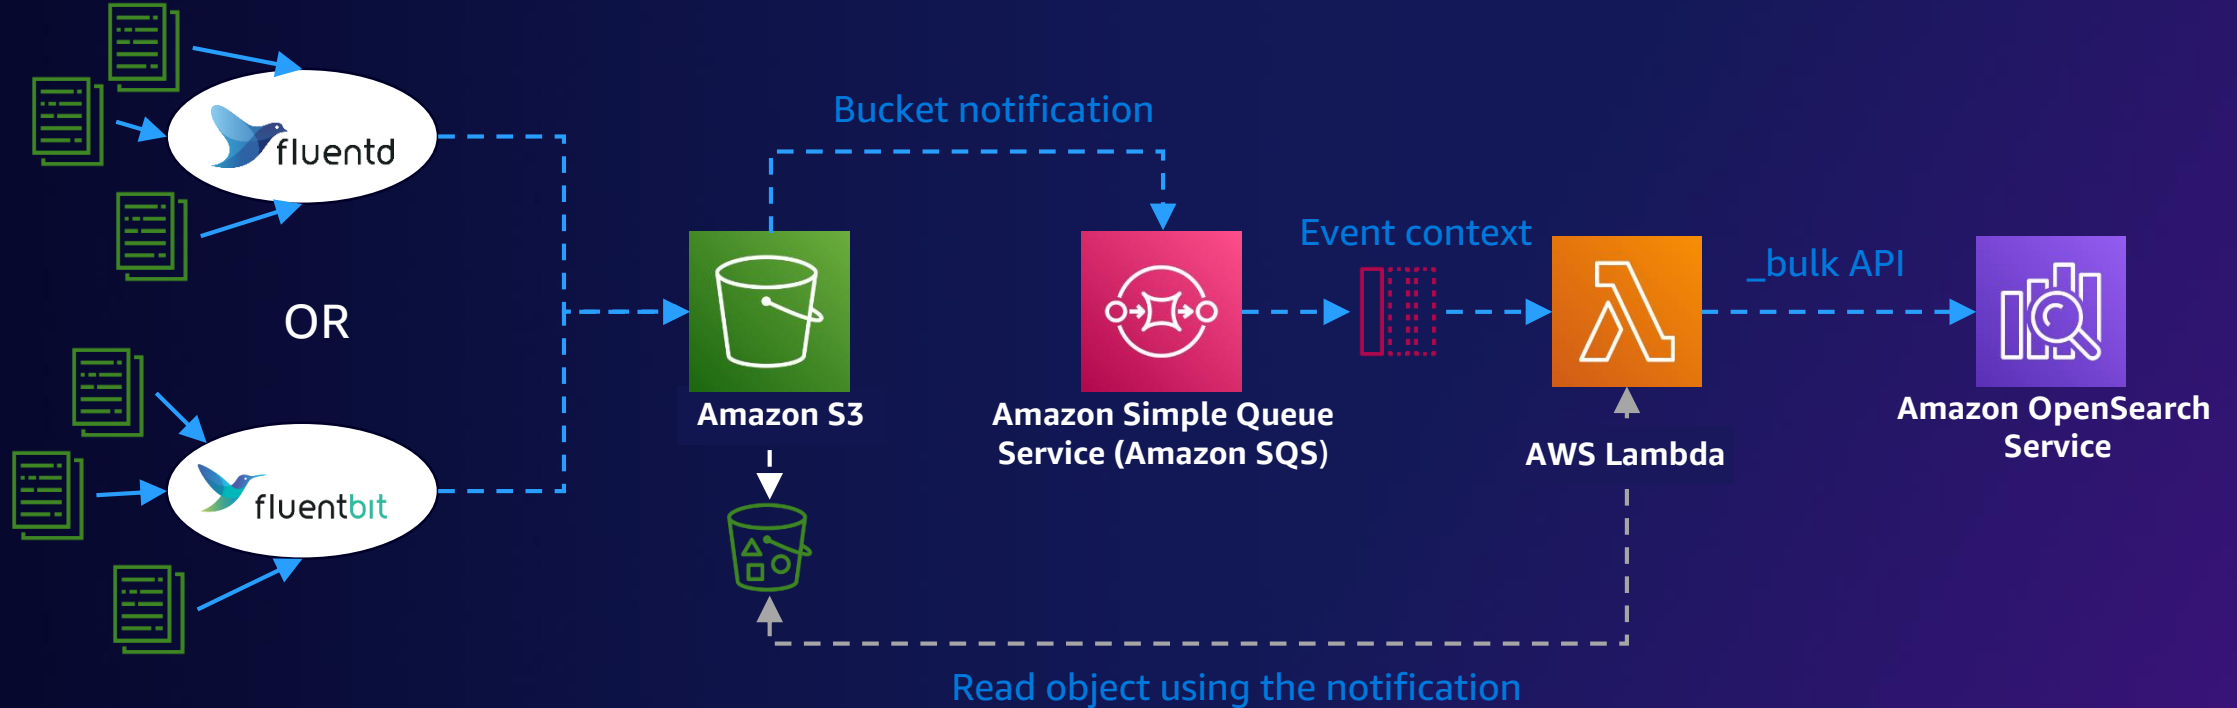
Container node data flow



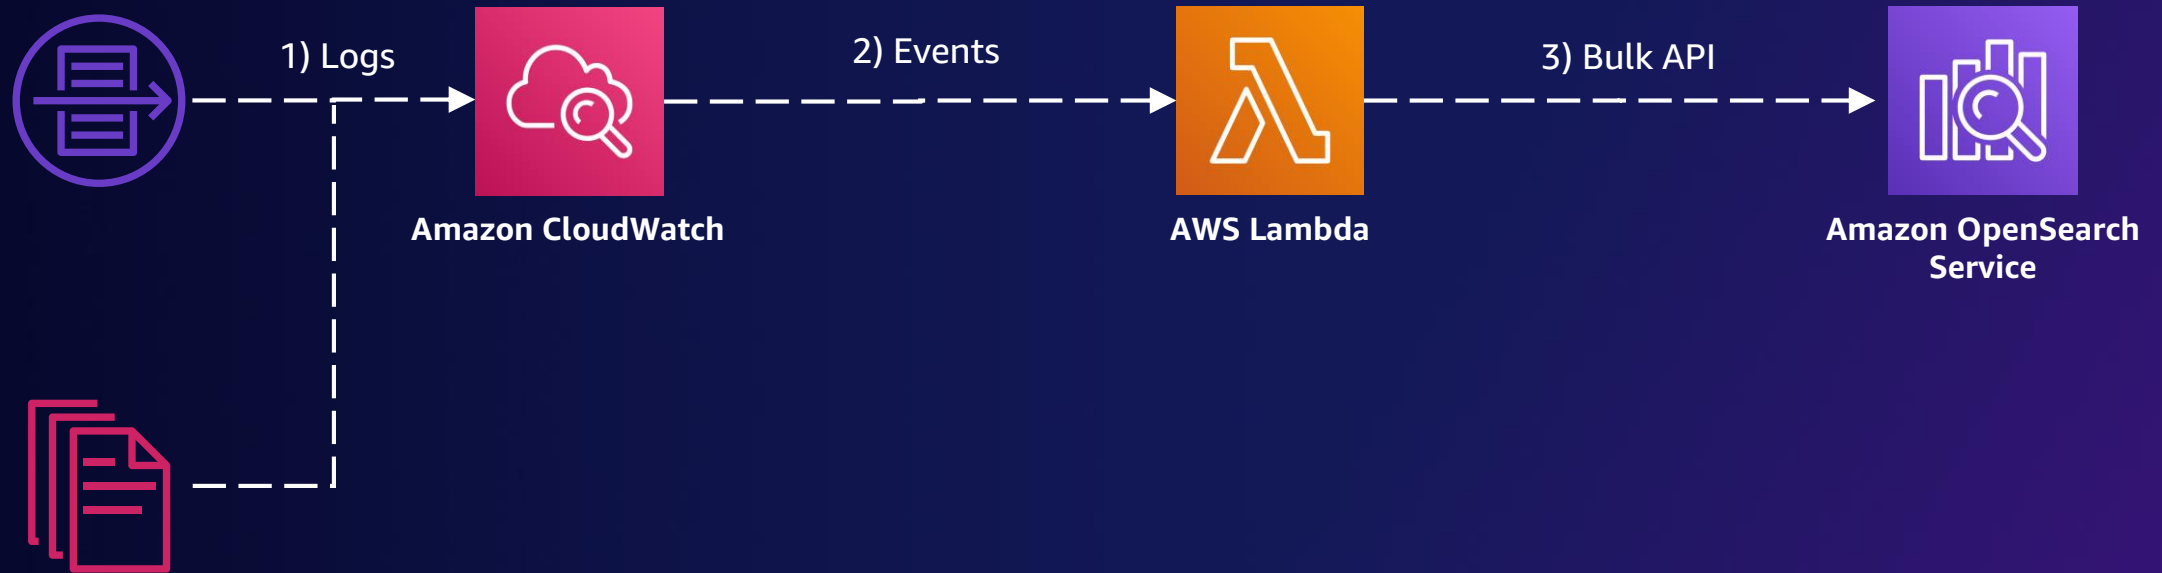
Open source observability architecture



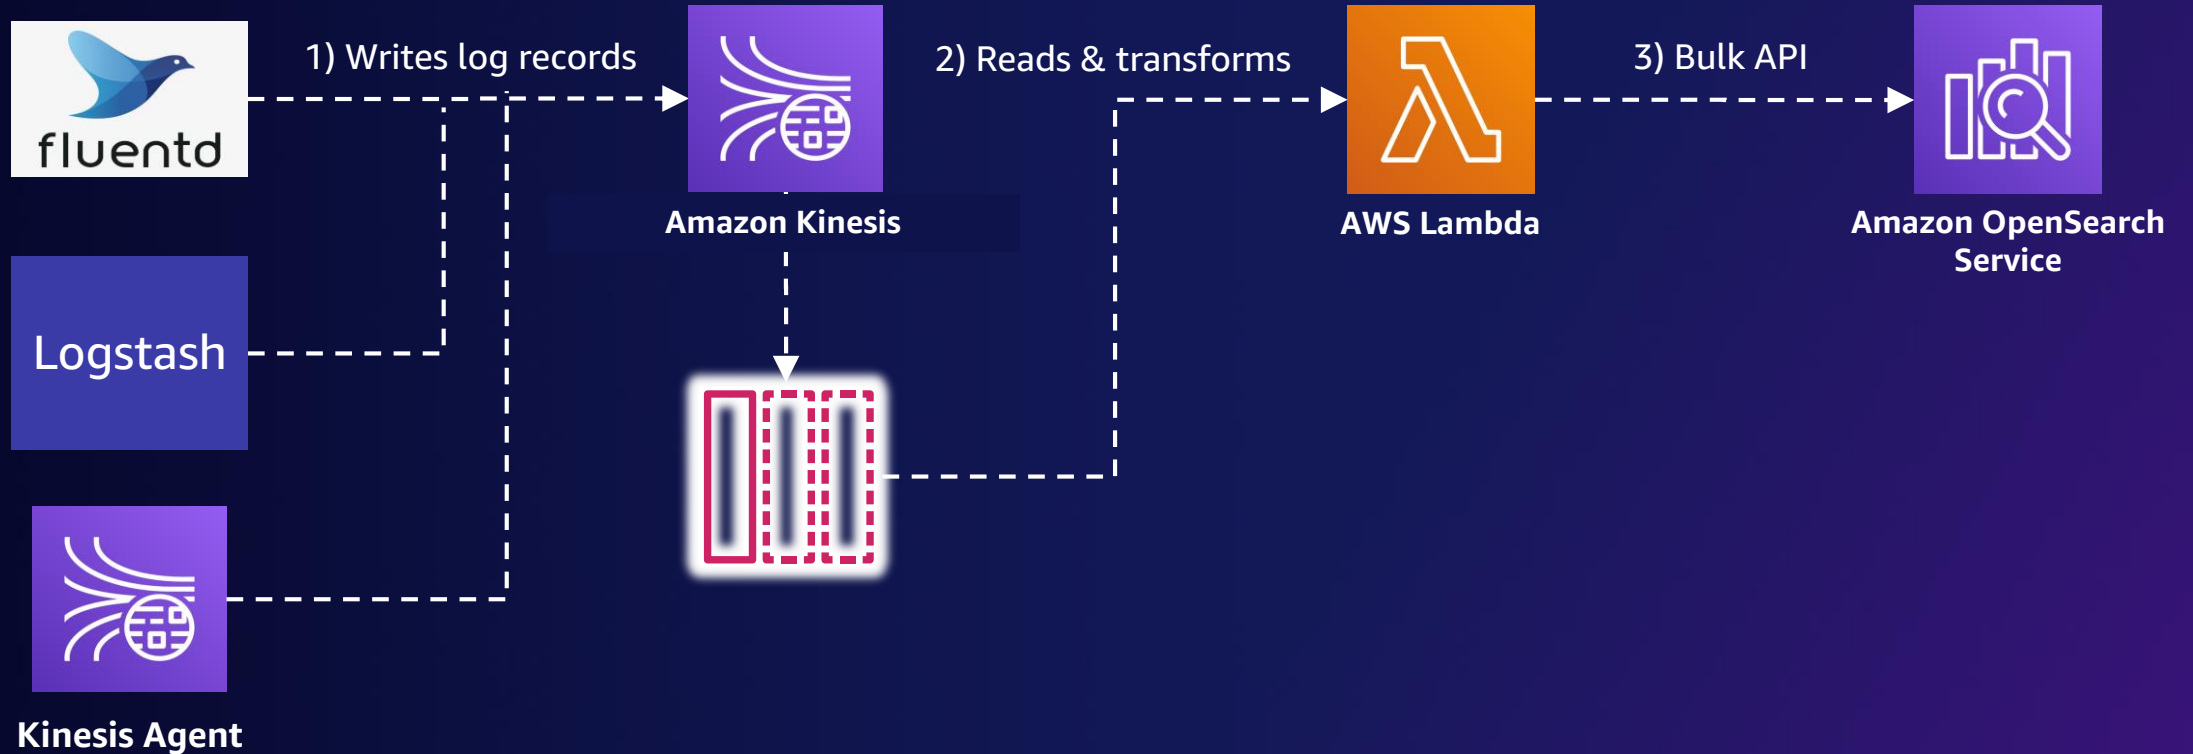
Low-cost buffering and delivery



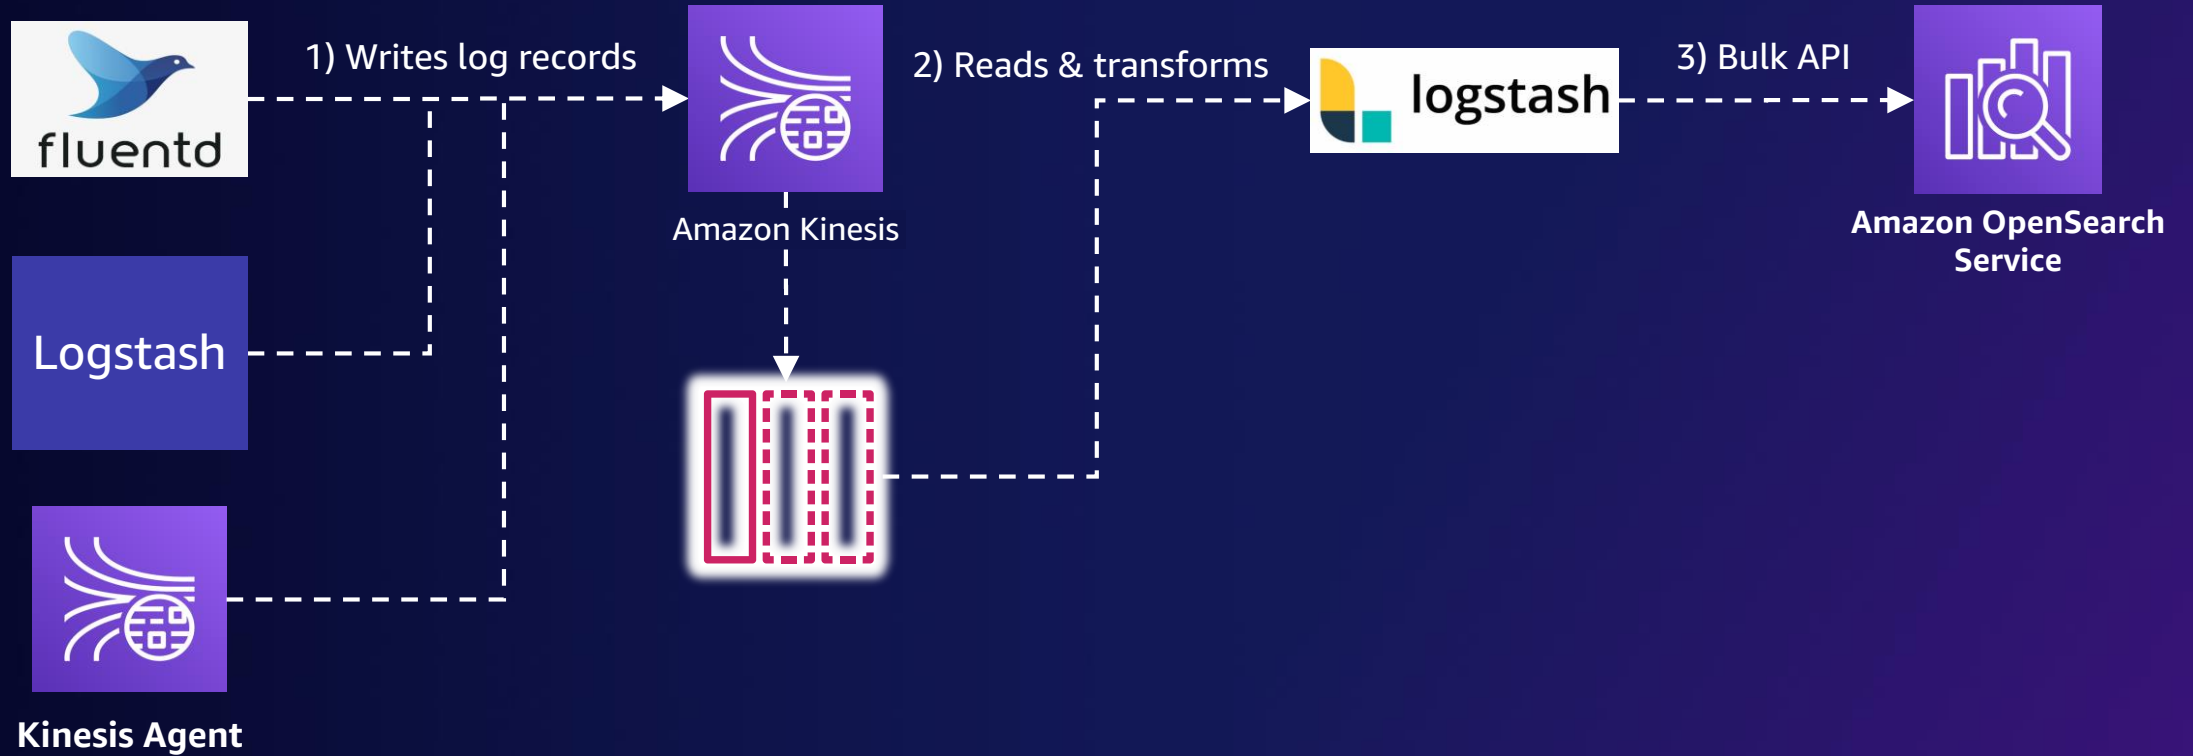
Amazon CloudWatch logs approach



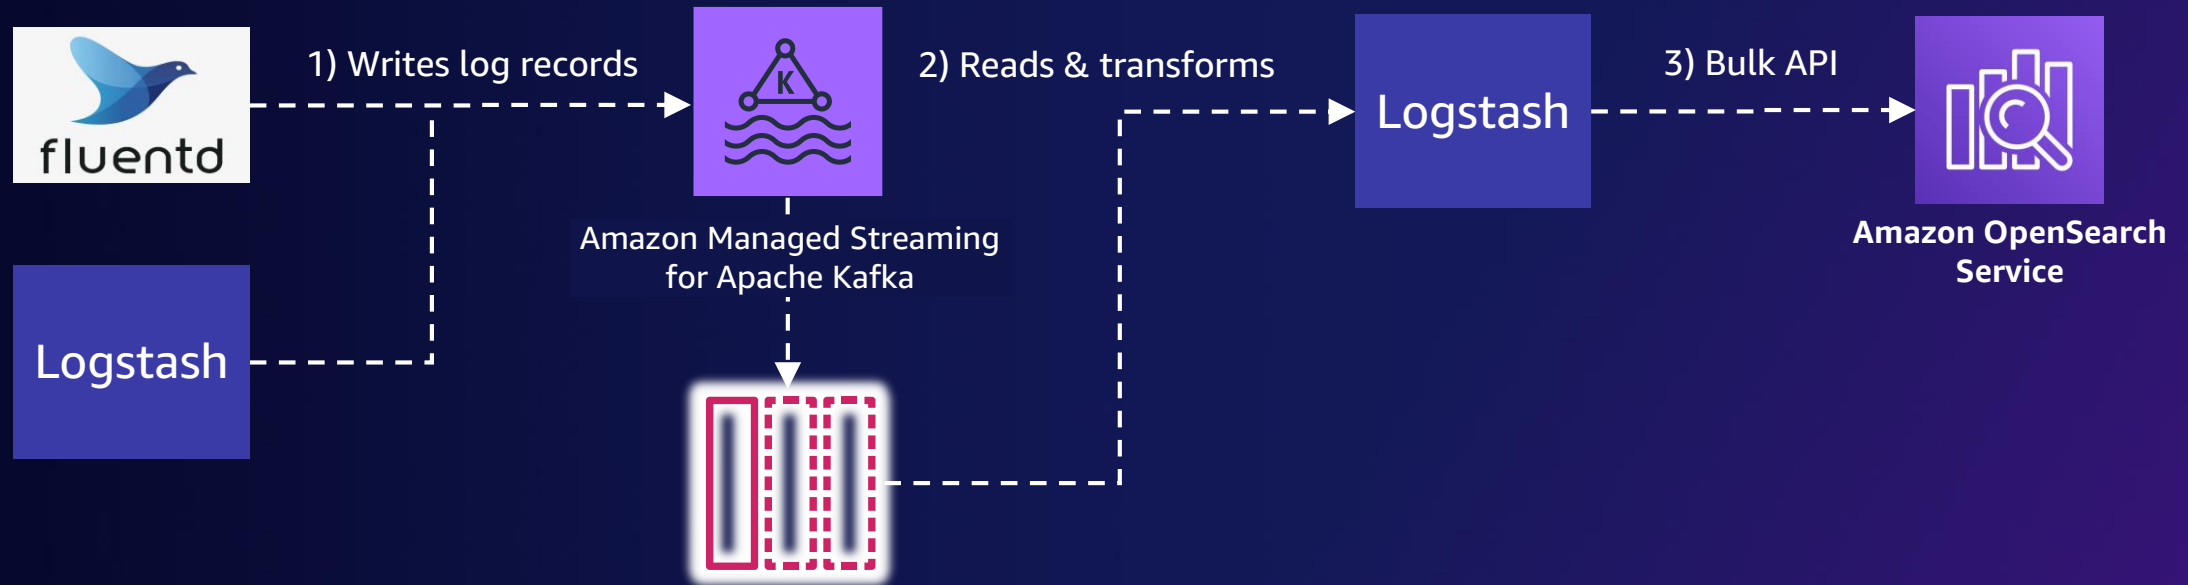
Amazon Kinesis – AWS Lambda approach



Amazon Kinesis – Logstash approach

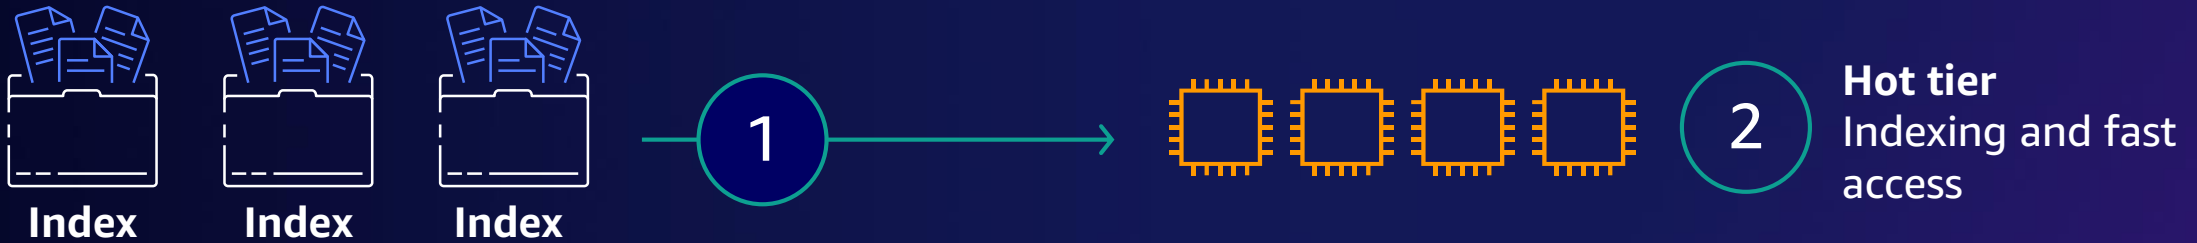


Amazon MSK – Logstash approach



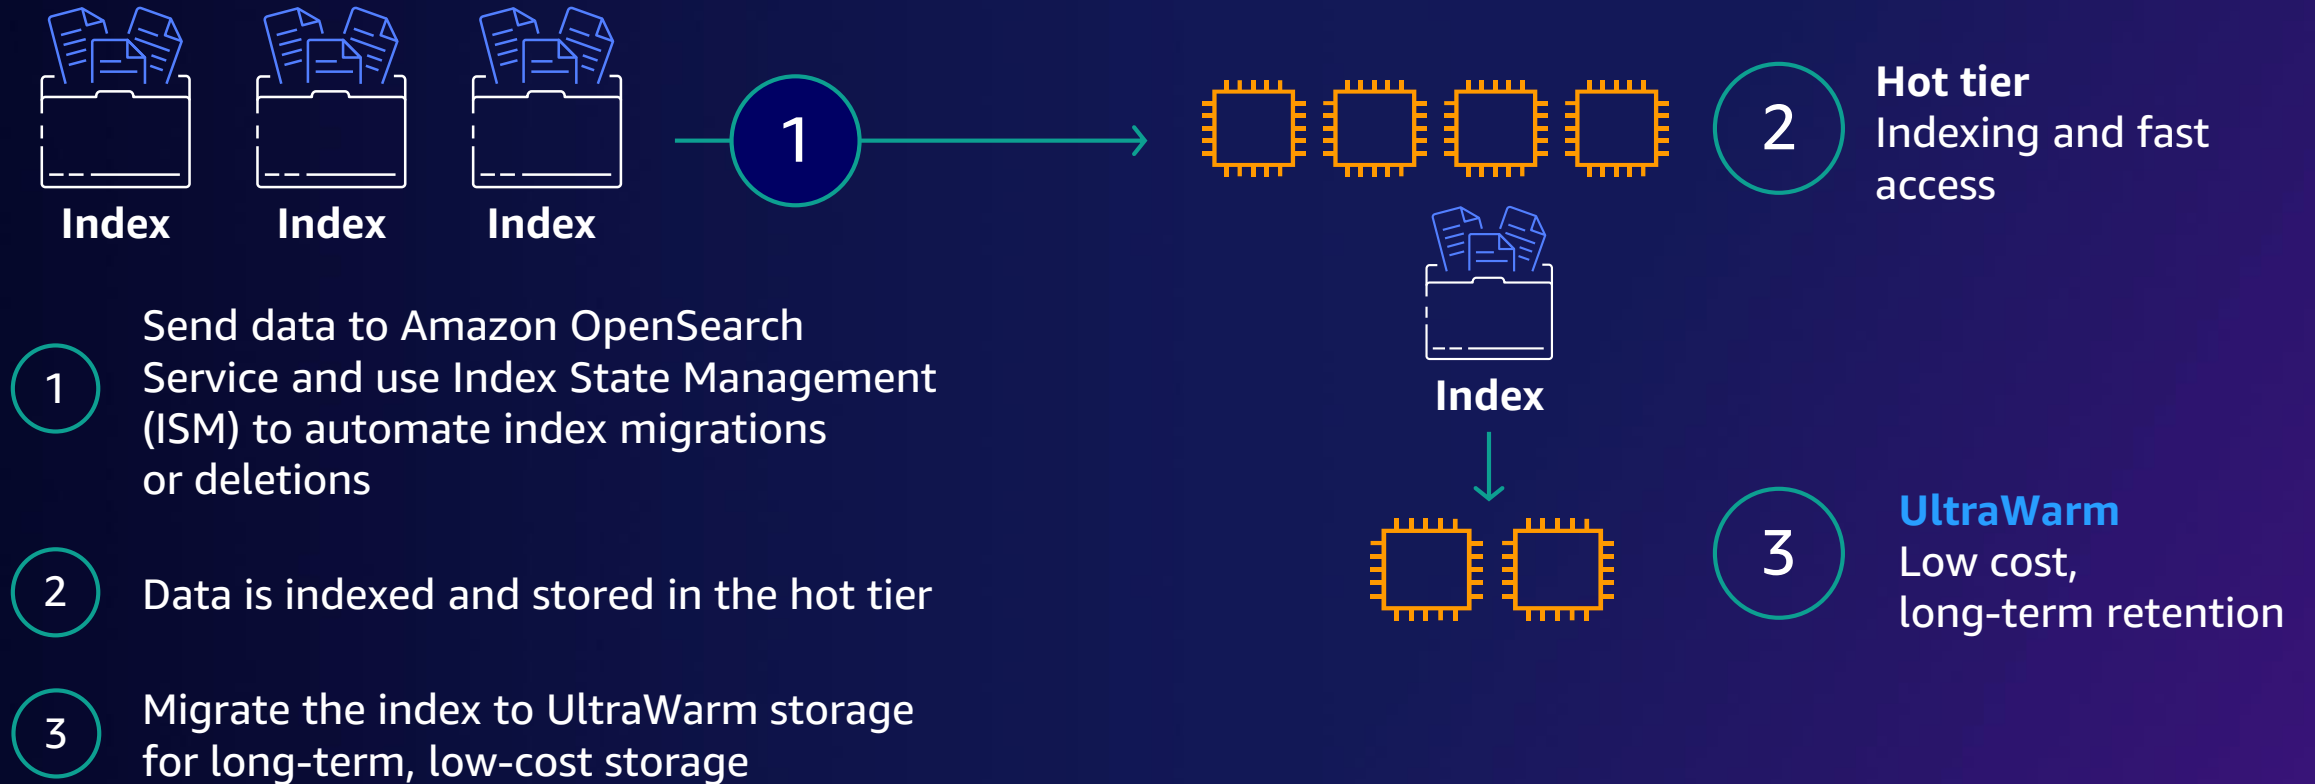
Data management

Data lifecycle in Amazon OpenSearch Service

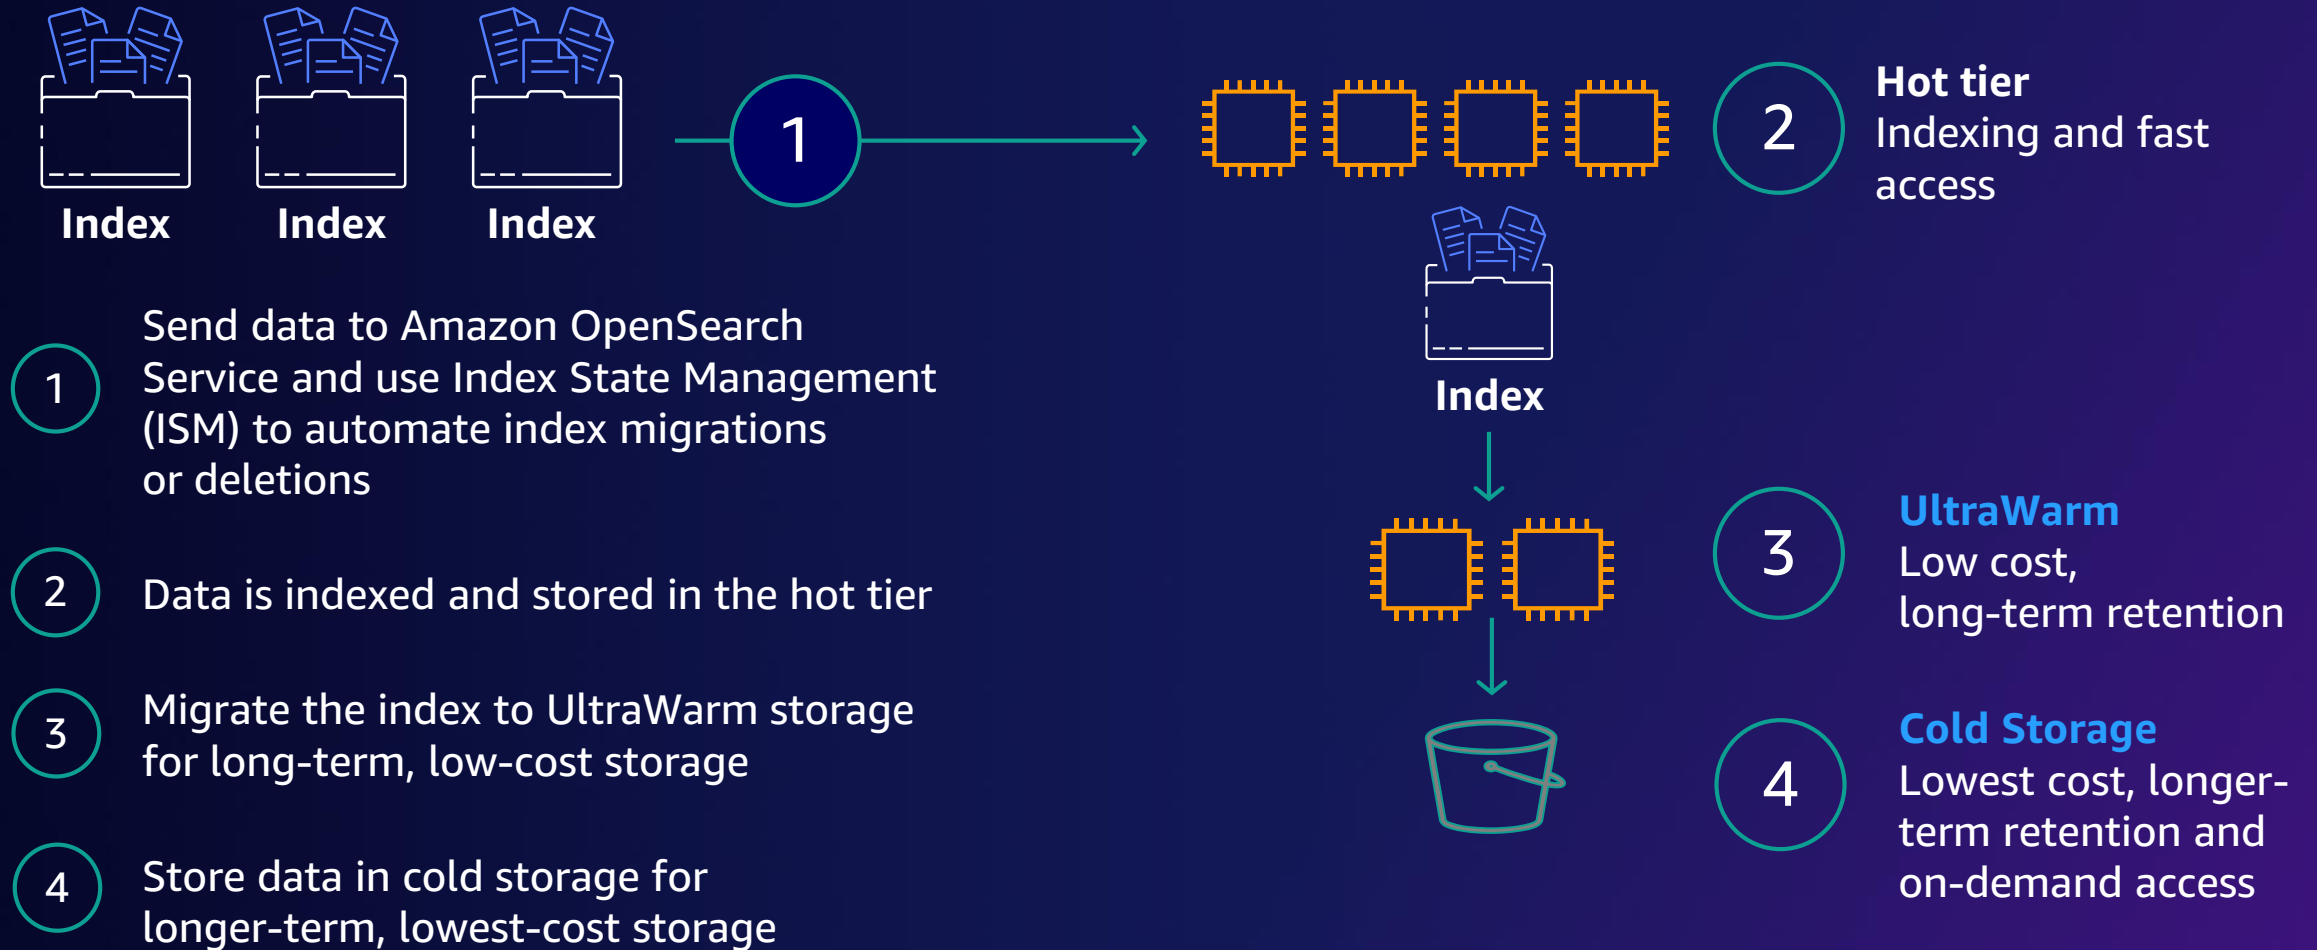


- 1 Send data to Amazon OpenSearch Service and use Index State Management (ISM) to automate index migrations or deletions
- 2 Data is indexed and stored in the hot tier

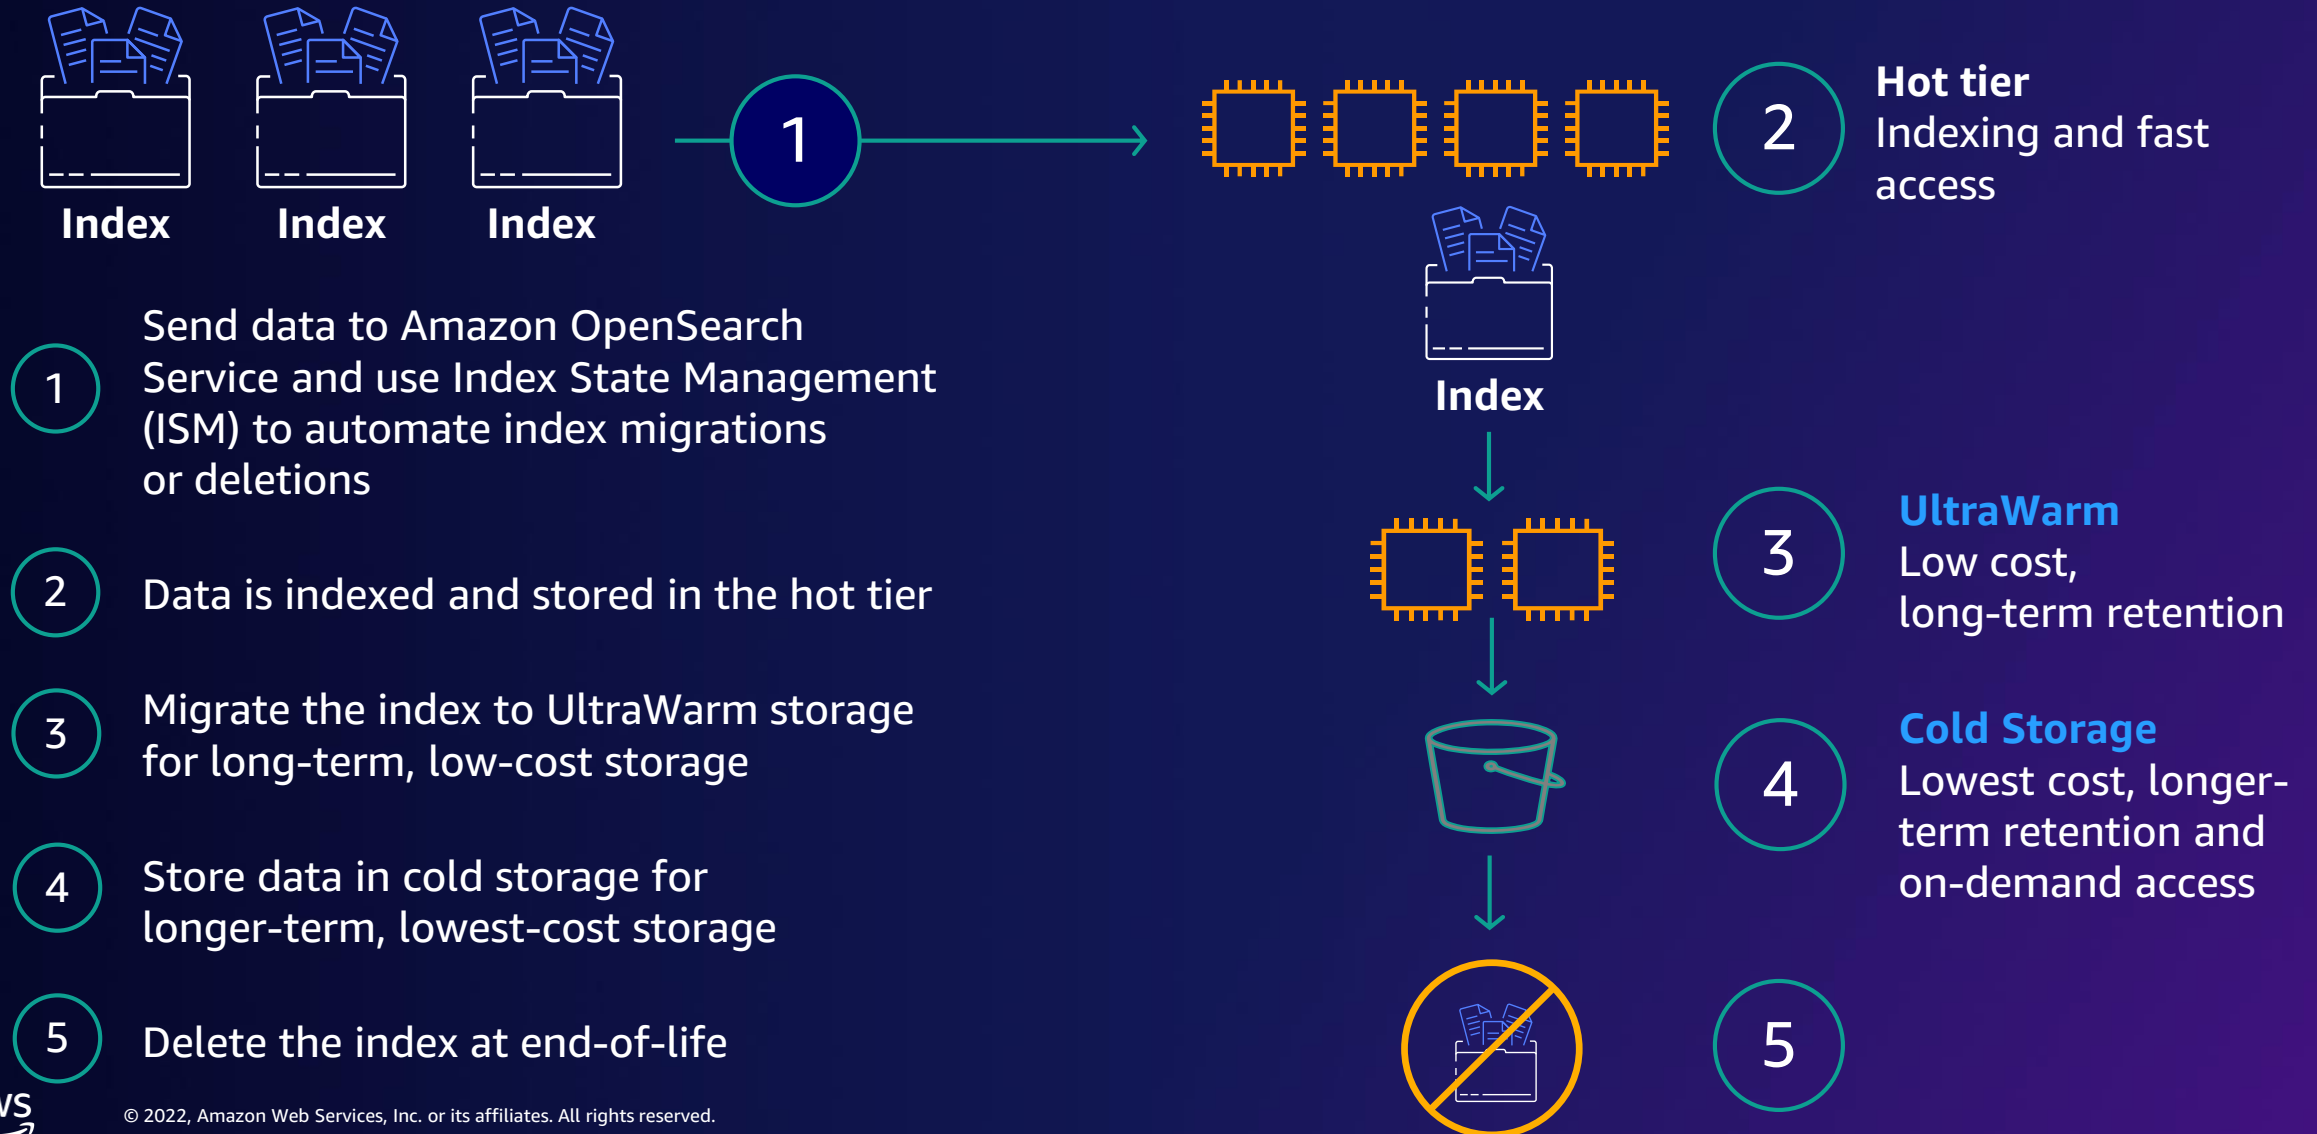
Data lifecycle in Amazon OpenSearch Service



Data lifecycle in Amazon OpenSearch Service



Data lifecycle in Amazon OpenSearch Service



Sizing and capacity planning

Capacity planning

Gather information

- Data per day
- Number of indexes
- Concurrency
- Rps (Docs, Qs)
- Index/query complexity
- Durability need

Plan for storage

- Ephemeral/EBS
- Replication
- Retention
- UltraWarm
- Cold storage

Storage is a solid starting point

Plan for compute

- Concurrency – number of readers/writers
- Active indexes/shards
- Total request counts and latencies

$\text{CPUs} = 1.5 * \text{active shards}$

Your configuration determines capacity

Instance type

- Deploy instances based on storage and compute needs

Instance count

- Add instances for increased parallelism

Storage

- Index data (primary and replica shards) is stored on disk

Shard count

- Shards are the units of work and storage

Logs workloads are storage-driven. Search workloads are CPU/JVM-driven

Step 1: figure out storage need

Storage needed =
 $\text{Source/day} * 1.1 * 2 * \text{retention} * 1.15$

Search: 100 GB of data needs 250 GB of storage

Logs: 1 TB daily of source data needs 18 TB of storage for 7 days of retention

Step 2: figure out shard count

Primary shards =
$$\text{Index size} / \text{target shard size}$$

Logs, use 50 GB max.
Search evaluate 20–30 GB

Step 3: Set a template

```
*PUT
<endpoint>/_template/template1
{
  "index_patterns": ["logs*"],
  "settings": {
    "number_of_shards": 50,
    "number_of_replicas": 1
  }
}
```

Step 4: Adjust for usage

$\text{vCPU} = 1.5 * \text{active shards}$

Active shards

Primaries for queries

Primaries and replicas for updates

E.g., 4 data streams at 1 TB daily
means 40 total shards (20 primary and
20 replica) active, so make sure to have
64 total CPUs

Example

100 GB of logs per day

1 index pattern

7 days hot

23 days UltraWarm

335 days Cold

Total hot storage

$$100 * 1.1 * 7 * 2 * 1.15 = 1.77 \text{ TiB}$$

Min CPUs hot = 6

Total Warm Storage

$$100 * 1.1 * 23 = 2.53 \text{ TiB}$$

Total Cold Storage

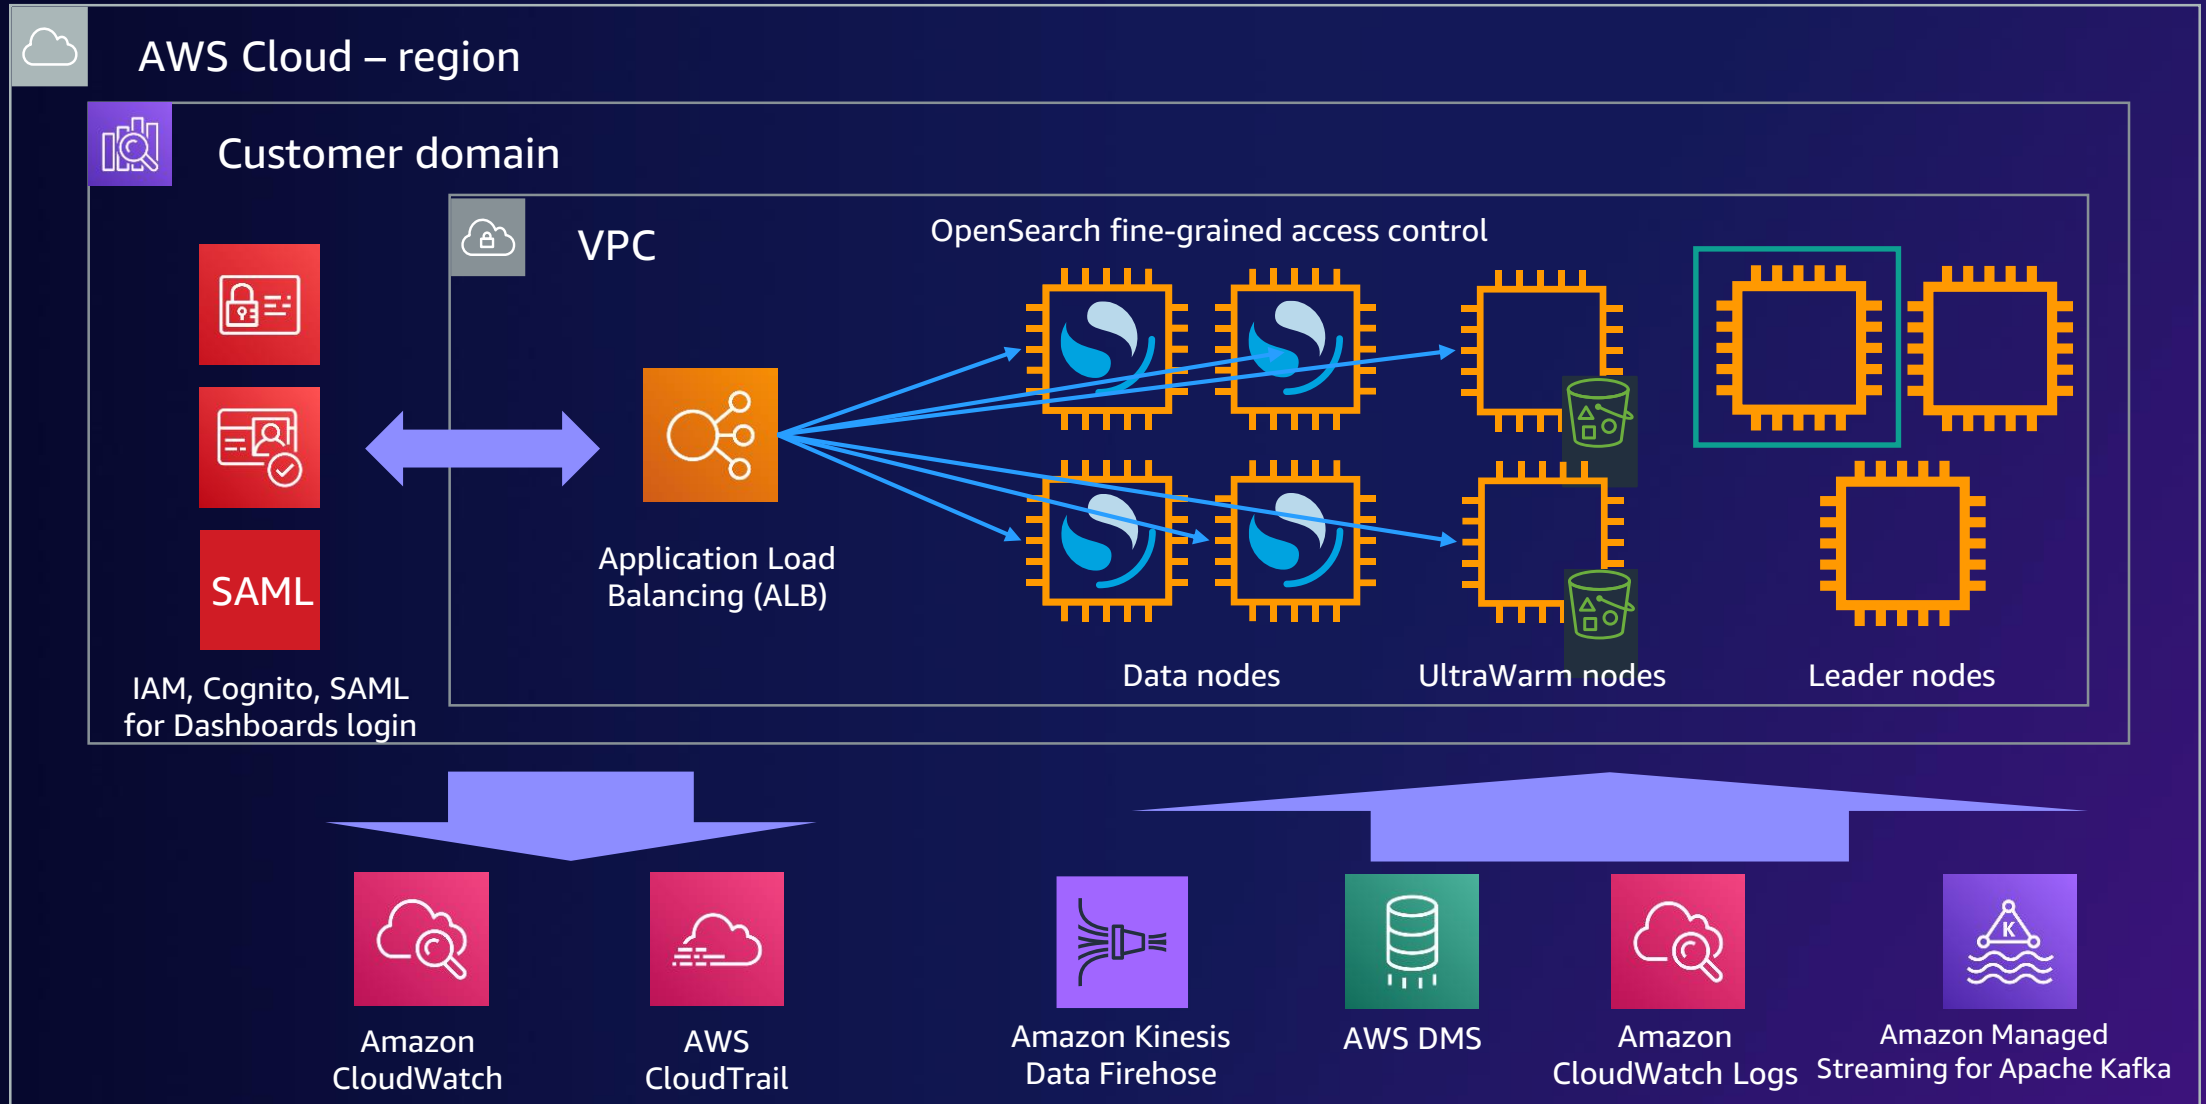
$$100 * 1.1 * 335 = 36.85 \text{ TiB}$$

3xM6g.xlarge hot (min for storage)

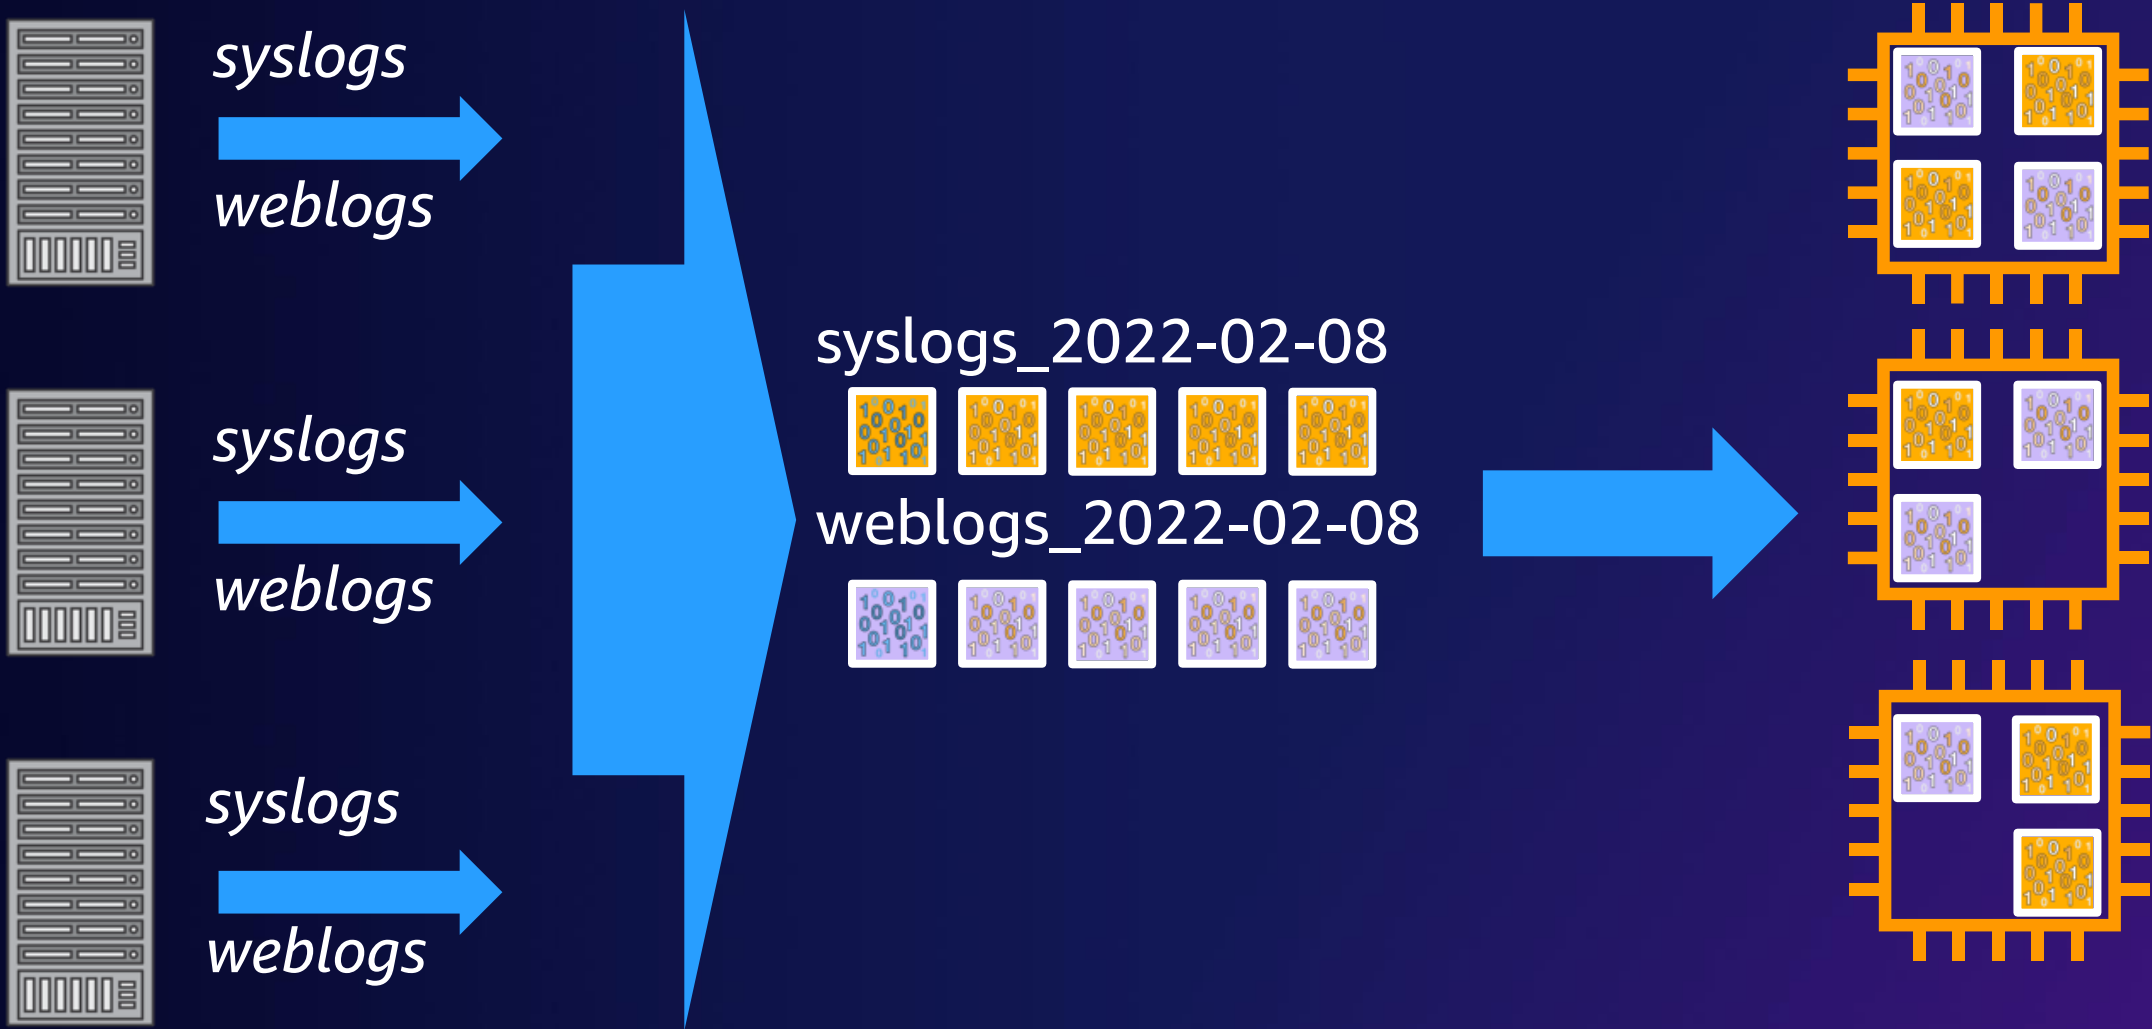
2xUltraWarm.medium

Request processing

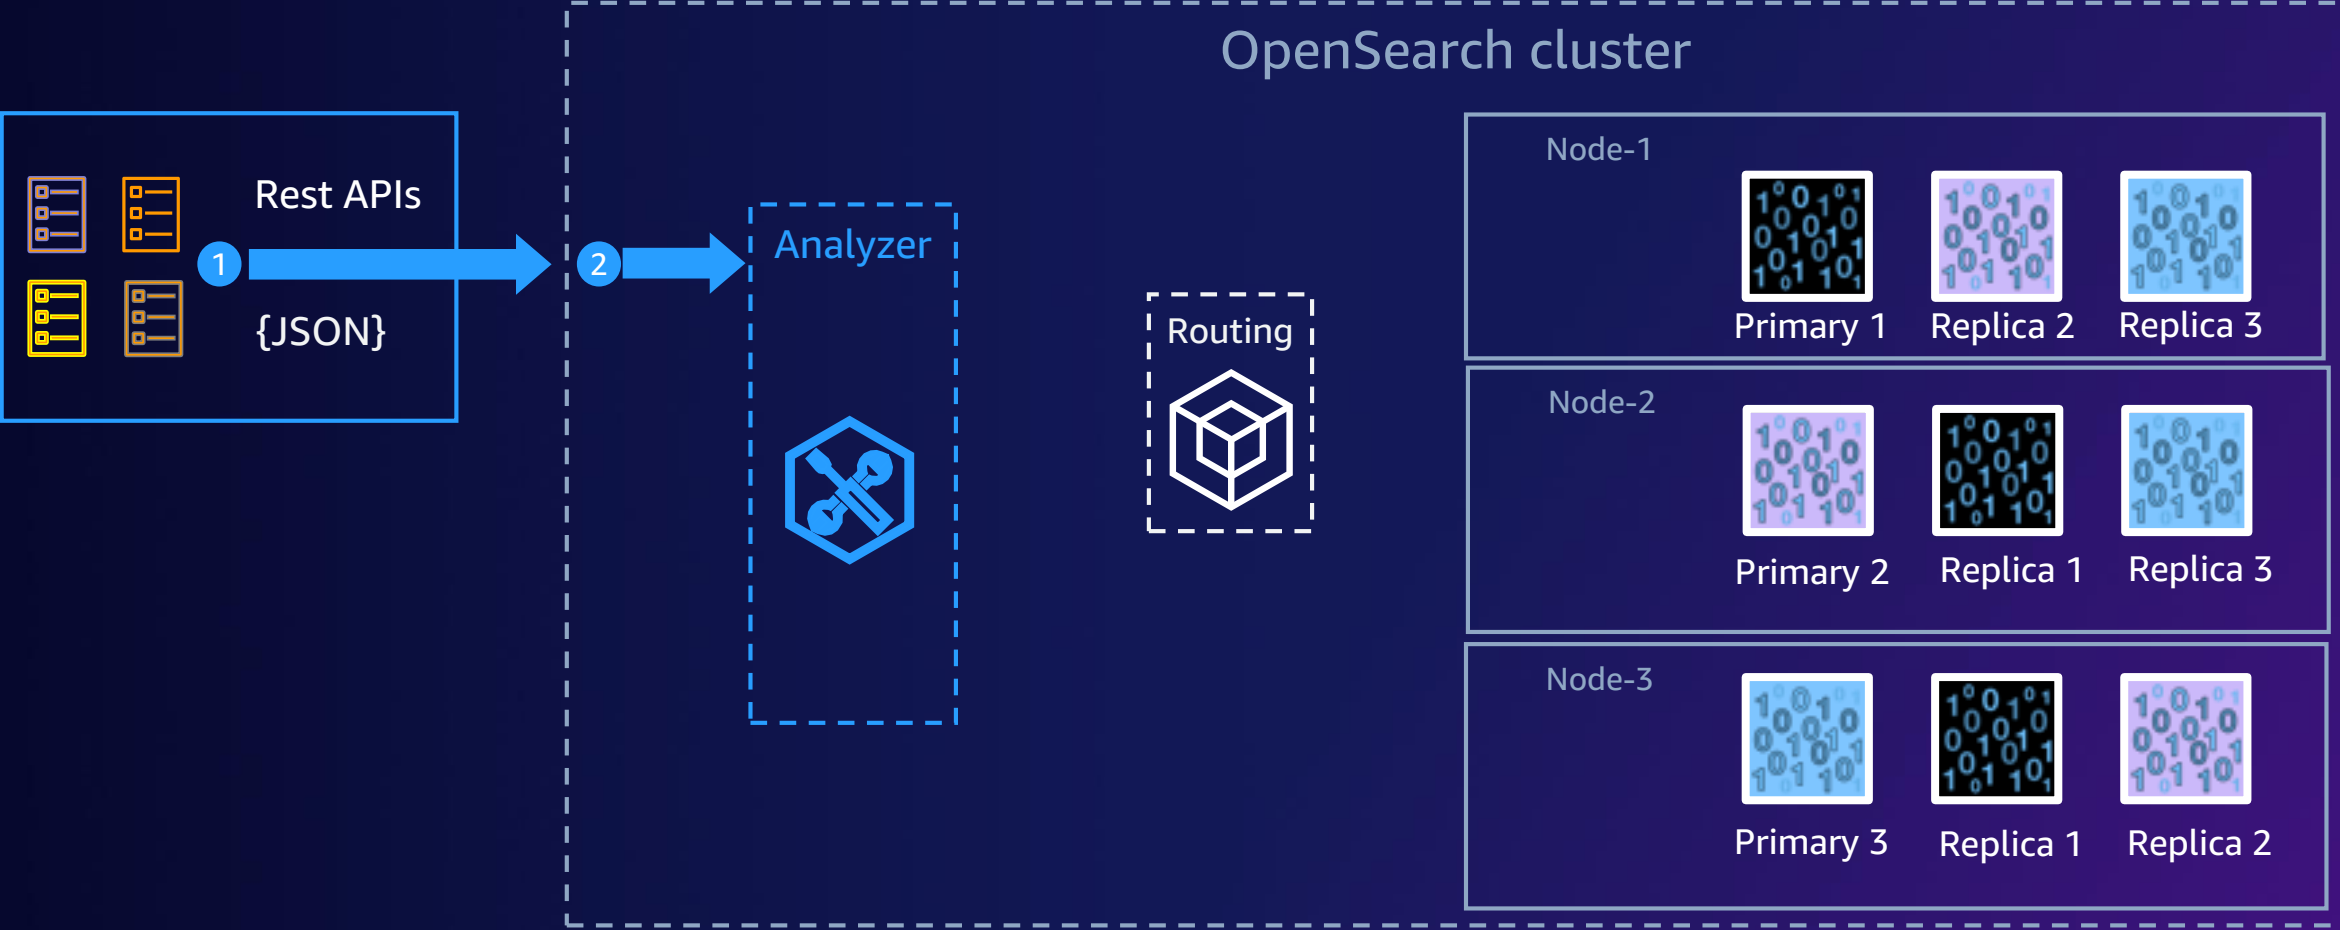
Amazon OpenSearch Service architecture



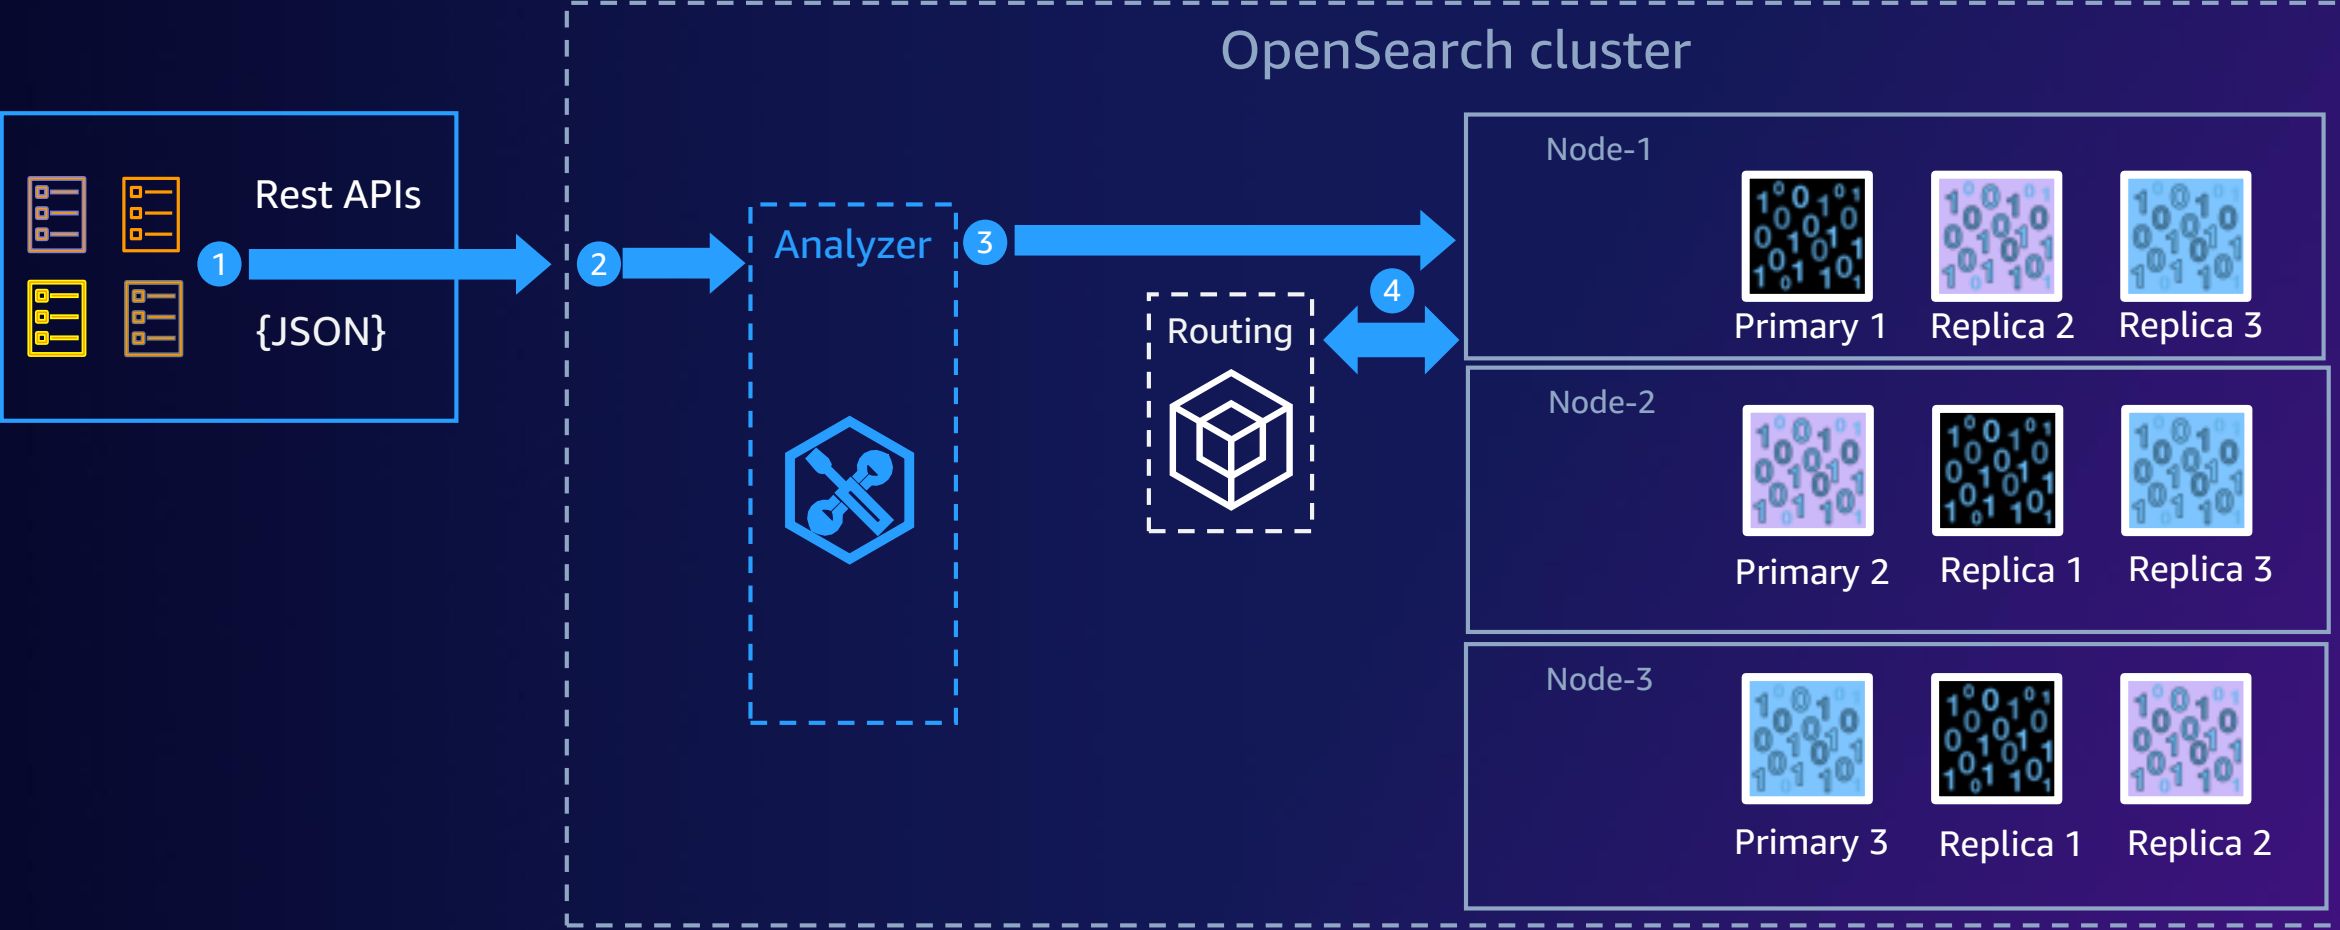
Indexes are distributed via shards (partitions)



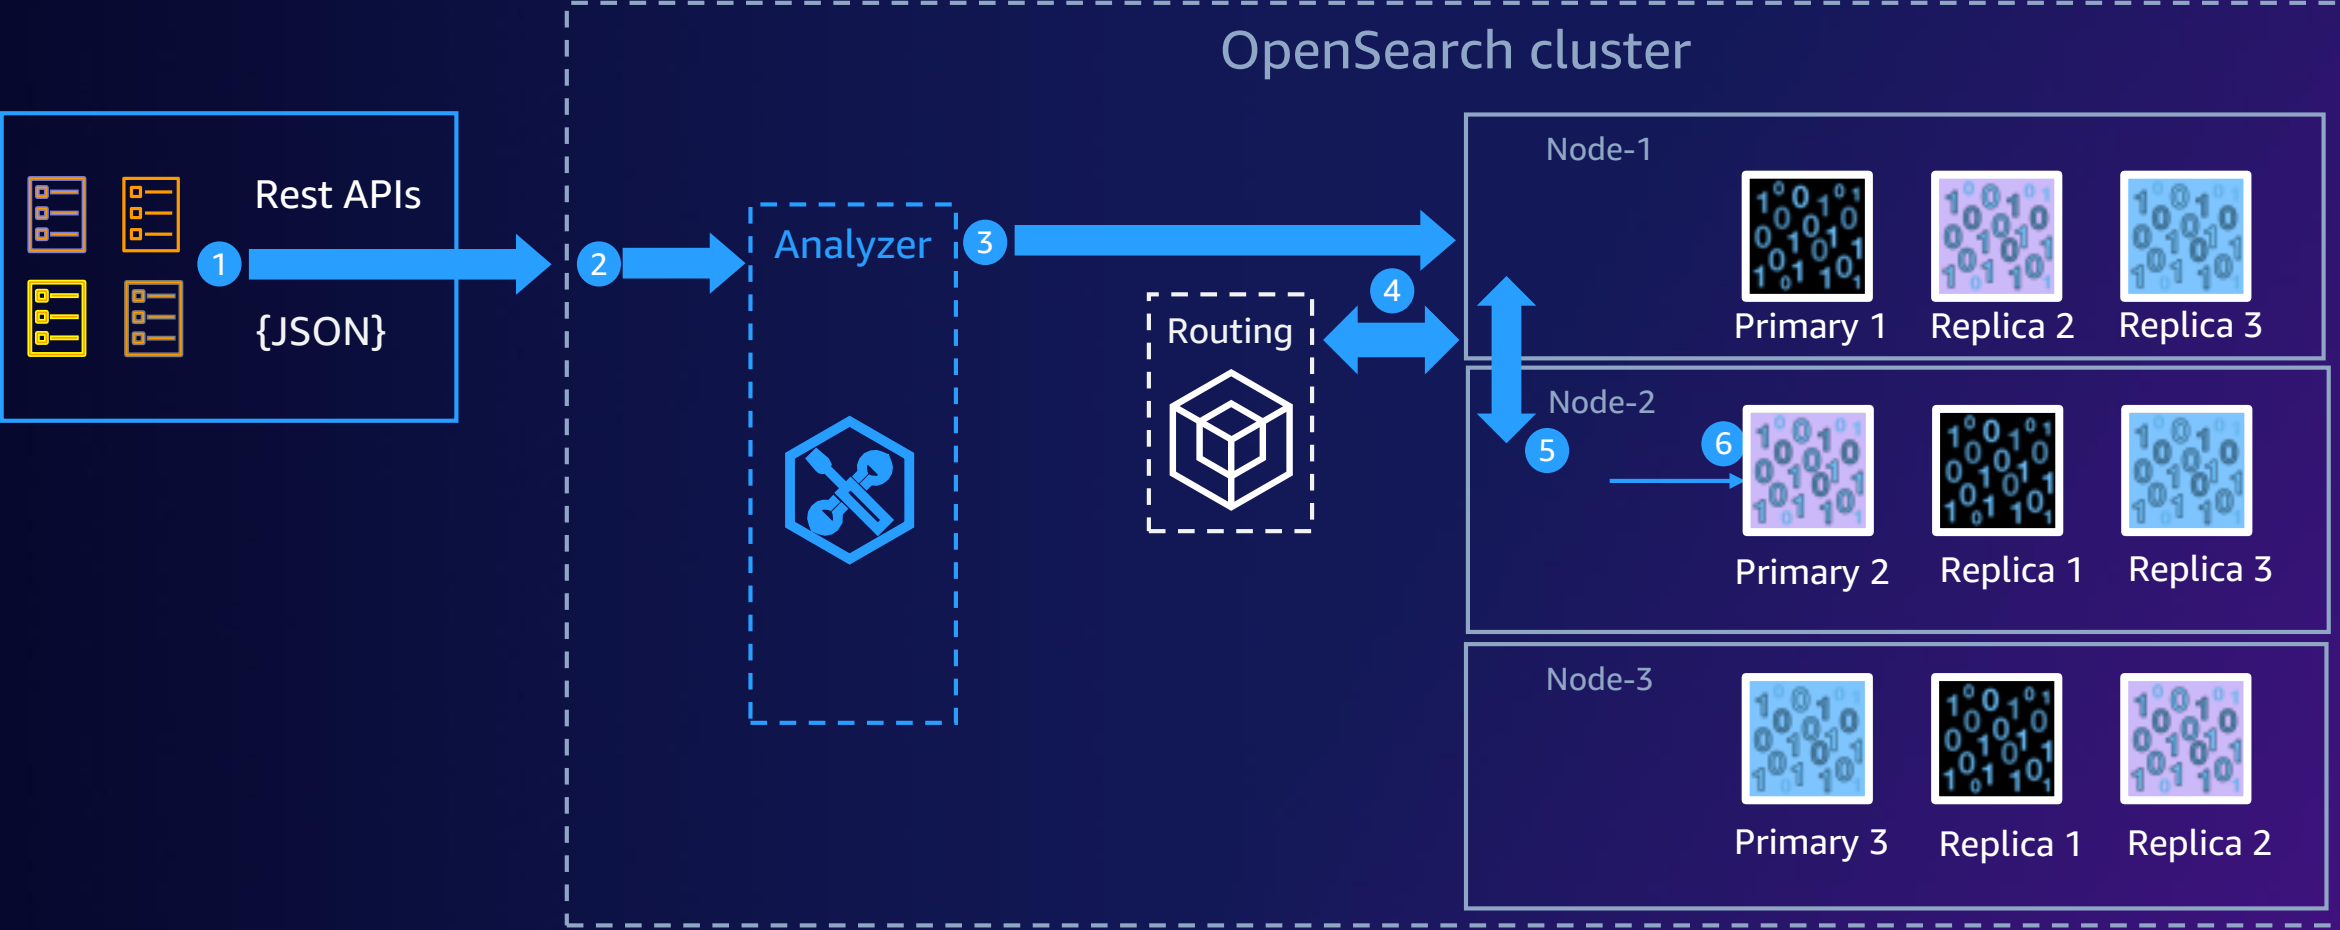
Shards are workers: indexing



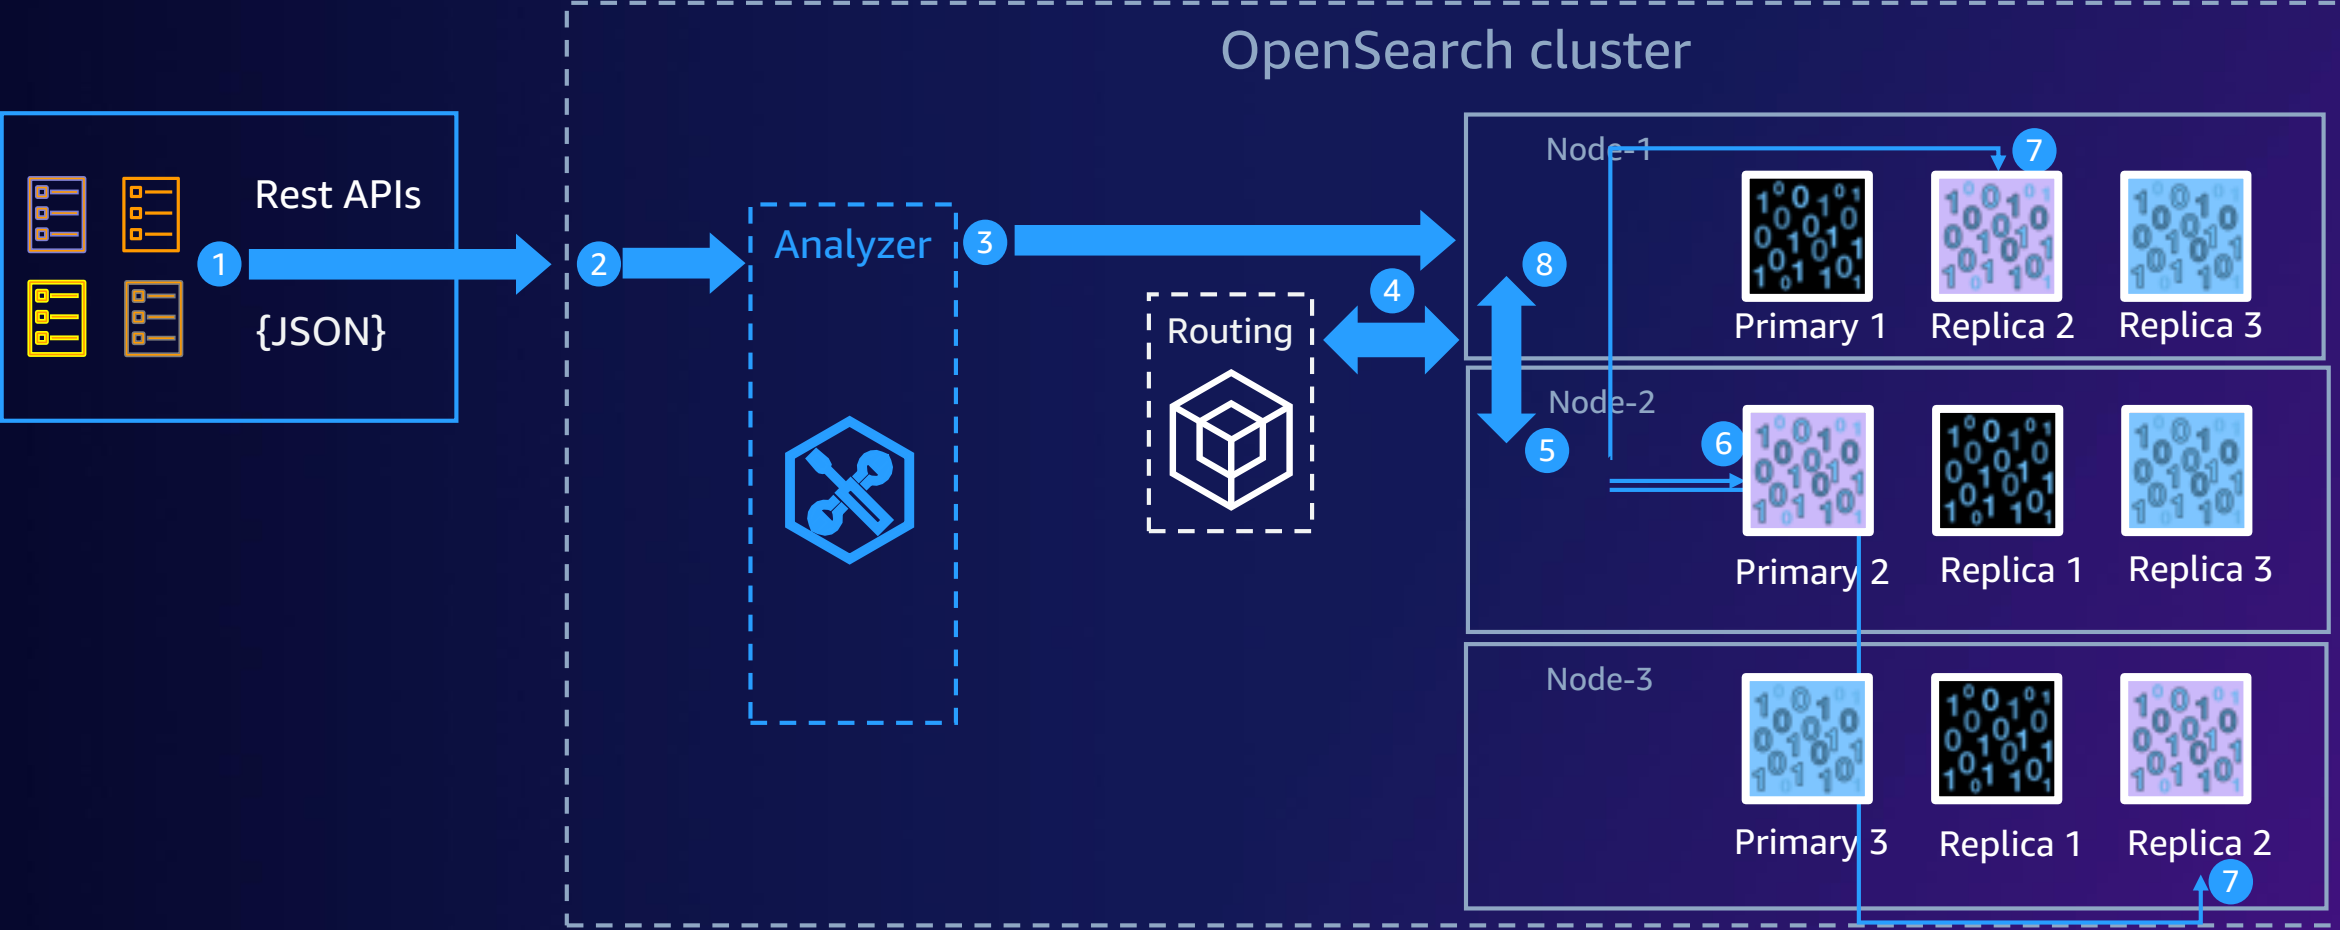
Shards are workers: Indexing



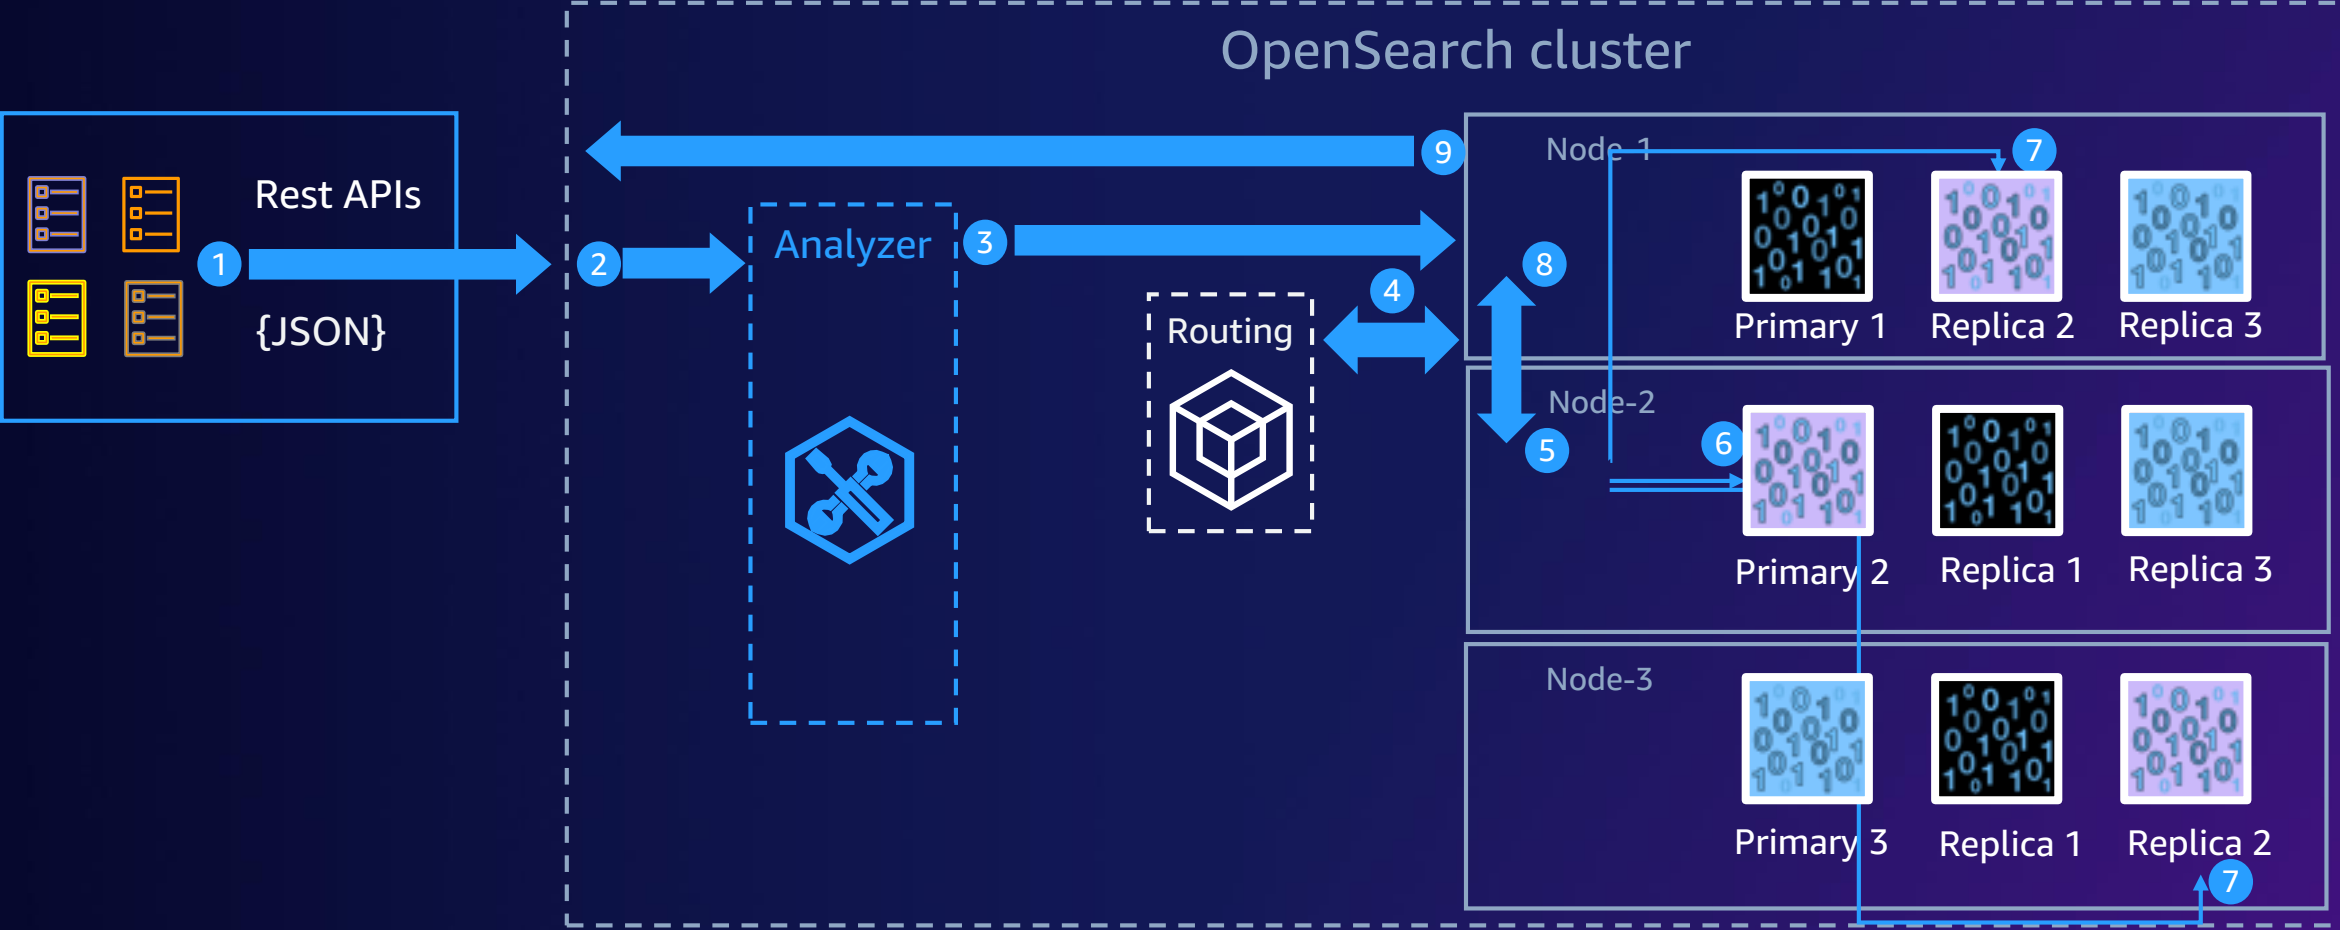
Shards are workers: Indexing



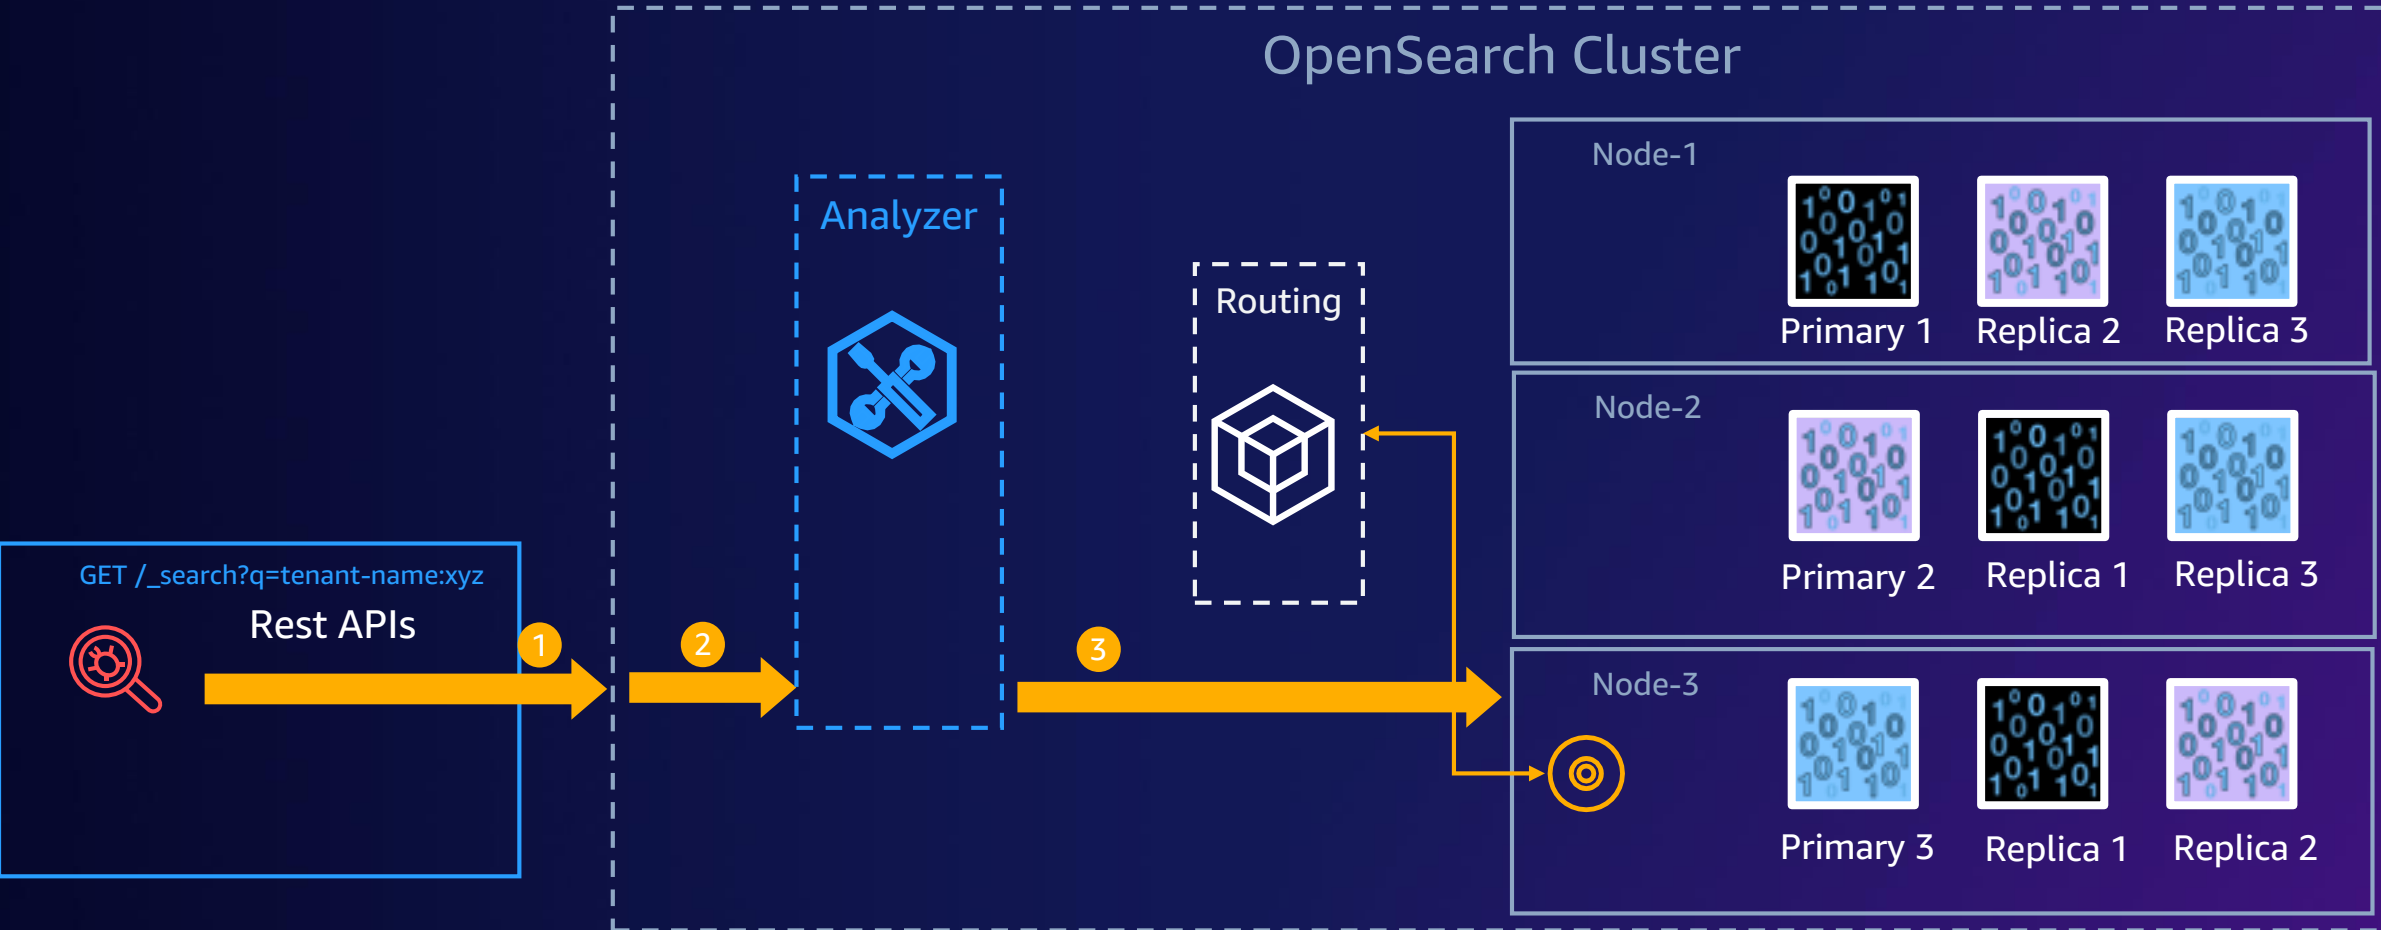
Shards are workers: Indexing



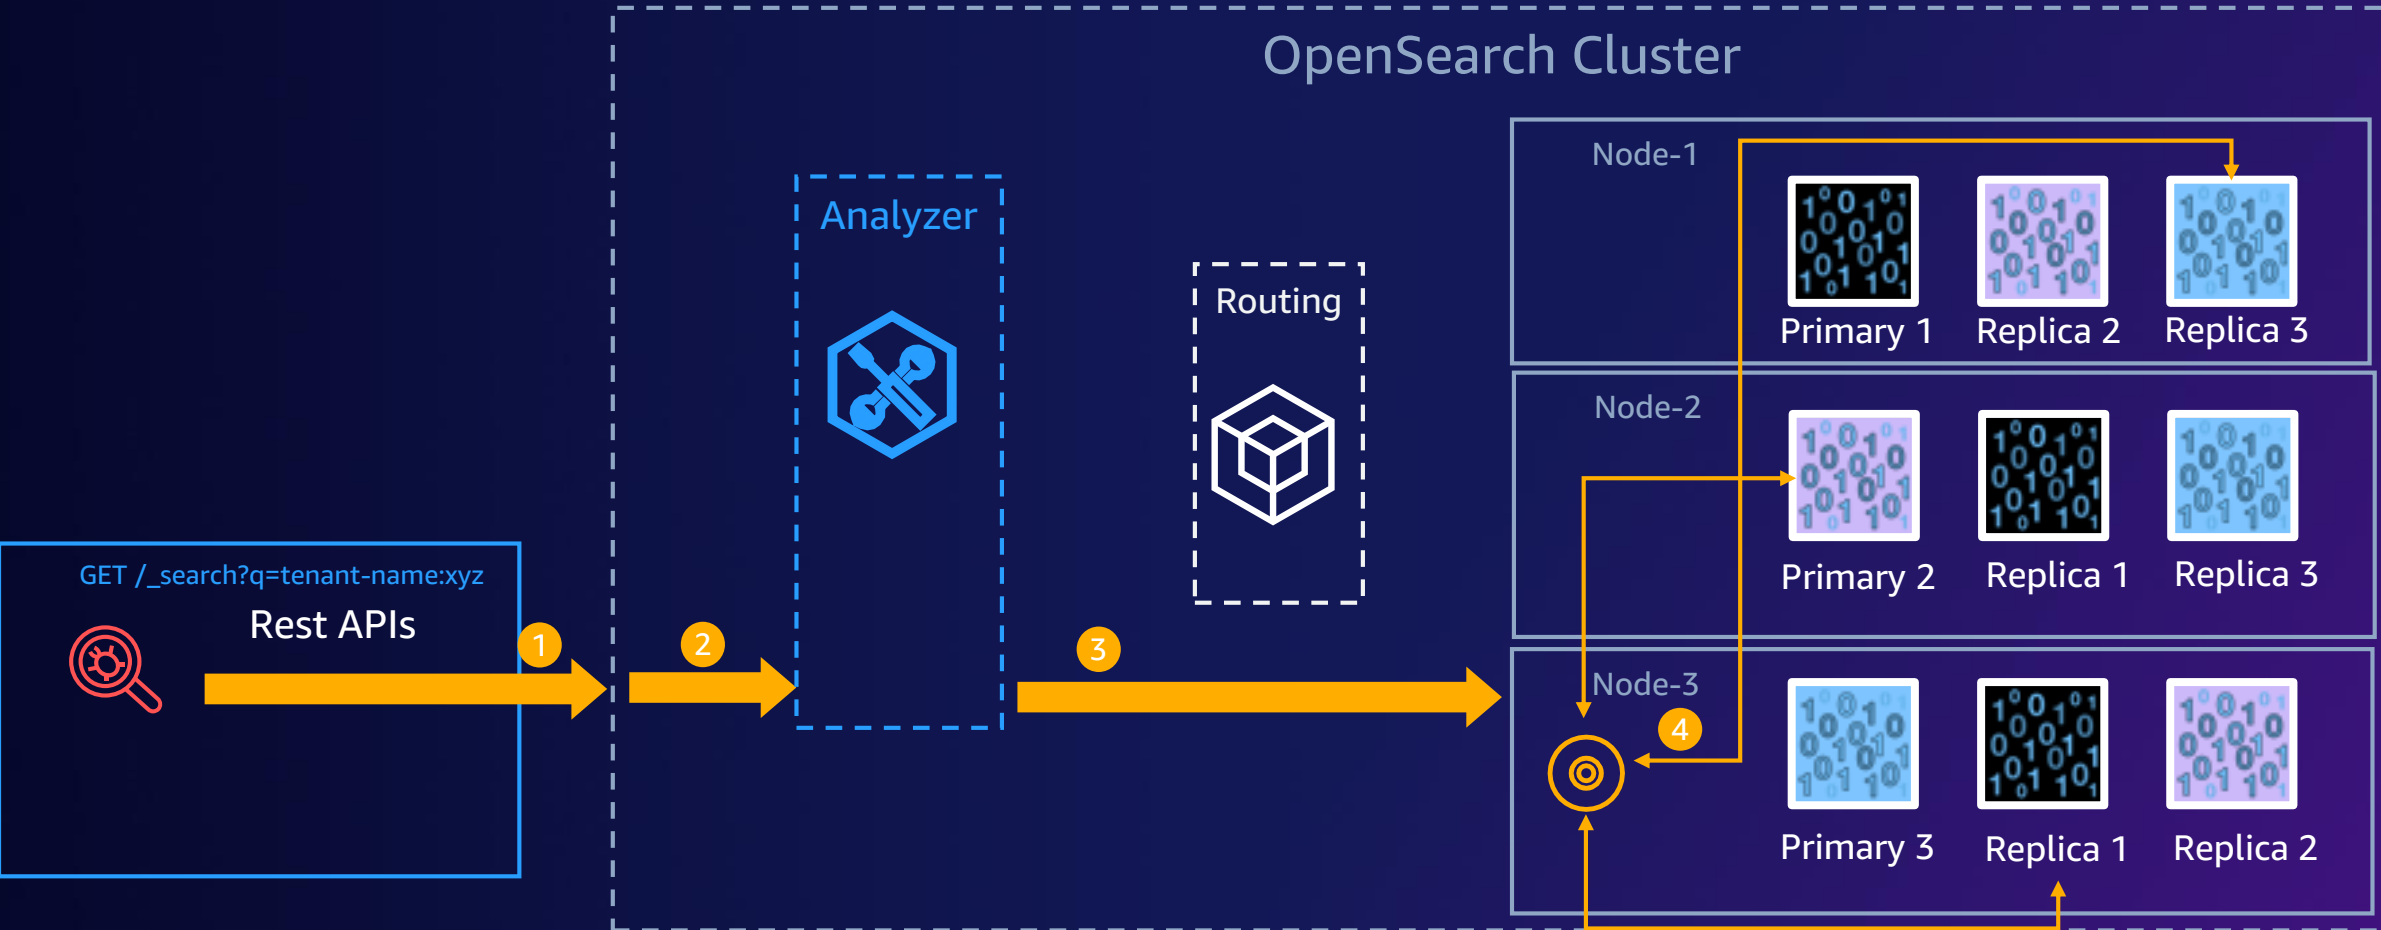
Shards are workers: Indexing



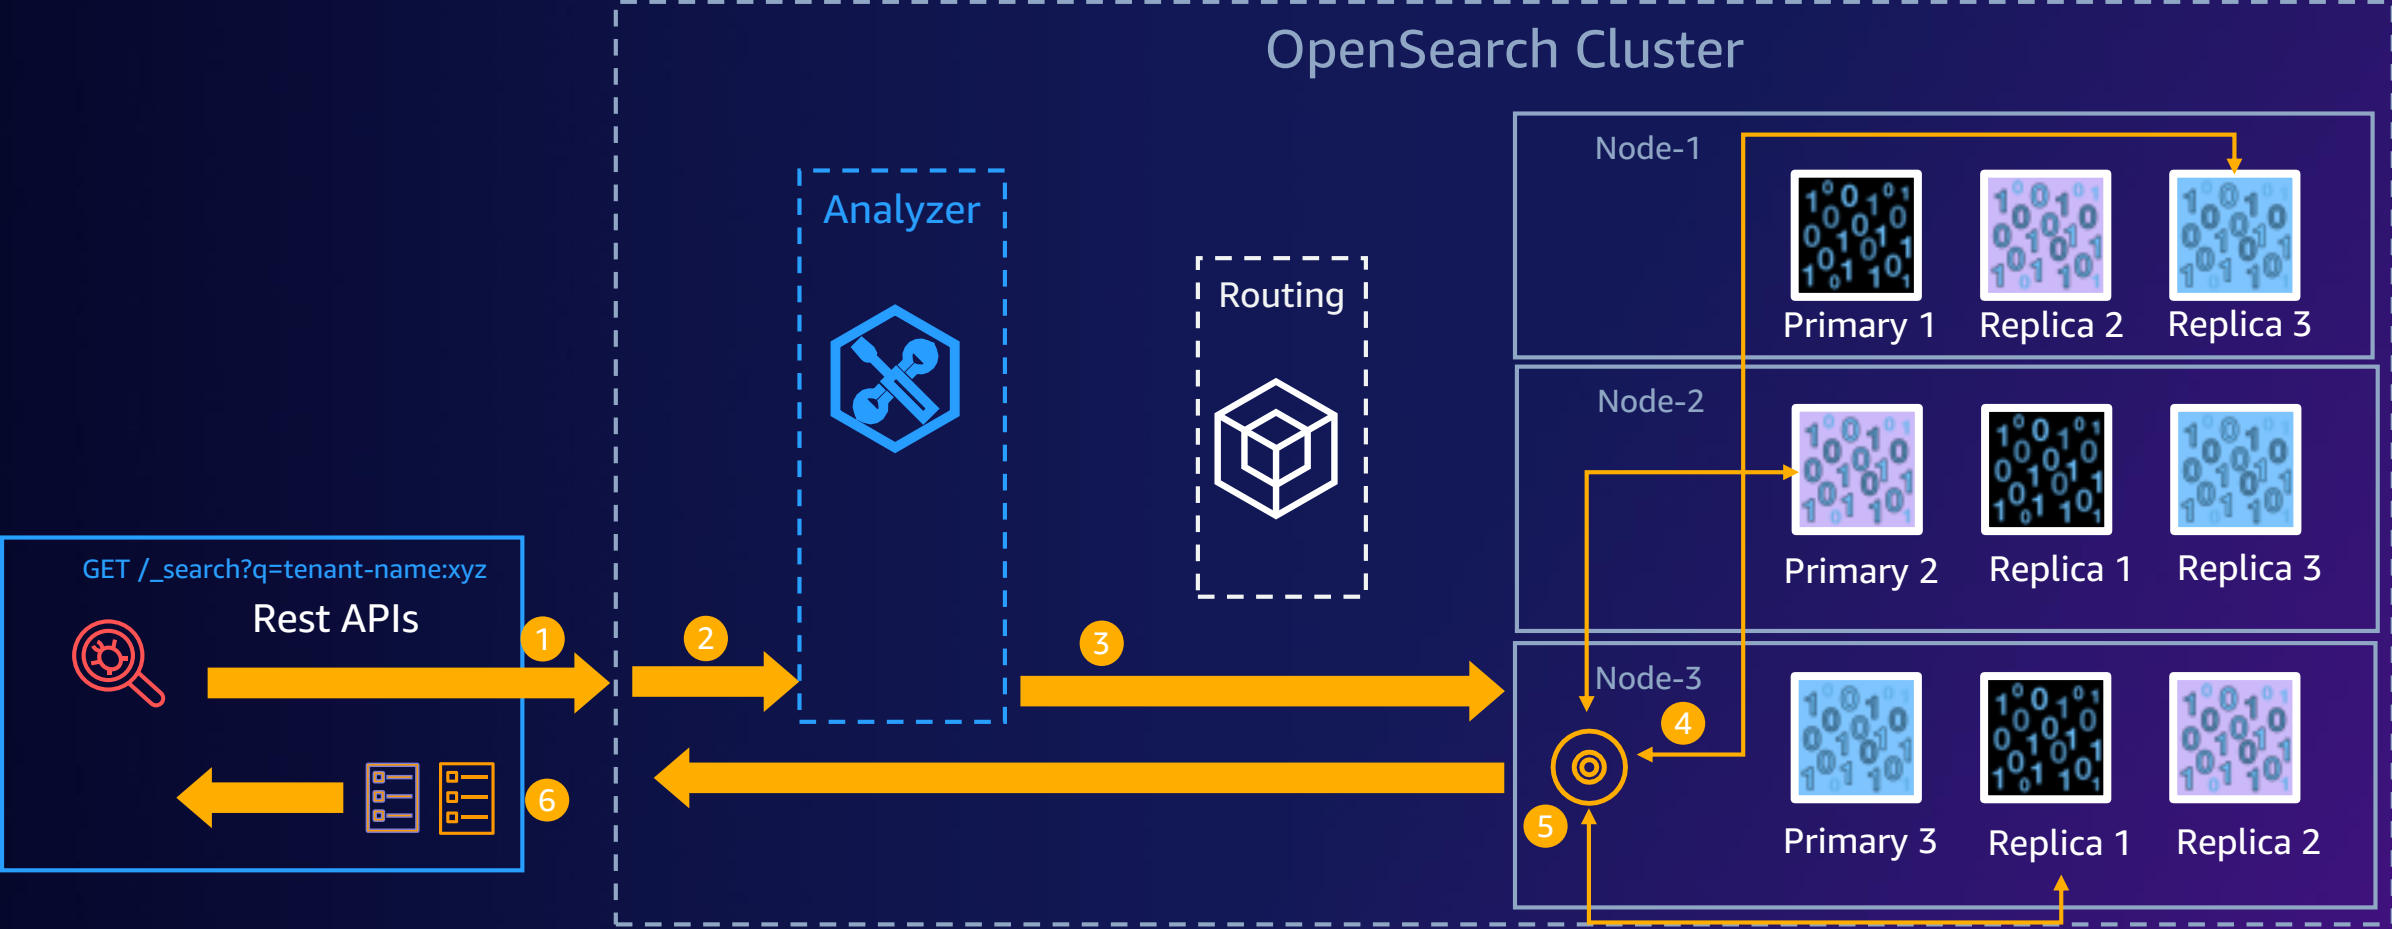
Shards are workers: Query



Shards are workers: Query

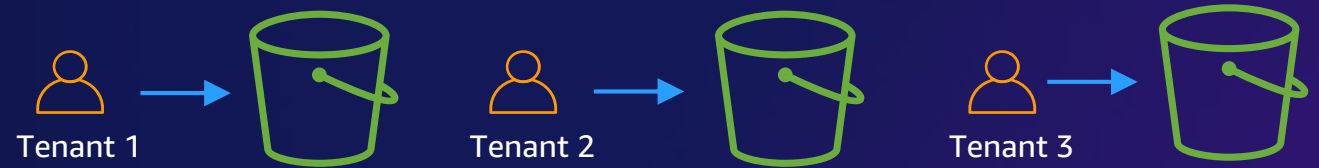


Shards are workers: Query

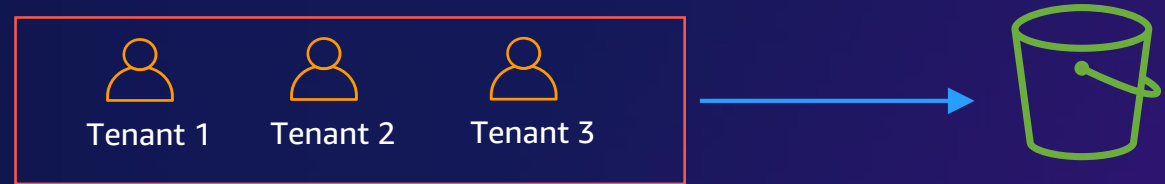


Different tenancy strategies

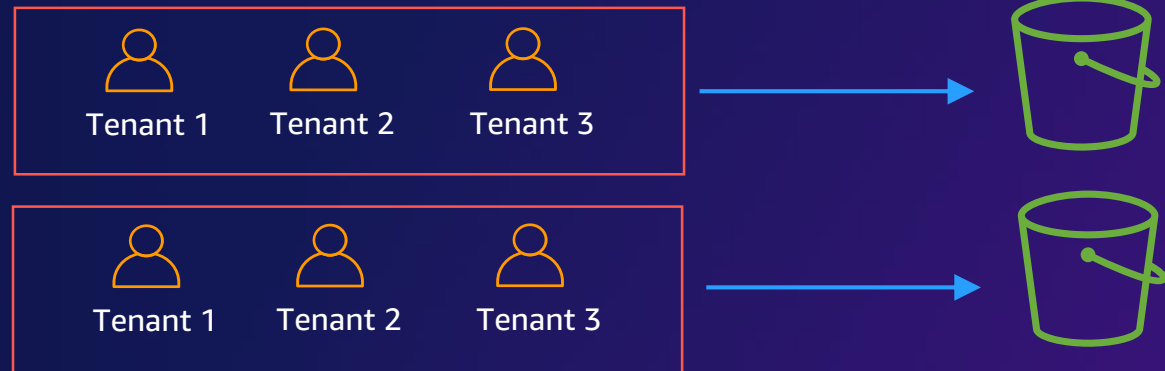
Siloed strategy



Pooled strategy



Hybrid strategy

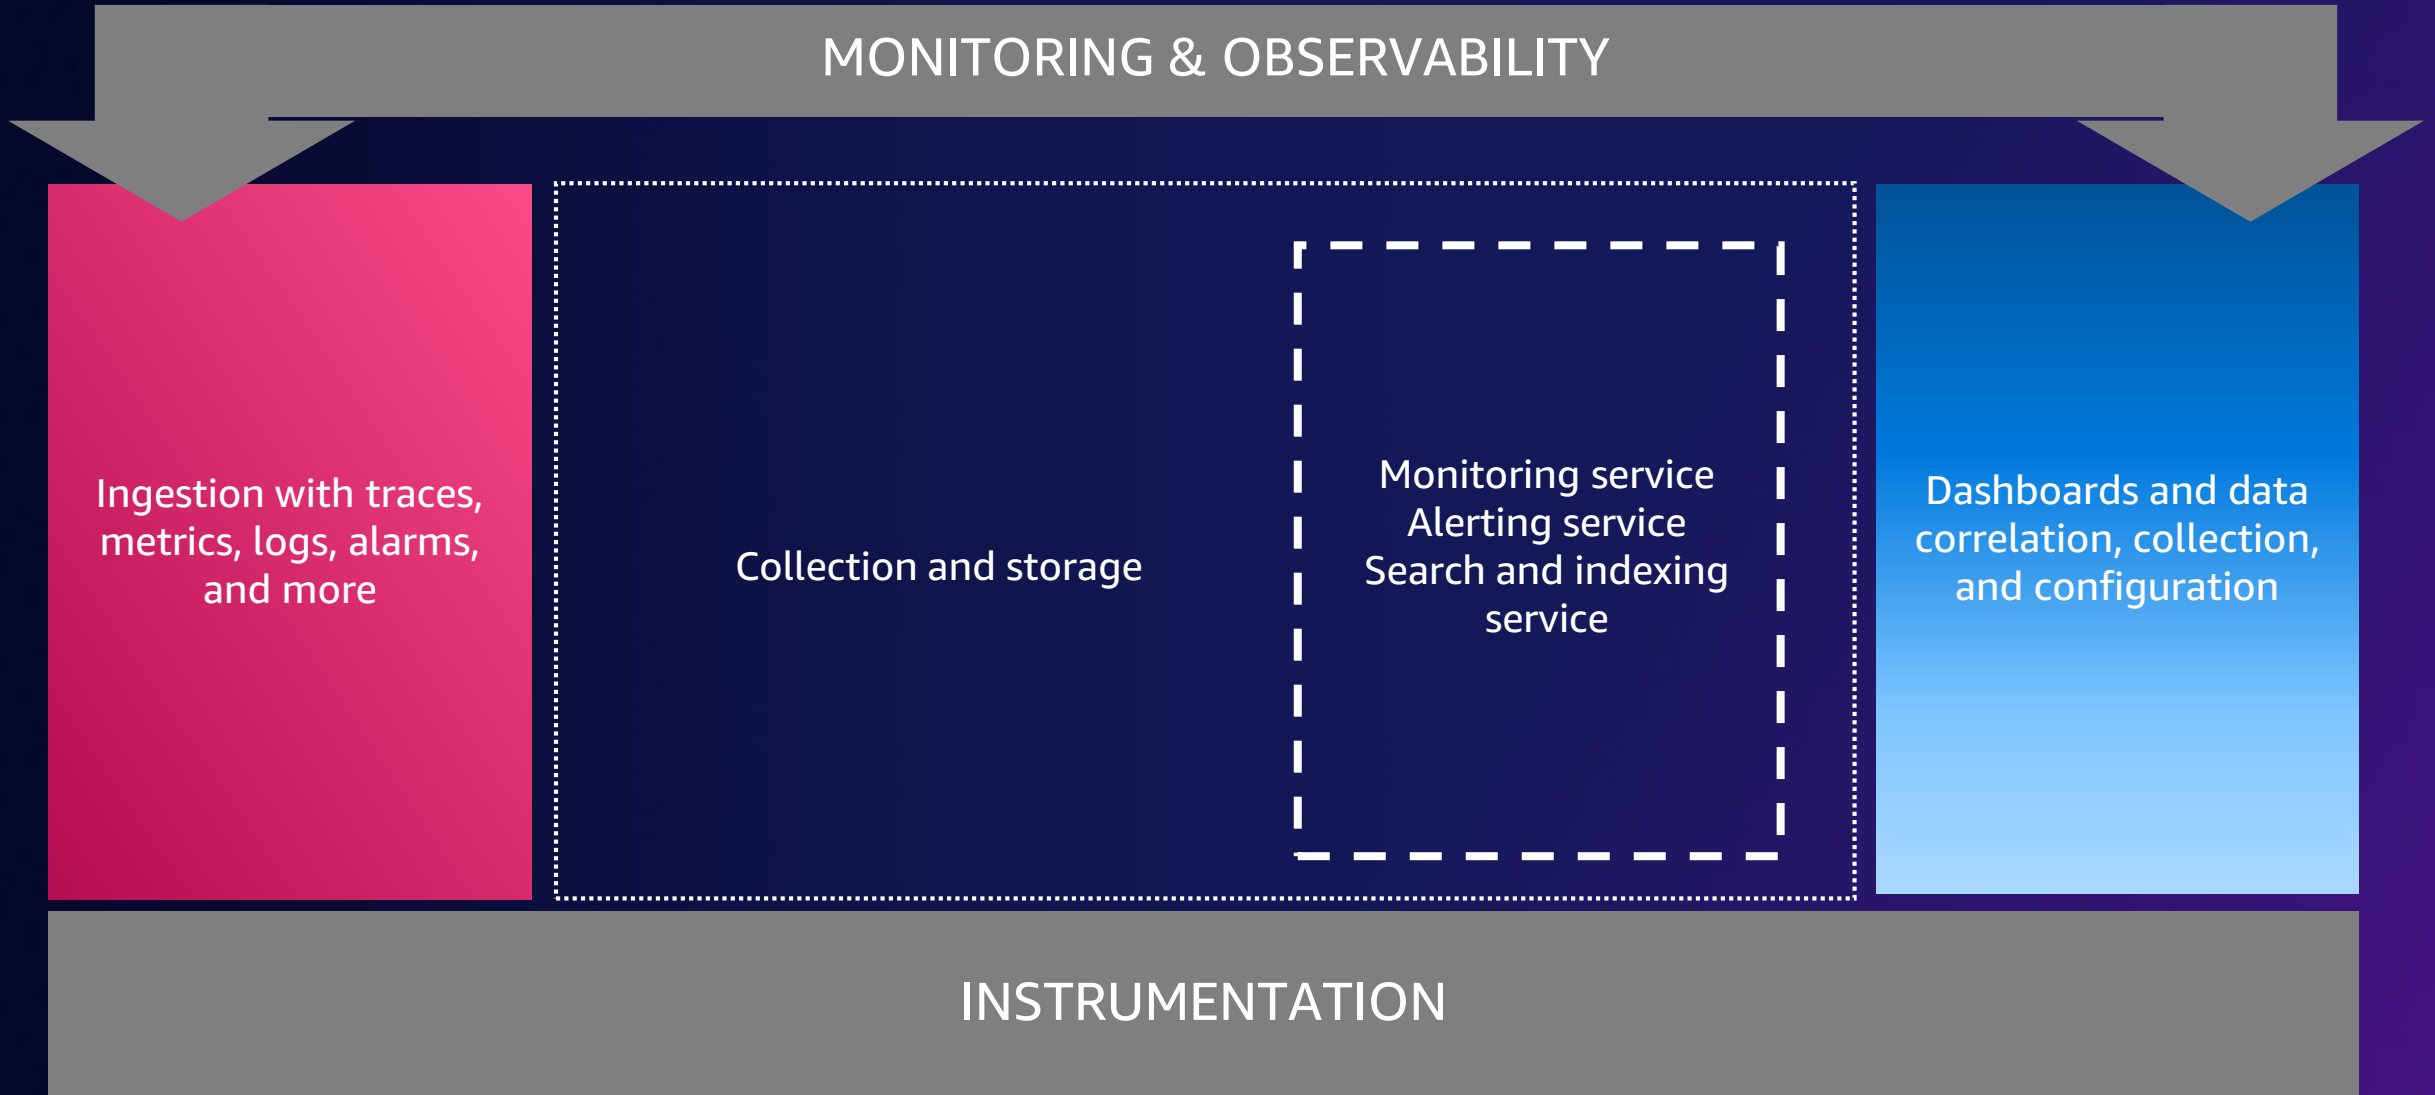


Tenants are isolated based on tenant characteristics

Observability in Amazon OpenSearch Service



The observability model: Overview

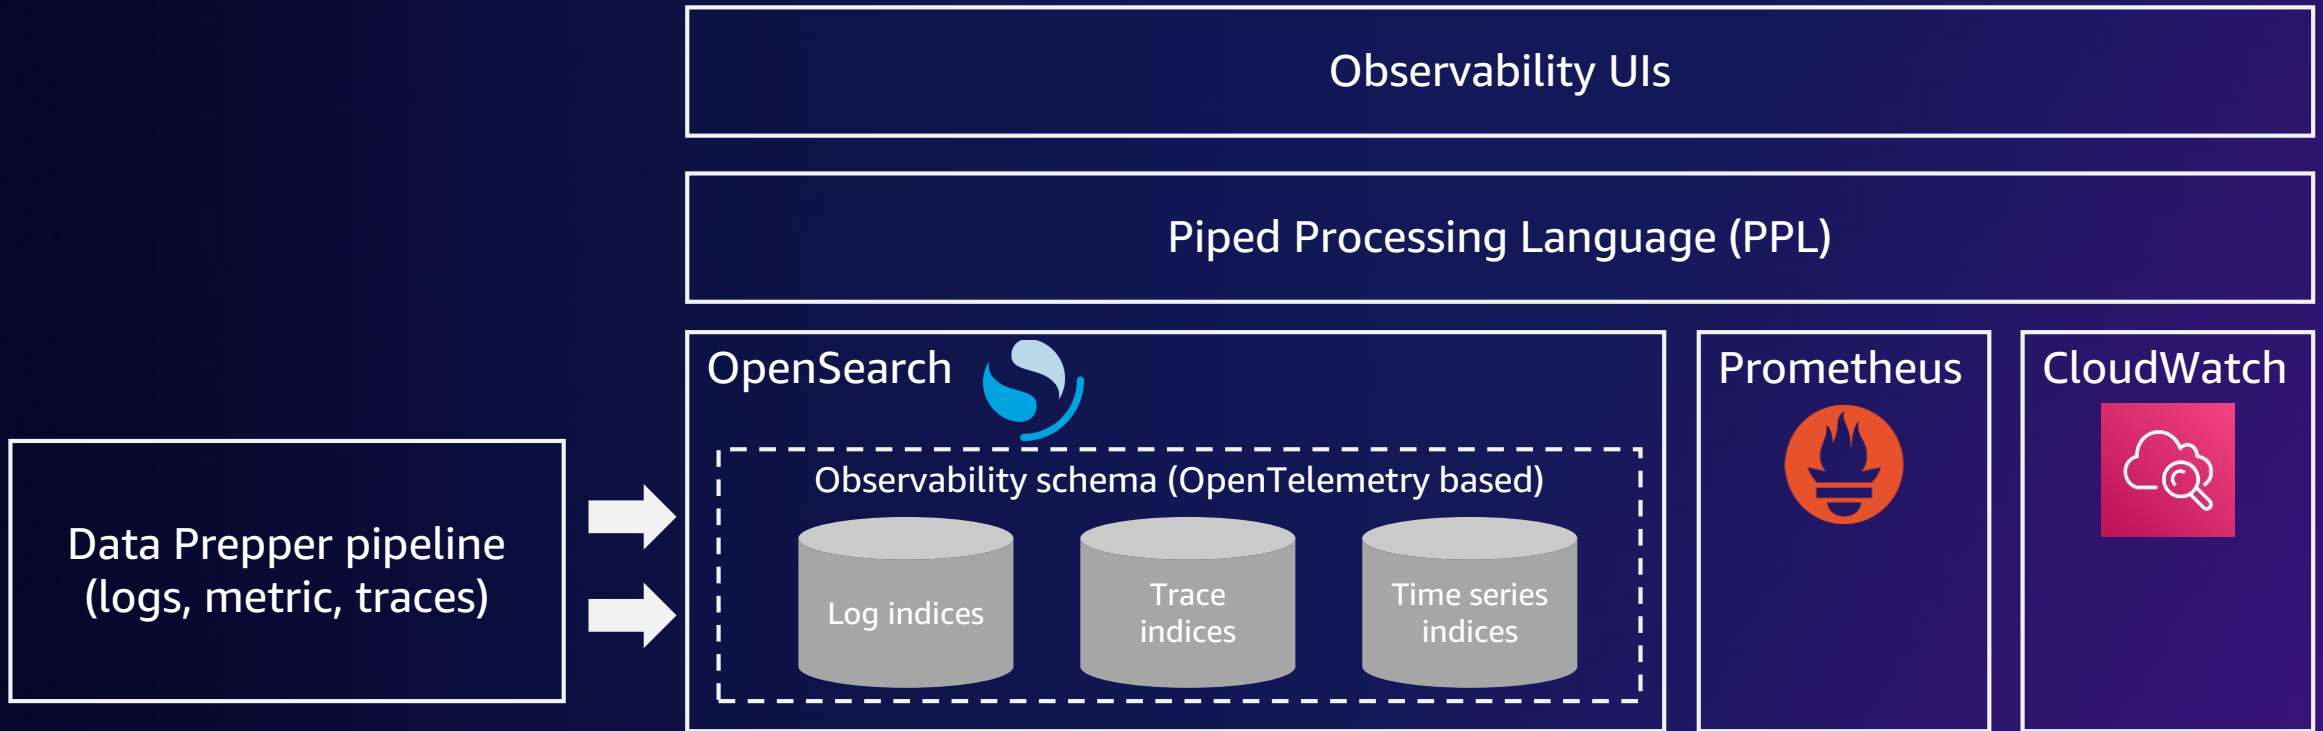


The observability model: Details

MONITORING & OBSERVABILITY



Observability implementation overview

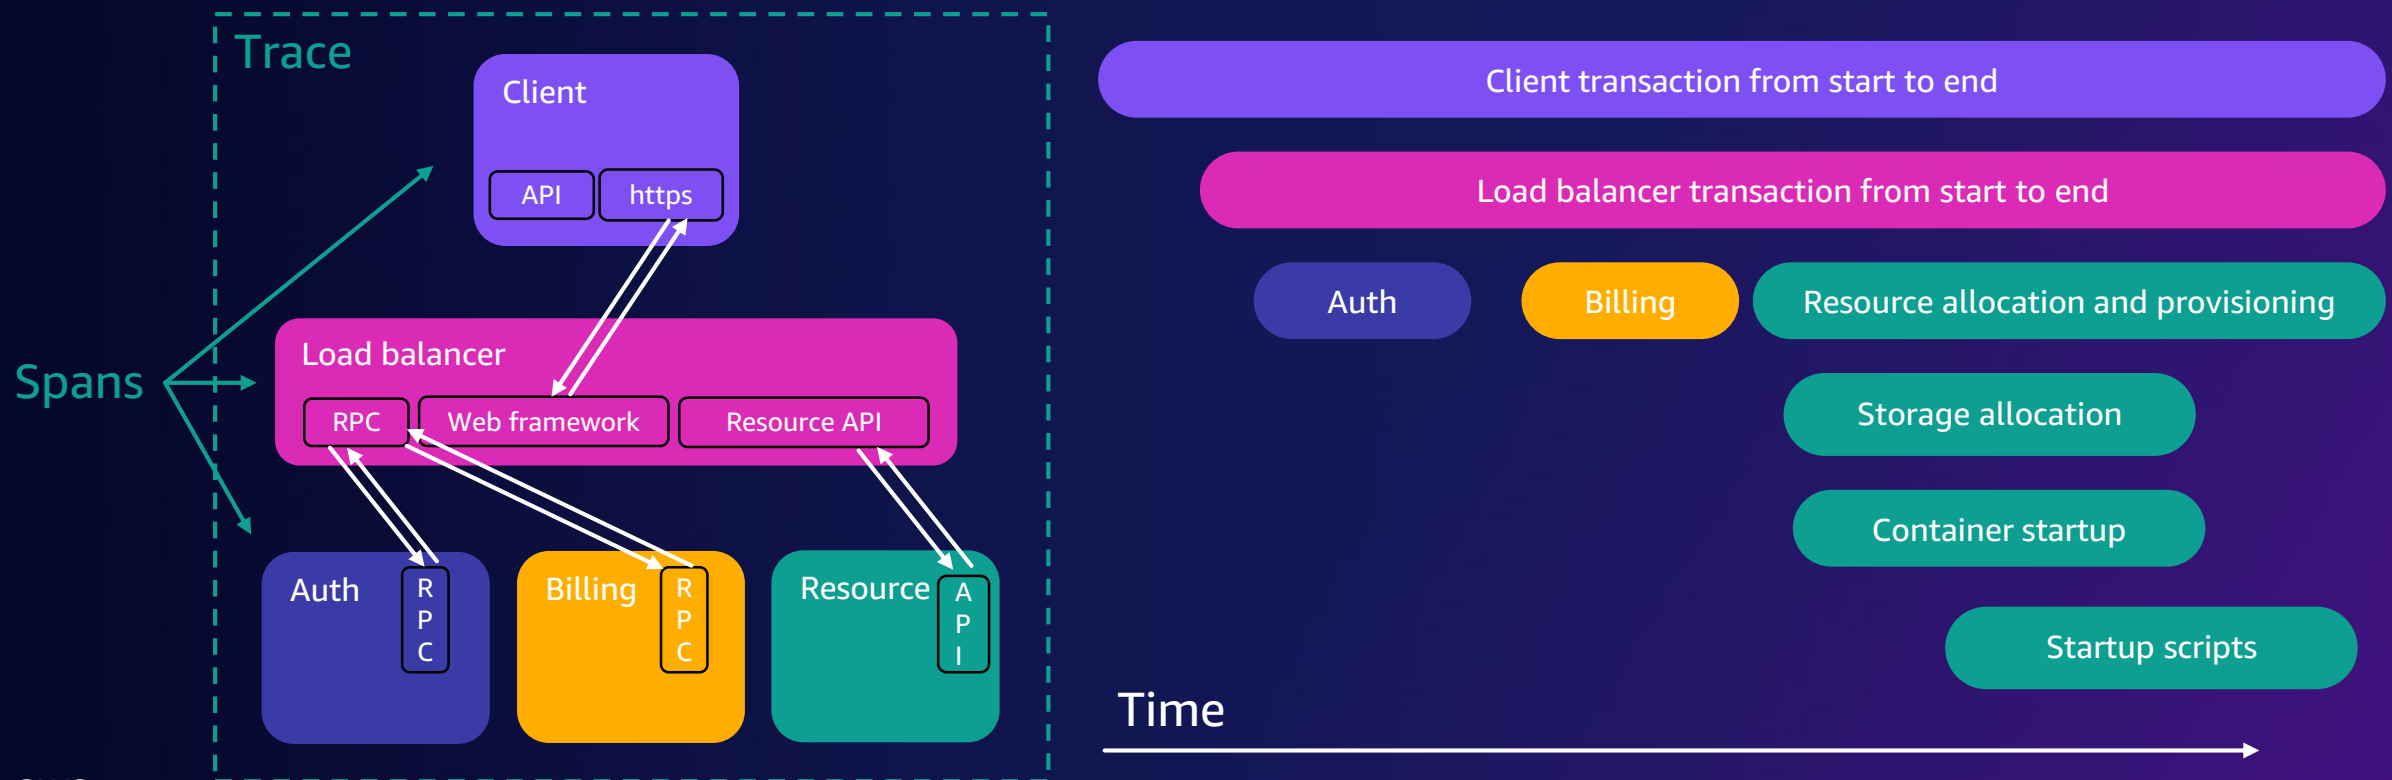


What is distributed tracing?

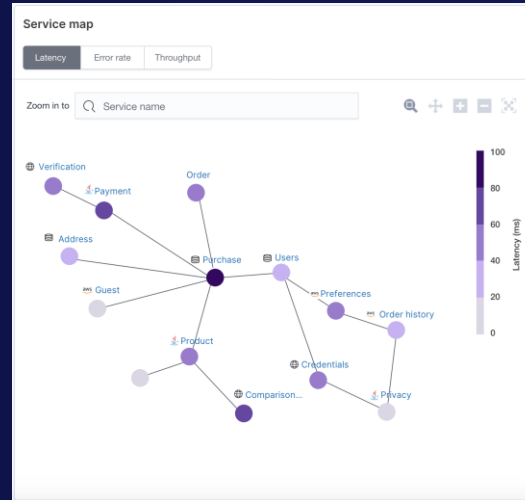
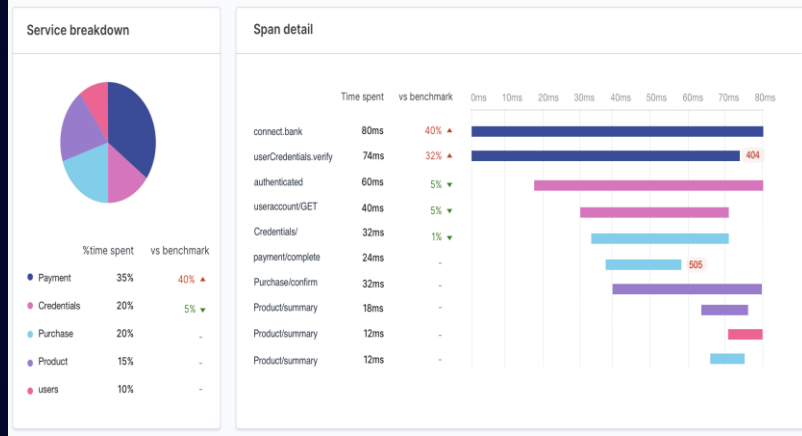
IDENTIFYING PROBLEMS IN CLOUD APPLICATIONS

A method of observing requests as they propagate through distributed systems

Trace: Hierarchical, end-to-end record of processing a request



Trace analytics



Latency by trace group

< 95th percentile

>= 95th percentile

Benchmark

This time last week

▼

Trace group name	Latency variance ↓	Average latency (ms)	Average latency vs benchmark	24-hour latency trend	Error rate	Traces
	020406080					
MakePayment.auto	<div><div></div><div></div></div>	45	30% ▲	<div><div></div><div></div></div>	20%	1,500
Order.confirmation	<div><div></div><div></div></div>	48	5% ▼	<div><div></div><div></div></div>	1%	2,000
MakePayment.oneoff	<div><div></div><div></div></div>	42	30% ▲	<div><div></div><div></div></div>	2%	1,200
Product.comparison	<div><div></div><div></div></div>	40	5% ▼	<div><div></div><div></div></div>	3%	1,000
Purchase.boynow	<div><div></div><div></div></div>	60	30% ▲	<div><div></div><div></div></div>	3%	800
MakePayment.auto	<div><div></div><div></div></div>	46	30% ▲	<div><div></div><div></div></div>	2%	900
Order.confirmation	<div><div></div><div></div></div>	64	15% ▼	<div><div></div><div></div></div>	0%	200
MakePayment.oneoff...	<div><div></div><div></div></div>	65	30% ▲	<div><div></div><div></div></div>	10%	400
Product.comparison...	<div><div></div><div></div></div>	43	10% ▼	<div><div></div><div></div></div>	10%	100
Purchase.boynow...	<div><div></div><div></div></div>	28	10% ▼	<div><div></div><div></div></div>	10%	1,100

Rows per page: 10 ▼

<

1

2

3

4

>

Trace-span details

- Single request performance
- Diagnose root cause

Service maps

- End-to-end view
- Isolate issues to services

Trace groups

- Monitor performance
- Identify issues early

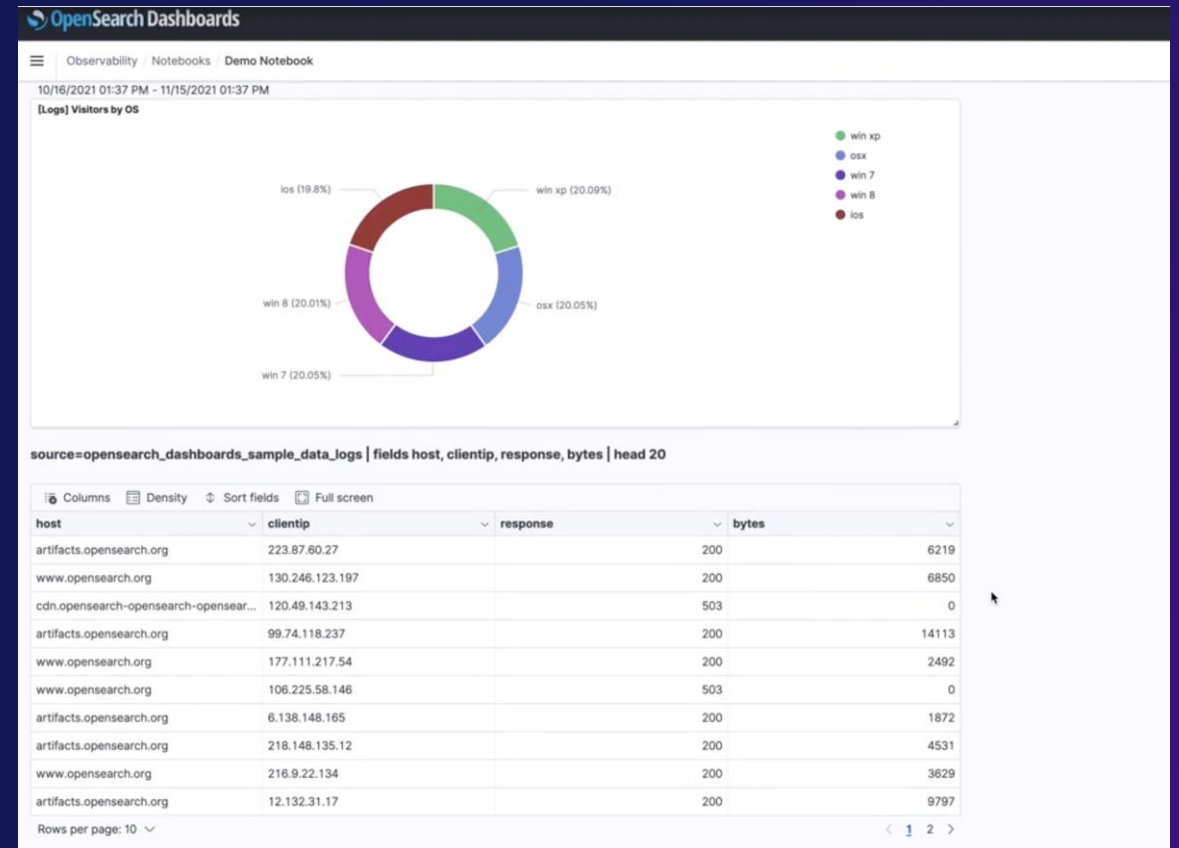
Storytelling with notebooks

A notebook is a document made up of cells or paragraphs that can combine markdown, SQL/PPL queries, and visualizations

Support for multi-timelines so that users can easily tell a story

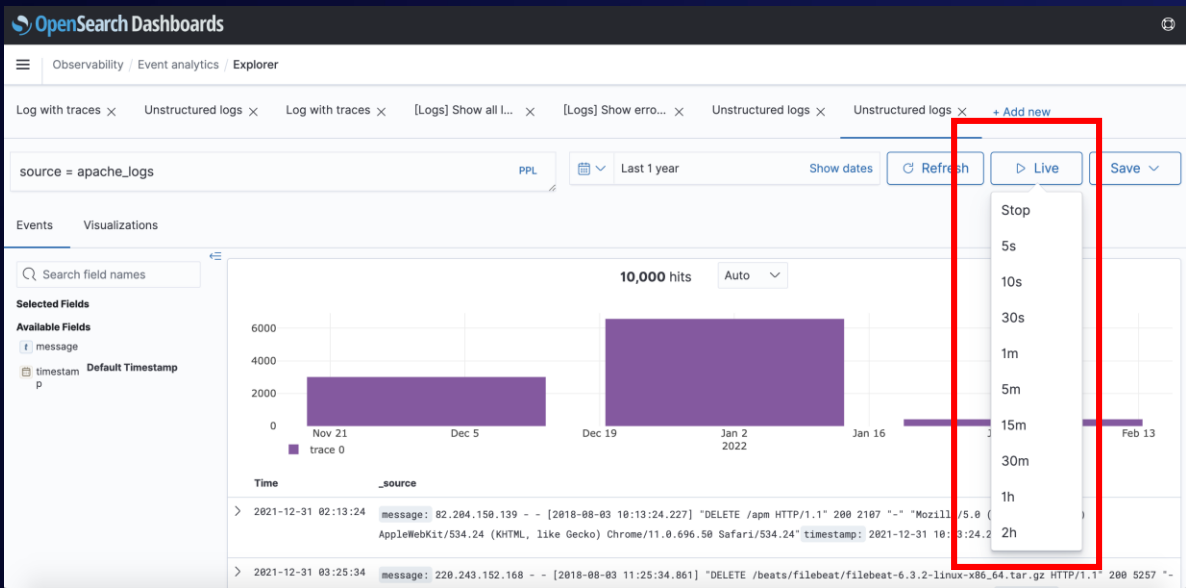
Can be shared as an OpenSearch Dashboards link, PDF, or PNG

- Common use cases include creating postmortem reports, designing runbooks, building live infrastructure reports, or even documentation

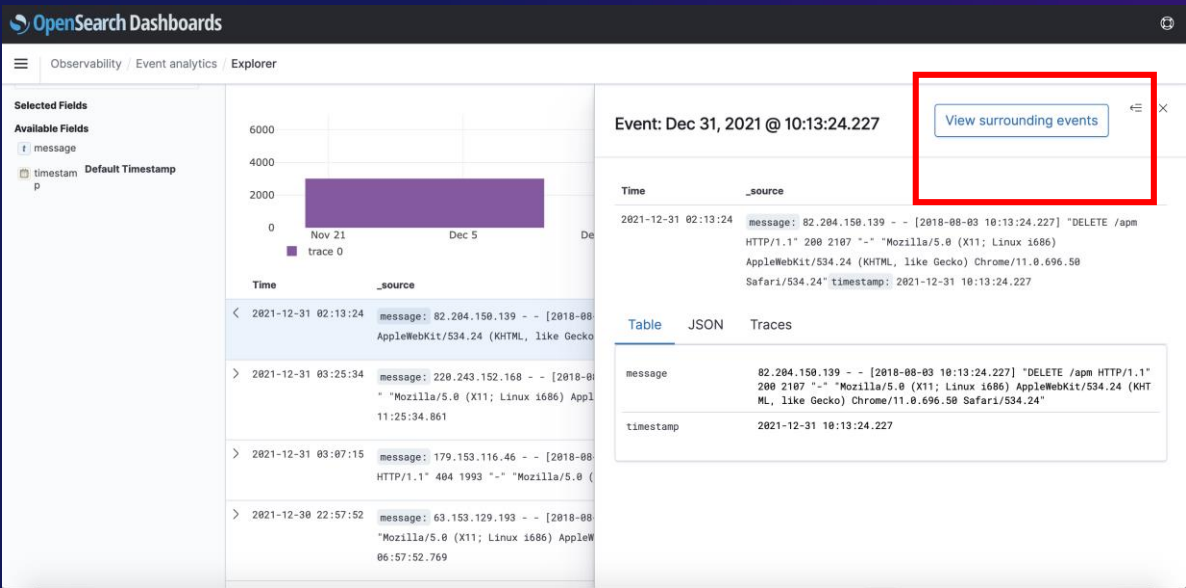


Log monitoring

Live tail



Log surround



Customer examples

Amazon OpenSearch Service customers

Software and internet



Education technology



BioTech and pharma



Financial services



Media and entertainment



Social media



Telecommunications



Travel and transportation



Real estate



Logistics and operations



Publishing



Other



Use case #1

Observability



Pinterest case study

(<https://go.aws/3wqGToQ>)

Why observability?

You need to correlate logs, metrics, and traces to gain insights into application health and performance and resolve issues across the business

How Amazon OpenSearch Service can help

Centralizes log analytics to identify or predict performance problems across your business. With cross-cluster search, you can analyze and query all of your log data via a single OpenSearch Dashboards interface

Use case #2

Application & infrastructure monitoring

Why application and infrastructure monitoring?

You need to proactively monitor your applications and infrastructure log data to find performance issues faster and improve operational health



Autodesk case study

(<https://go.aws/3cmnHSu>)

How Amazon OpenSearch Service can help

Provides real-time search and log analytics capabilities to identify or predict performance problems and enable your teams to do real-time root cause and forensic analysis, therefore reducing mean time to detect (MTTD) and mean time to resolve (MTTR) issues

Use case #3

Search

COMPASS

[Compass blog](#)

(<https://go.aws/3CuNB00>)

Why search?

You need a fast search experience for your applications, websites, and data lake catalogs, allowing your users to quickly find relevant data

How Amazon OpenSearch Service can help

Delivers high-quality and personalized search results to customers. You get access to all of Elasticsearch's search APIs, supporting natural language search, auto-completion, faceted search, adjustable ranking, and location-aware search

Use case #4

Security monitoring



Pearson case study

(<https://go.aws/3QTINqg>)

Why security monitoring?

You need to keep your data safe, preventing security threats such as data breaches, unauthorized login attempts, DoS attacks, and fraud

How Amazon OpenSearch Service can help

Accelerate security incident detection, forensic analysis, and response by being able to quickly analyze logs from disparate applications and systems across your network

Thank you!

Jon Handler

handler@amazon.com

Muthu Pitchaimani

muthupmk@amazon.com

